

# Oracle SQL Tuning

3 parts (BASIC, INTERMEDIATE, ADVANCED)

[www.lingaro.com](http://www.lingaro.com)



# Training Agenda

- **BASIC**
  - SQL Tuning Introduction
  - Instrumentation & Rewrite SQL Text
  - SQL Plan Reading
- **INTERMEDIATE**
  - Optimizer & Statistics
  - SQL Plan Tuning
  - SQL Plan Operations
- **ADVANCED**
  - SQL Processing & Cursor Sharing
  - Advanced Plan Operations
  - SQL Plan Transformations
  - SQL Plan Management

# Training Agenda

- Theory
- Quiz Comp
- Workshop
- Workshop Comp
- Winner Awards On WINS

# Oracle SQL Tuning

BASIC

[www.lingaro.com](http://www.lingaro.com)

# Training Agenda

- SQL Tuning Introduction
- Instrumentation & Rewrite SQL Text
- SQL Plan Reading

# Topic Agenda

## SQL Tuning Introduction

- Dynamic Performance Views - V\$
- Automatic Workload Repository - AWR
- Performance Metrics
- Time Model
- Wait Events
- Active Session History - ASH
- SQL Monitoring
- Optimizer Hints
- Bind Variables

# Dynamic Performance Views

## V\$ Overview



- Based on x\$ dynamic tables (based on SGA)
- Listed in v\$fixed\_table view
- Contains performance statistics (not only)
- Not guarantee read consistency
- Source for almost all performance tools
- Owned by SYS user (real names v\_\$)
- Public synonyms with v\$ names
- SELECT\_CATALOG\_ROLE needed to access all
- Many are usable during SQL tuning
- Use gv\$ in RAC database - union from all instances c\$

# Dynamic Performance Views

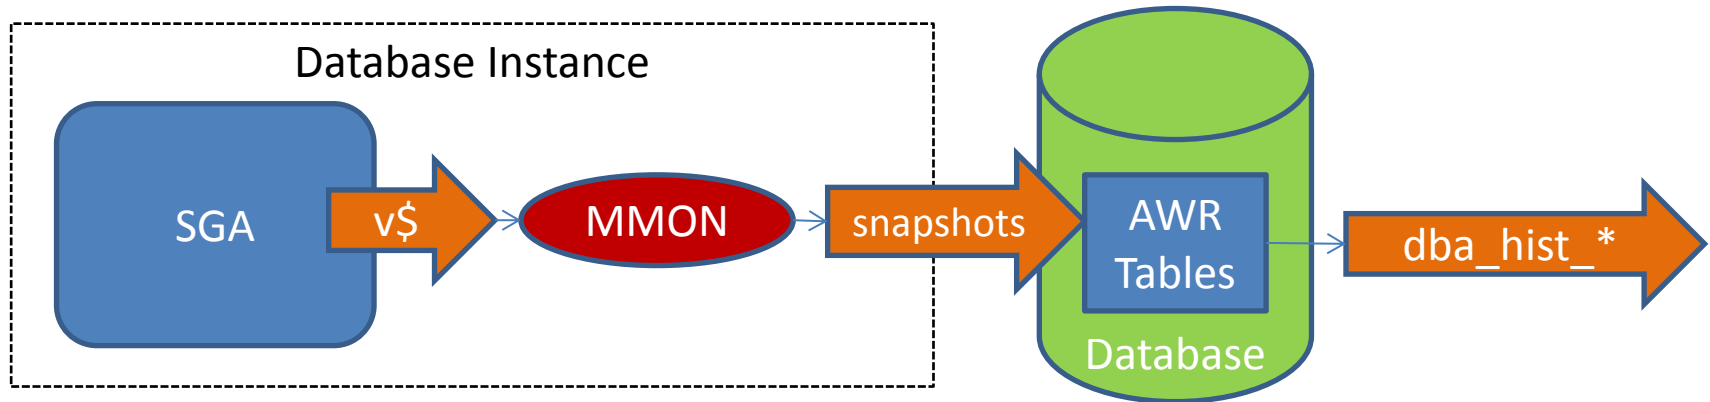
## V\$ Examples

- v\$sqltext, v\$sqltext\_with\_new\_lines - sql text
- v\$sql\_bind\_capture - bind variables definition and values
- v\$sqlarea, v\$sqlstat - cursor (statement) statistics
- v\$sql - child cursor (cursor version/plan) statistics
- v\$sql\_plan - execution plan operations
- v\$session - detailed sessions description
- v\$sesstat - sessions statistics
- v\$mystat - current session statistics
- v\$service - active database services
- v\$service\_stats - services statistics
- v\$instance - one row about database instance
- v\$sysstat - instance statistics
- v\$osstat - operating system statistics



# Automatic Workload Repository

## AWR Overview



```
SELECT snap_interval, retention, topnsql FROM dba_hist_wr_control;
```

	SNAP_INTERVAL	RETENTION	TOPNSQL
Default	+00 01:00:00.000000	+08 00:00:00.000000	DEFAULT
ADW	+00 00:30:00.000000	+120 00:00:00.000000	DEFAULT

Number of top SQL captured for each criteria  
(Elapsed Time, CPU Time, Parse Calls, -- Shareable Memory, Version Count).  
DEFAULT = 30 for TYPICAL and 100 for ALL statistics level

## Dictionary views

dba\_hist\_sqltext  
dba\_hist\_sqlbind  
dba\_hist\_sqlstat  
dba\_hist\_sql\_plan

dba\_hist\_sess\_time\_stats  
dba\_hist\_sessmetric\_history  
dba\_hist\_service\_stat  
dba\_hist\_database\_instance  
dba\_hist\_sysstat  
dba\_hist\_sysmetric\_history  
dba\_hist\_osstat

SQL Tuning Introduction

[www.lingaro.com](http://www.lingaro.com)

# Performance Metrics



## Overview

- Calculated automatically from base stats by MMON
- Calculated as delta in time - rate of change
- E.g. “Physical Reads Per Sec”
- Available in
  - v\$metric - most recent values (only few metrics)
  - v\$metricgroup - calculation intervals
  - v\$metric\_history - one hour history (60 and 15 seconds intervals)
  - v\$sessmetric
  - v\$servicemetric
  - v\$servicemetric\_history
  - v\$sysmetric - many metrics
  - v\$sysmetric\_history
  - dba\_hist\_sessmetric\_history
  - dba\_hist\_sysmetric\_history

# Time Model

## Overview



- Database Time
  - Sum of all sessions time in ACTIVE state
  - Consists of processing (CPU) time and wait time (but no idle time)
- Time model
  - Time hierarchy with DB Time as root
  - Describing where time is used in database
  - Enable to drill down into most time consuming operations
  - Leaf can help to find very detailed performance problem reason
  - Leaf e.g. “hard parse (sc) time”

```
v$sess_time_model  
v$sys_time_model  
dba_hist_sys_time_model
```



STAT_NAME	VALUE
DB time	14093932161
DB CPU	4895968750
parse time elapsed	2584226426
hard parse elapsed time	1769777748
sql execute elapsed time	14278200204
connection management call elapsed time	92870541
failed parse elapsed time	121601539
failed parse (out of shared memory) elapsed time	0
hard parse (sharing criteria) elapsed time	732782215
hard parse (bind mismatch) elapsed time	3558585

# Wait Events



## Overview

- Occurred when active session is blocked by another process
    - See dba\_waiters, dba\_blockers, v\$lock, v\$locked\_object, v\$access, v\$wait\_chains
  - v\$session view display current (if in progress) or last wait
    - status - ACTIVE or INACTIVE
    - state - WAITING or WAITED KNOWN TIME
    - event - wait event name
    - wait\_time - last wait time or 0 if currently waiting [1/100 sec]
    - seconds\_in\_wait
      - If WAIT\_TIME = 0, then SECONDS\_IN\_WAIT is the seconds spent in the current wait
      - If WAIT\_TIME > 0, then SECONDS\_IN\_WAIT is the seconds since the start of the last wait
      - SECONDS\_IN\_WAIT - WAIT\_TIME / 100 is the active seconds since the last wait ended
    - p1text, p1, p1raw, p2text, p2, p2raw, p3text, p3, p3raw - wait parameters (v\$event\_name)
  - When event is finished time is added to aggregates v\$ stats in SGA
    - v\$waitstat
    - v\$waitclassmetric
    - v\$waitclassmetric\_history
    - v\$service\_wait\_class
    - v\$session\_wait
    - v\$session\_wait\_class
    - v\$session\_wait\_history
    - v\$system\_wait\_class
    - v\$session\_event
    - v\$service\_event
    - v\$system\_event
    - v\$eventmetric
- and later written to AWR snapshots
- dba\_hist\_system\_event
  - dba\_hist\_service\_wait\_class
  - dba\_hist\_waitclassmet\_history
  - dba\_hist\_waitstat

# SQL Monitoring

## Overview



- Real time statistic for statements during execution
- Active execution plan with current operation indication
- Real time statistics for operations in plan
- By default parallel and longer then 5 sec SQLs are monitored
- Enabled/disabled by hints /\*+ MONITOR \*/ /\*+ NOMONITOR \*/
- Results visible in
  - v\$sql\_monitor
  - v\$sql\_plan\_monitor
  - In 11r2 are not available in AWR
  - Flash/html/text monitoring report for single SQL

```
SELECT dbms_sqltune.report_sql_monitor( sql_id => '526mvccm5nfy4',  
                                       type => 'ACTIVE', -- TEXT, HTML, XML  
                                       report_level => 'ALL') AS report  
  
FROM dual;
```

# Active Session History

## ASH Overview








- **Every second** all active sessions are sampled
  - Session info, stats and wait event info from v\$session are gathered
  - Only here single event parameters values are preserved (not like in agg)
- ASH visible in
  - SGA - v\$active\_session\_history - most recent samples
  - AWR - **dba\_hist\_active\_sess\_history** - **only 10%** of samples from SGA
  - Same columns, e.g. **session\_state** - "ON CPU" or "WAITING"
  - Not contains exact results (only samples)
  - Columns useful during SQL tuning

```
SAMPLE_TIME SESSION_ID SESSION_SERIAL# USER_ID SESSION_STATE
SQL_ID SQL_OPNAME SQL_PLAN_HASH_VALUE SQL_PLAN_LINE_ID
SQL_EXEC_ID SQL_EXEC_START PLSQL_OBJECT_ID PLSQL_SUBPROGRAM_ID
EVENT WAIT_TIME TIME_WAITED BLOCKING_SESSION BLOCKING_SESSION_SERIAL#
CURRENT_OBJ# IN_PARSE PROGRAM MODULE ACTION MACHINE
TM_DELTA_TIME TM_DELTA_CPU_TIME TM_DELTA_DB_TIME [microsec]
```

# SQL Monitoring

## Lingaro SQL Monitor

- Java tool used to monitor SQL in Oracle DB
- Available in <http://wiki.lingaro.local/display/PG/Oracle+SQL+Monitor>
- Lists monitored statements

STATUS	USERNAME	SQL	EXECUTION_...	SQL_ID
Executing ▼	▼	▼	▼	
	ADWG_IZOOMT2_BI	select /*+ parallel(4) */ ""    sc.node_trim_val    ""	2015-12-25	2xkyp743wz8up
	ADWGP_CBLS_RPT	select /*+ PARALLEL(2) */ T4269346.TIME_NAME as	2015-12-27	bu5jq5k5fu7wn
	ADWG_IZOOMT2_BI	select /*+ parallel(4) */ ""    sc.node_trim_val    ""	2015-12-23	4k59hd49n2w7h
	ADWGP_CBLS_RPT	select /*+ PARALLEL(2) */ T4269278.TIME_NAME as	2015-12-27	aujpwk322h5mp
	ADWGP_CBLS_RPT	select /*+ PARALLEL(2) */ T4269278.TIME_NAME as	2015-12-27	9x4tkus6jspp1

- Generates ACTIVE/HTML/TEXT monitoring report on demand
- Requires access to many v\$ / gv\$ views - best is select\_catalog\_role
- Version 2.0 requires:
  - java 7 JRE installed on laptop
  - Adobe Flash Player MSIE
- Example SQL Monitoring report



Example SQL  
Monitoring Report

# Optimizer Hints



## Introduction

- Used by developer to change SQL execution method (plan)
- Placed in comment after statement first keyword
- Hint comment have to start from + (plus) character

- Examples

```
SELECT /*+ USE_HASH(f d) FULL(f) SWAP_JOIN_INPUTS(f) */ ...  
INSERT /*+ APPEND PARALLEL(4) */ INTO ...
```

- Possible strong influence on performance
- Using hints is kind of hardcoding
- Is not recommended - use as last resort
- Instead use:
  - Fresh optimizer statistics
  - SQL Profile
  - SQL Plan Baseline



# Topic Agenda

## Tuning Techniques (not prepared yet)

- Using Automatic SQL Tuning
- Using SQL Tuning & SQL Access Advisors (**dbms\_sqltune** or OEM)
- Using Oracle Enterprise Manager Console OEM
- **Using DBA\_HIST and V\$ performance views (described here)**
- Using Test Executions
- Using SQL Decomposition
  - Test Execution on SQL parts

# Q & A

## + Quiz

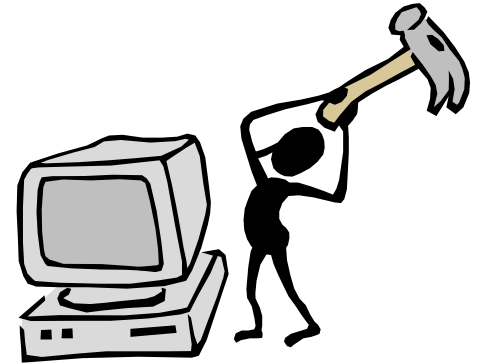
# Introduction Workshop

- Login to Enterprise Manager

<https://<virtual host IP>:1158/em/>

User: sys Pass: oracle (as sysdba)

default IP = 192.168.56.101



- Enter „Performance” page (and later) „SQL Monitoring” page
- Use script `SQL_Tuning_BASIC.sql` in Sqldeveloper
- Execute workload `sh_workload_pro(TRUE)`, make observations in EM
- Execute query on Dynamic Performance Views
- Execute query on `v$mystat` before and after workload on your session using

```
exec SH_WORKLOAD_PRO(TRUE,1,'prod_categ_pro',false);  
exec SH_WORKLOAD_PRO(TRUE,1,'chanl_class_pro',false);
```

# Introduction

## Workshop

### Performance Reporting Examples

```
SELECT /*+ NO_MERGE(s) */ sid, serial#, username,
      plsql_entry_object_id, object_type, object_name,
      sql_id, sql_text
FROM (SELECT sid, serial#, sql_id, username, plsql_entry_object_id, sql_text
      FROM v$sqlsession NATURAL JOIN v$sqltext WHERE piece = 0
      ) s JOIN dba_objects o ON o.object_id = plsql_entry_object_id;

SELECT COUNT(*) samples, sql_id FROM v$active_session_history WHERE session_state = 'ON CPU'
      AND user_id = (SELECT user_id FROM all_users WHERE username = 'TRAIN')
GROUP BY sql_id ORDER BY samples DESC;

SELECT COUNT(*) samples, session_state, event, in_parse
FROM v$active_session_history h WHERE h.sql_id = 'fdd01xqmvtkat'
GROUP BY h.sql_id, session_state, event, in_parse ORDER BY samples DESC;

SELECT h.*, sql_text FROM (
      SELECT snap_id, begin_interval_time,
            COUNT(*) samples, user_id, sql_id, in_parse
      FROM dba_hist_active_sess_history NATURAL JOIN dba_hist_snapshot
      WHERE sql_id = 'b1thg2wm16t5b'
      GROUP BY sql_id, user_id, in_parse, snap_id, begin_interval_time
      ) h JOIN dba_hist_sqltext t ON h.sql_id = t.sql_id
ORDER BY samples DESC;
```

# Q & A

## + Workshop Comp

# Topic Agenda

## Instrumentation & Rewrite SQL Text

- Statistic Level & Module/Action Info
- Critical Server Resource
- Resource Consumer SQL
- SQL Performance Issue Reasons
- Rewrite Problematic SQL Workshop

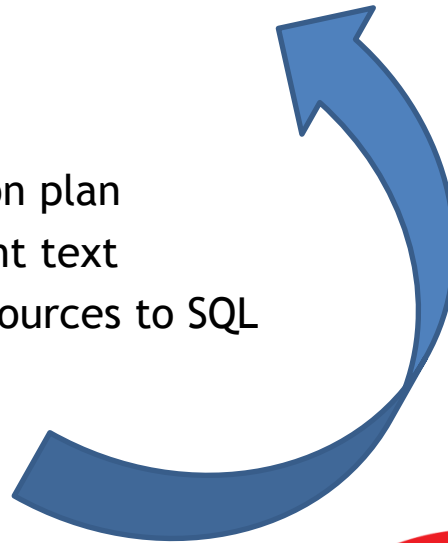
# Prerequisites

## SQL Tuning Steps

- 1. Set statistic level & module/action info (optional)
- 2. Find most critical (limited) resource in database server
  - CPU, Disk, Wait Time (in performance issue time range)
- 3. Find top SQL statements in workload by critical resource

(Will be described later)

- 4. Analyze SQL plan operations (start from top SQL)
  - Rank operations in SQL
  - Find top cost reason
- 5. Tune SQL by
  - Influencing on execution plan
  - Rewriting SQL statement text
  - Adding more server resources to SQL



Instrumentation

# Prerequisites

## Optional Settings

- Set on workload session before workload (e.g. logon trigger)
- **Statistic Level**
  - Gather current statistics about row processed to execution plan by
  - Setting on session

```
ALTER SESSION SET statistic_level = ALL      -- default TYPICAL
```
  - In statement using hint

```
/*+ GATHER_PLAN_STATISTICS */
```
- **Module, Action, Client Information**
  - Mark session with module/action/client info
  - It will simplify finding interesting statements in statistics
  - It enables module/action/client rankings

```
dbms_application_info.set_module('Loading FCT', 'STEP010')
dbms_application_info.set_action('STEP020')
```



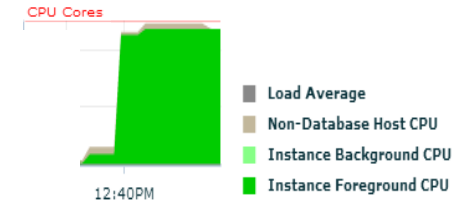
# Critical Resource

## Is CPU Overloaded?

- Near 100% of host CPU utilization ?

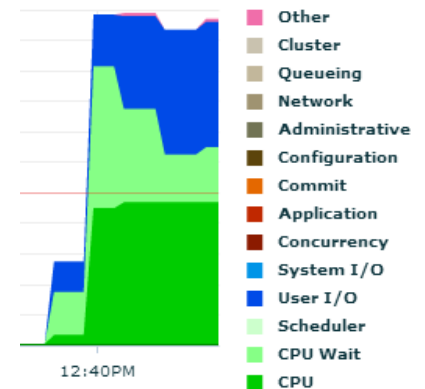
```
SELECT value os_cpu, begin_time,
       numtodsinterval(end_time - begin_time, 'DAY') INTERVAL
FROM v$sysmetric_history
WHERE metric_name = 'Host CPU Utilization (%)'
ORDER BY value desc; -- % Busy/(Idle+Busy)
```

```
SELECT round(b.VALUE/(i.VALUE+b.VALUE),3)*100 cpu_pct,
       begin_interval_time,
       end_interval_time - begin_interval_time INTERVAL
FROM dba_hist_osstat i JOIN dba_hist_osstat b USING (snap_id,dbid,instance_number)
JOIN dba_hist_snapshot USING (snap_id)
WHERE b.stat_name = 'BUSY_TIME' AND i.stat_name = 'IDLE_TIME'
ORDER BY cpu_pct DESC;
```



- Number of CPU session near number of cores ?

```
SELECT count(DISTINCT session_id) sess_cnt, sample_time
FROM v$active_session_history WHERE session_state = 'ON CPU'
GROUP BY sample_time ORDER BY sess_cnt DESC;
```



- CPU Wait is Over 10% of CPU time ?

# Critical Resource

## Is Storage Overloaded?

- IO Calibration

```
DBMS_RESOURCE_MANAGER.CALIBRATE_IO(num_physical_disks => in_dsk, max_latency => in_max_lat  
    max_iops => out_max_iops, max_mbps => out_max_mbps, actual_latency => out_actual_lat)
```

```
SELECT * FROM dba_rsrc_io_calibrate;
```

START_TIME	END_TIME	MAX_IOPS	MAX_MBPS	MAX_PMBPS	LATENCY	NUM_PHYSICAL_DISKS
15/12/29 14:34:15,995000000	15/12/29 14:41:17,066000000	2019	289	289	0	1

- Check how far is current workload from calibrated maximum

```
select round(value,2) per_sec, begin_time,  
    numtodsinterval(end_time - begin_time, 'DAY') INTERVAL  
from v$sysmetric_history  
where metric_name = 'I/O Megabytes per Second'  
order by value desc;
```

```
SELECT round(VALUE,2) req_per_sec, begin_time,  
    numtodsinterval(end_time - begin_time, 'DAY') INTERVAL  
FROM dba_hist_sysmetric_history  
WHERE metric_name = 'I/O Requests per Second'  
    AND begin_time between sysdate - 1 and sysdate  
ORDER BY VALUE DESC;
```

# TOP SQL Ranking

## By Critical Resource in Time Range

- CPU rank based on CPU time

```
SELECT h.*, sql_text FROM (SELECT MAX(cpu_time_total) cpu_time,  
                                MAX(disk_reads_total) disk_rds, sql_id  
                                FROM dba_hist_sqlstat NATURAL JOIN dba_hist_snapshot  
                                WHERE begin_interval_time > sysdate - 1 GROUP BY sql_id  
    ) h JOIN dba_hist_sqltext t ON h.sql_id = t.sql_id ORDER BY cpu_time DESC;
```

- CPU rank based on number of CPU samples in ASH

- Number of samples = seconds in v\$ (but in dba\_hist\_ seconds/10)

```
SELECT h.*, sql_text FROM ( SELECT COUNT(*) samples, user_id, sql_id  
                                FROM dba_hist_active_sess_history  
                                WHERE sample_time > sysdate - 1  
                                AND session_state = 'ON CPU' GROUP BY sql_id, user_id  
    ) h JOIN dba_hist_sqltext t ON h.sql_id = t.sql_id JOIN all_users u ON h.user_id = u.user_id  
    ORDER BY samples desc;
```

- Wait events rank for particular SQL

```
SELECT COUNT(*) samples, session_state, event, in_parse  
    FROM dba_hist_active_sess_history h  
    WHERE h.sql_id = '3f35td3yb01m6'  
    GROUP BY h.sql_id, session_state,  
            event, in_parse  
    ORDER BY samples DESC;
```

SAMPLES	SESSION_STATE	EVENT	IN_PARSE
32	ON CPU	(null)	N
25	WAITING	db file scattered read	N
6	WAITING	db file sequential read	N

# Issue Reasons

## Introduction

- Stale or missing optimizer statistics
- Missing access structures
  - Indexes, Partitions, Materialized Views ...
- Suboptimal execution plan selection
  - Even statistics are correct
- Poorly constructed SQL
  - Logical, architectural and design mistakes
  - Poopy mistakes like unwanted cross join, missed filters
  - Not optimal processing method e.g. correlated subquery on large volume
  - Costly PL/SQL functions used in SQL

# Issue Reasons

## Inefficient SQL: Examples

correlation  
on large  
volume

```
SELECT COUNT(*) FROM prod_dim p
  WHERE prod_list_price <
        1.15 * (SELECT avg(unit_cost) FROM cost_fct c
                WHERE c.prod_id = p.prod_id)
```

expression  
not column  
on left side

```
SELECT * FROM job_history jh, employees e
  WHERE substr(to_char(e.employee_id),2) =
        substr(to_char(jh.employee_id),2)
```

conversion

```
SELECT * FROM orders WHERE order_id_char = 1205
```

function  
moved left

```
SELECT * FROM employees
  WHERE to_char(salary) = :sal
```

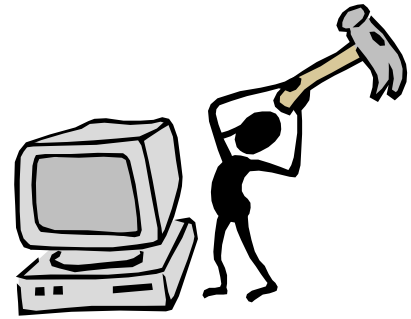
should be  
UNION ALL

```
SELECT * FROM parts_old
UNION
SELECT * FROM parts_new
```

# Q & A

## + Quiz

# Instrumentation Workshop



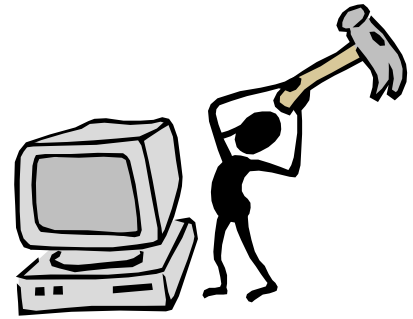
- Test `GATHER_PLAN_STATISTICS` hint
- Use module and action SQL indication
- Workshop statemets included in `SQL_Tuning_BASIC.sql`

# Q & A

## + Workshop Comp



# Rewrite SQL Text Workshop Comp



- Modify SQL text to archive better performance
- Workshop statemets included in `SQL_Tuning_BASIC.sql`

# Topic Agenda

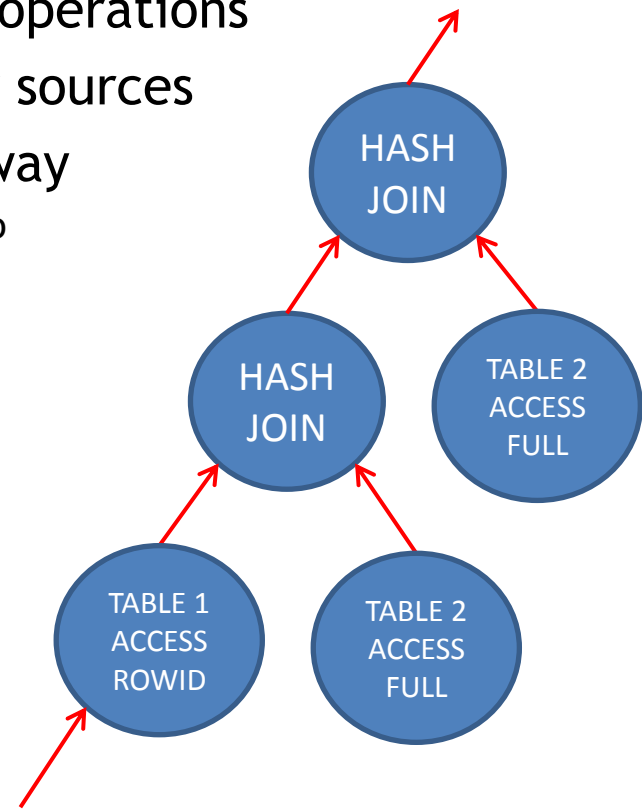
## Execution Plan Reading

- Introduction
- Methods to See Plan
- Plan Interpretation

# SQL Execution Plan

## Introduction

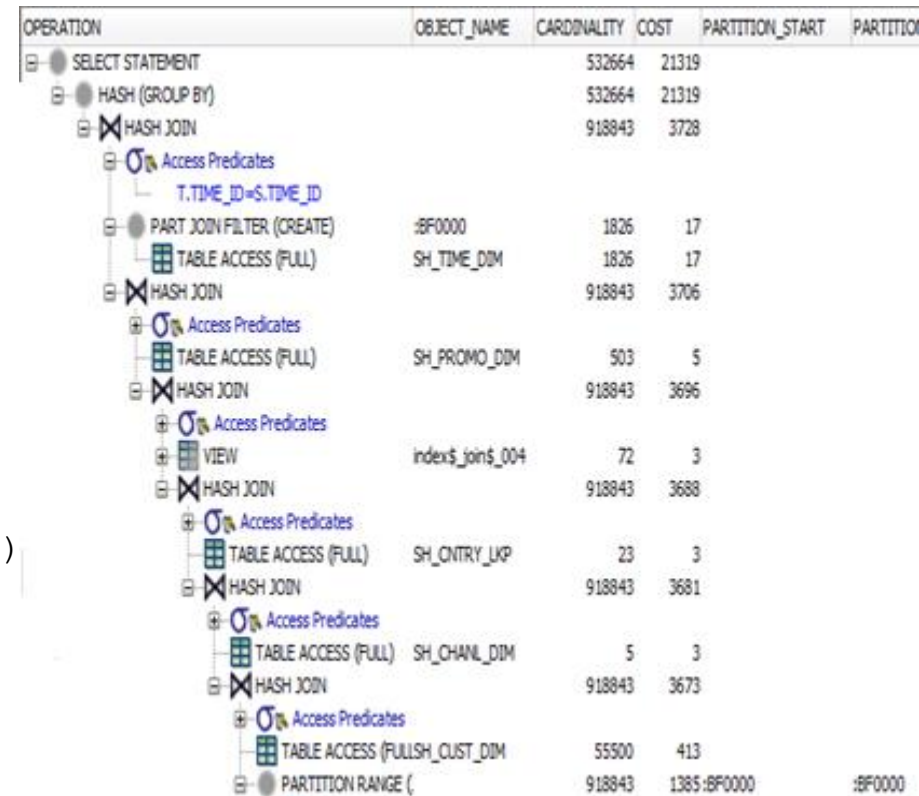
- Consists of plan **Row Source Operations** related in hierarchical form
- Child operation is **Row Source** (set of rows producer)
- Parent operations consumes rows from child operations
- JOIN parent operation receives data two row sources
- Is specification of SQL statement execution way
  - Needed because SQL code not contains all “how to do” info
  - One way from many possible ways
  - Way selected by optimizer during SQL hard parse phase
- Plan describes
  - Order of operations
  - Object access methods
  - Join methods
  - Join order
  - And more
- PL/SQL do not have execution plans
  - PL/SQL code describes how to do work itself



# See Execution Plan

## Tools & Sources

- **PLAN\_TABLE** (not really used plan)
  - EXPLAIN PLAN command and  
SELECT from PLAN\_TABLE  
or DBMS\_XPLAN.DISPLAY()
  - SQL Developer - F10
  - Autotrace
- **V\$SQL\_PLAN\_MONITOR (11g)**  
DBMS\_SQLTUNE.REPORT\_SQL\_MONITOR()
- **V\$SQL\_PLAN (Library Cache)**  
DBMS\_XPLAN.DISPLAY\_CURSOR()
- **DBA\_HIST\_SQL\_PLAN (AWR)**  
DBMS\_XPLAN.DISPLAY\_AWR()
- **SQL management base (SQL plan baselines)**  
DBMS\_XPLAN.DISPLAY\_SQL\_PLAN\_BASELINE()
- **SQL tuning set (STS)**  
DBMS\_XPLAN.DISPLAY\_SQLSET()

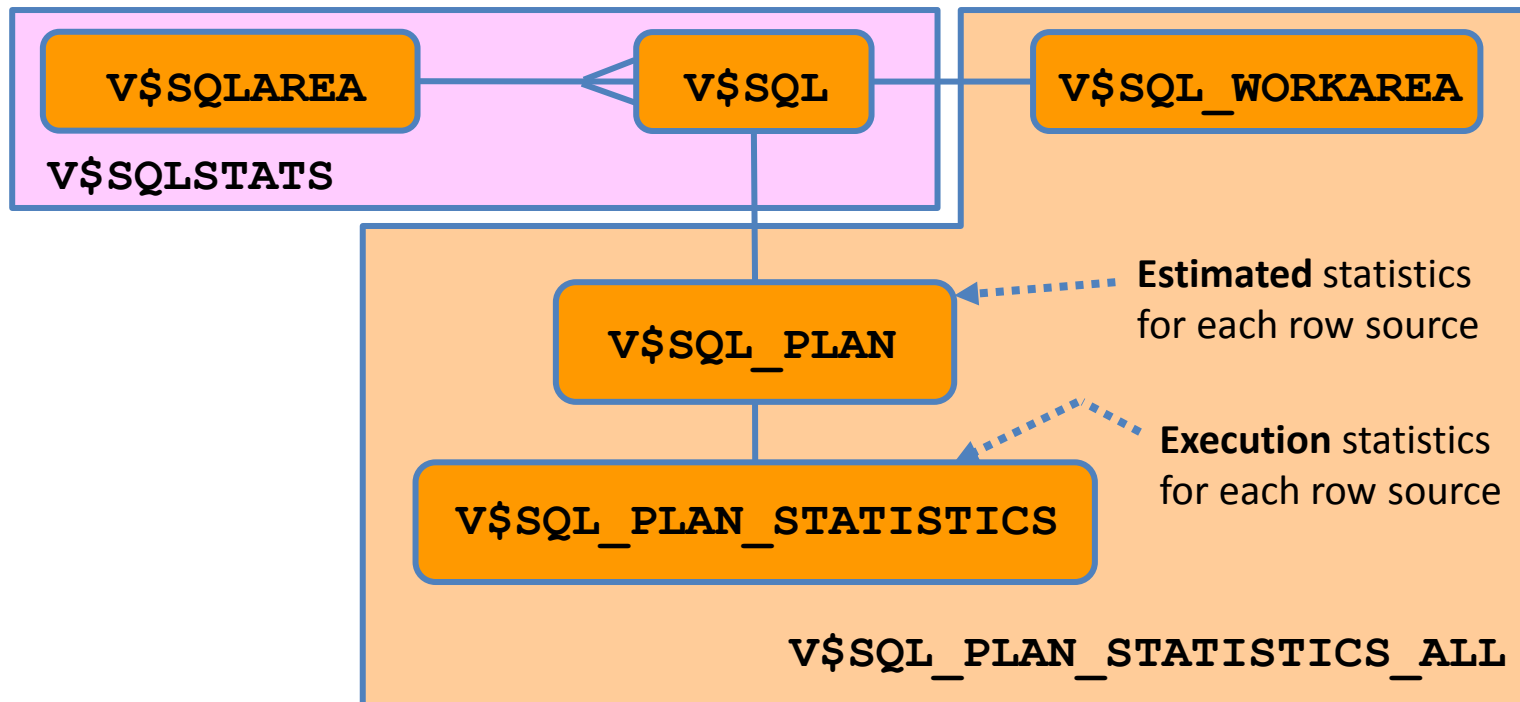


OPERATION	OBJECT_NAME	CARDINALITY	COST	PARTITION_START	PARTITION
SELECT STATEMENT		532664	21319		
HASH (GROUP BY)		532664	21319		
HASH JOIN		918843	3728		
Access Predicates					
T.TIME_ID=S.TIME_ID					
PART JOIN FILTER (CREATE)	:BF0000	1826	17		
TABLE ACCESS (FULL)	SH_TIME_DIM	1826	17		
HASH JOIN		918843	3706		
Access Predicates					
TABLE ACCESS (FULL)	SH_PROMO_DIM	503	5		
HASH JOIN		918843	3696		
Access Predicates					
VIEW	index\$_join\$_004	72	3		
HASH JOIN		918843	3688		
Access Predicates					
TABLE ACCESS (FULL)	SH_CNTRY_LKP	23	3		
HASH JOIN		918843	3681		
Access Predicates					
TABLE ACCESS (FULL)	SH_CHANL_DIM	5	3		
HASH JOIN		918843	3673		
Access Predicates					
TABLE ACCESS (FULL)	SH_CUST_DIM	55500	413		
PARTITION RANGE (		918843	1385:BF0000		:BF0000

# See Execution Plan

## Actual Execution Statistics

- `V$SQL_PLAN_STATISTICS`  
    `STATISTICS_LEVEL` set to `ALL`  
    The `GATHER_PLAN_STATISTICS` hint
- `V$SQL_PLAN_STATISTICS_ALL`
  - Enables side-by-side comparisons of estimates with the actual statistics.



# PLAN\_TABLE

- Automatically created but you can create - utlxplan.sql
- 👍 SQL is not executed. 👎 May not be the actual execution plan
- Row = one plan operation, hierarchy by ID and PARENT\_ID columns.
- Many columns are the same like in v\$sql\_plan

Statement and plan columns:

ADDRESS	Address of the handle to the parent for this cursor
HASH_VALUE	Hash value of the parent statement
SQL_ID	SQL identifier of the parent cursor in the library cache
PLAN_HASH_VALUE	Numerical representation of the SQL plan for the cursor.
CHILD_ADDRESS	Address of the child cursor
CHILD_NUMBER	Number of the child cursor that uses this execution plan.
STATEMENT_ID	Parameter specified in the EXPLAIN PLAN statement
PLAN_ID	Unique identifier of a plan in the database
TIMESTAMP	Date and time when the execution plan was generated
REMARKS	Indicate whether an outline or SQL Profile was used

Operations columns:

ID	
PARENT_ID	
DEPTH	Depth (or level) of the operation in the tree.
POSITION	Order of processing for have the same PARENT_ID
OPERATION	
OPTIONS	e.g. object access method
OBJECT_NODE	
OBJECT_OWNER	
OBJECT_NAME	
OBJECT_ALIAS	
OBJECT_INSTANCE	
OBJECT_TYPE	

Estimates columns:

CARDINALITY  
COST  
BYTES  
PARTITION\_ID  
OTHER  
OTHER\_XML  
DISTRIBUTION  
CPU\_COST  
IO\_COST  
TEMP\_SPACE  
TIME

Other columns:

OTHER\_TAG  
PARTITION\_START  
PARTITION\_STOP  
OPTIMIZER  
SEARCH\_COLUMNS  
ACCESS\_PREDICATES  
FILTER\_PREDICATES  
PROJECTION  
QBLOCK\_NAME

## Execution Plan Reading

# PLAN\_TABLE

## Inserting Plan Rows

→ **EXPLAIN PLAN** →  
    [ SET STATEMENT\_ID  
      = 'text' ]  
→  
    [ INTO *your plan table* ]  
→ **FOR *statement*** →

Example

```
SQL> EXPLAIN PLAN
      2  SET STATEMENT_ID = 'demo01' FOR
      3  SELECT * FROM emp;
```

Execution Plan Reading

# Show Plan

## from PLAN\_TABLE using DISPLAY function

```
SELECT plan_table_output FROM TABLE(  
  dbms_xplan.display(table_name => NULL, statement_id => 'demo01',  
    format => 'ADVANCED -Projection -Predicate -Aliases, filter_preds => NULL));
```

Plan hash value: 3956160932

-----								
Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time		
-----								
0	SELECT STATEMENT		1	37	3 (0)	00:00:01		
1	TABLE ACCESS FULL	EMP	1	37	3 (0)	00:00:01		
-----								

Outline Data

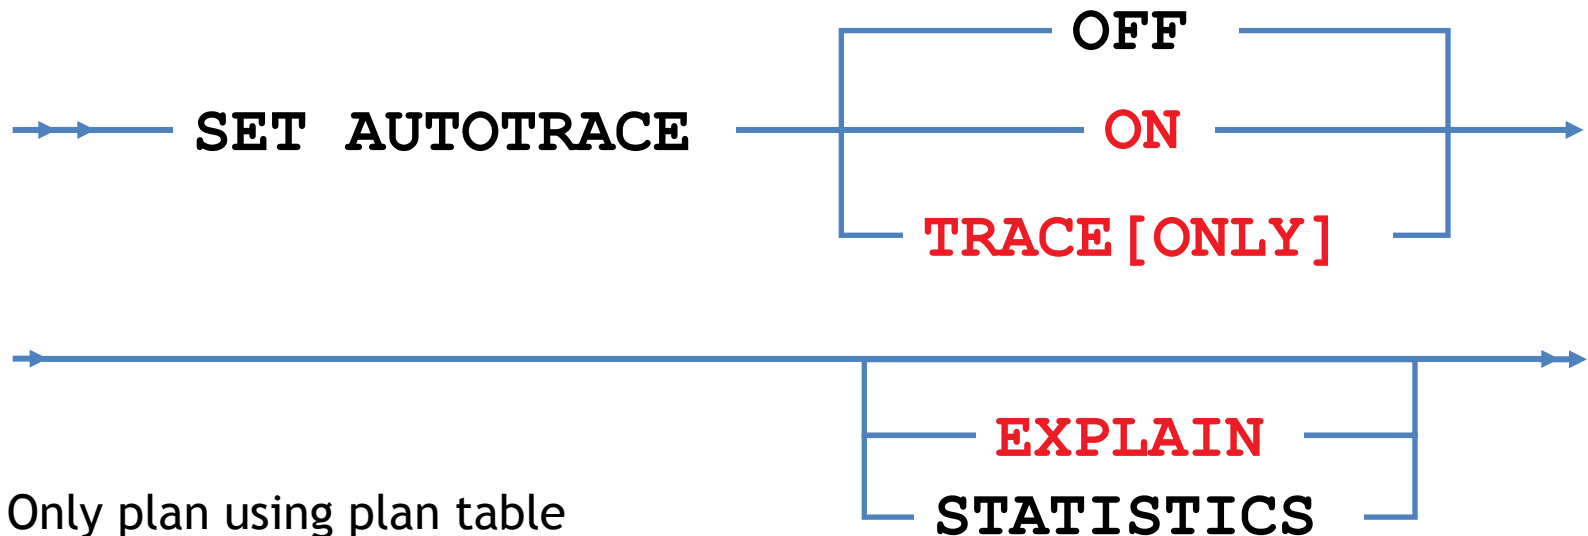
```
-----  
/*+  
  BEGIN_OUTLINE_DATA  
  FULL(@"SEL$1" "EMP"@"SEL$1")  
  OUTLINE_LEAF(@"SEL$1")  
  ALL_ROWS  
  DB_VERSION('11.1.0.6')  
  OPTIMIZER_FEATURES_ENABLE('11.1.0.6')  
  IGNORE_OPTIM_EMBEDDED_HINTS  
  END_OUTLINE_DATA  
*/
```



# Show Plan

## Autotrace

- Autotrace Is a SQL\*Plus and SQL Developer facility
- Needs a `PLAN_TABLE`
- Needs the `PLUSTRACE` role to retrieve statistics from some `V$` views
- By default, execution plan and statistics after running the query



- Only plan using plan table
  - `SET AUTOTRACE TRACE EXPLAIN`

# Autotrace

## SQL Execution Statistics

```
SQL> show autotrace
autotrace OFF
SQL> set autotrace traceonly statistics
SQL> SELECT * FROM prod_dim;
```

288 rows selected.

### Statistics

```
-----
      1334 recursive calls
         0 db block gets
       686 consistent gets
       394 physical reads
         0 redo size
    103919 bytes sent via SQL*Net to client
       629 bytes received via SQL*Net from client
        21 SQL*Net roundtrips to/from client
        22 sorts (memory)
         0 sorts (disk)
       288 rows processed
```

# DBMS\_XPLAN

## DISPLAY\_CURSOR Function Example

```
SQL> SELECT plan_table_output
      FROM TABLE(dbms_xplan.display_cursor(sql_id => '3vjxpmhhzngu4',
                                           child_number => NULL
                                           format = 'ALL +Projection'));
```

```
SQL_ID 3vjxpmhhzngu4, child number 0
```

```
-----
SELECT * FROM dual
```

```
Plan hash value: 272002086
```

```
-----
| Id  | Operation                               | Name | Rows  | Bytes | Cost (%CPU)| Time     |
-----+-----+-----+-----+-----+-----+-----+
|   0 | SELECT STATEMENT                       |      |       |       |    2  (100)|          |
|   1 |  TABLE ACCESS STORAGE FULL            | DUAL |     1 |      2 |    2   (0)| 00:00:01 |
-----+-----+-----+-----+-----+-----+-----+
```

```
Query Block Name / Object Alias (identified by operation id):
```

```
-----
1 - SEL$1 / DUAL@SEL$1
```

```
Column Projection Information (identified by operation id):
```

```
-----
1 - "DUAL"."DUMMY" [VARCHAR2,1]
```

## Execution Plan Reading

# DBMS\_XPLAN

## DISPLAY\_AWR Function Example

```
SQL> SELECT plan_table_output
        FROM TABLE(dbms_xplan.display_awr(sql_id => ' 78zm0s46bbm5h',
                                           plan_hash_value => NULL
                                           db_id => NULL
                                           format => 'TYPICAL +Note'));
```

PLAN\_TABLE\_OUTPUT

-----  
SQL\_ID 78zm0s46bbm5h

...

Plan hash value: 376465169

-----  
| Id | Operation | Name | Rows | Bytes | Cost (%CPU) | Time |  
-----  
...  
-----

Note

-----

- dynamic sampling used for this statement (level=2)
- cardinality feedback used for this statement

# DBMS\_XPLAN

## Output Format

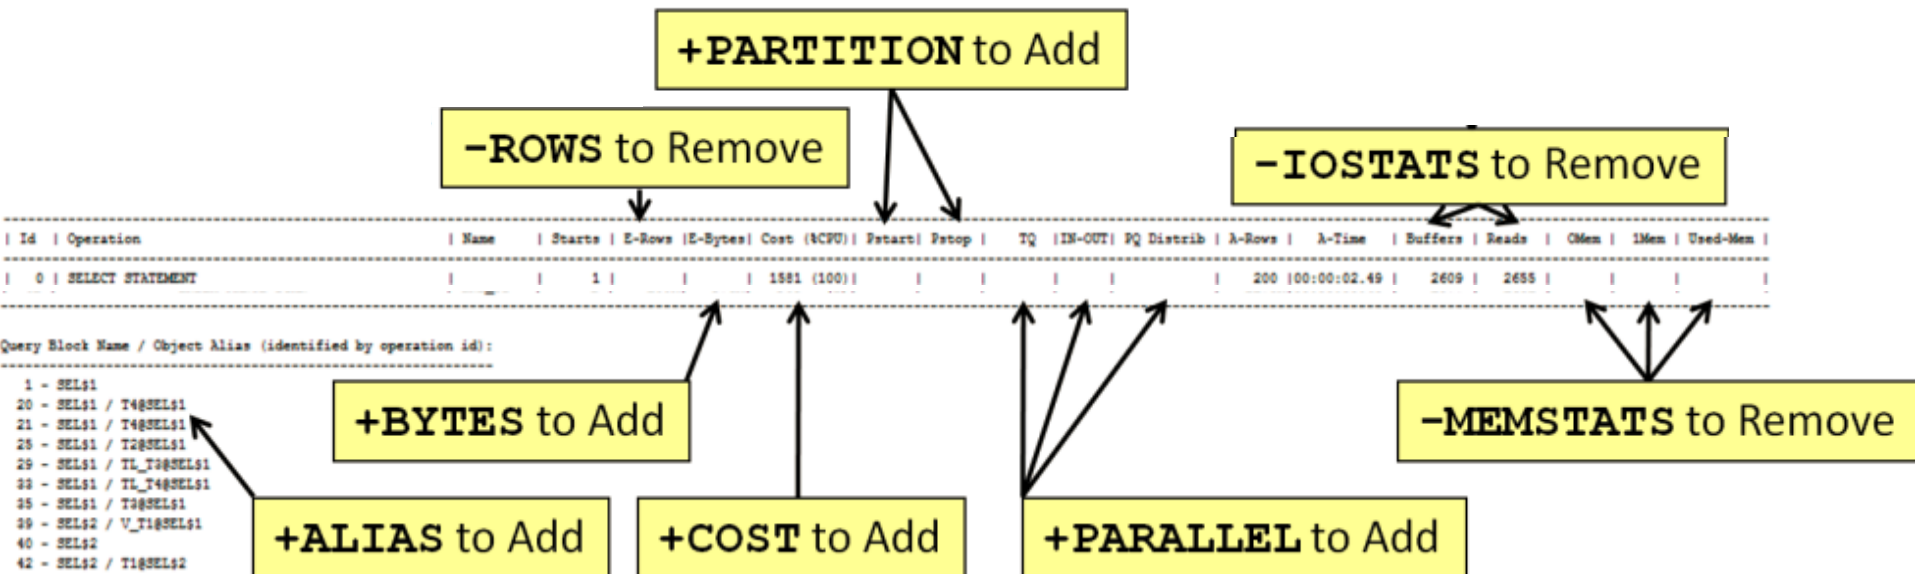
**dbms\_xplan.display(NULL, NULL, format => '<FORMAT +-flag +-flag ...>')**

FORMAT: BASIC, TYPICAL, ALL, ADVANCED = ALL+Outline

flags: columns: iostat, memstat, allstats=iostat & memstat, bytes, parallel, cost, rows, partition

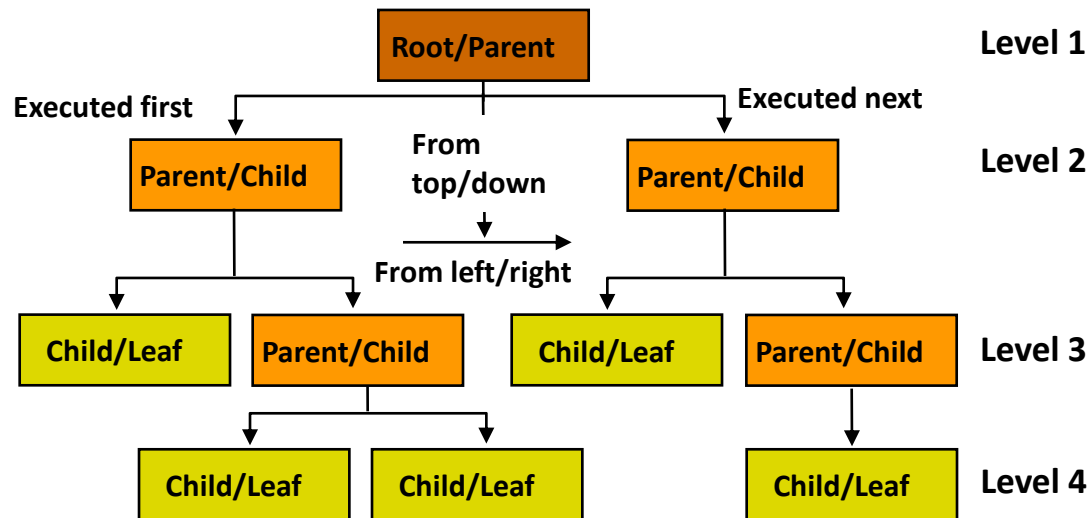
sections: Alias, Outline, Predicate, Projection, Note

executions: last, all (default)



# Interpretation

id= 1	(pid= )		root/parent
id= 2	(pid=1)	(pos=1)	parent/child
id= 3	(pid=2)	(pos=1)	child/leaf
id= 4	(pid=2)	(pos=2)	parent/child
id= 5	(pid=4)	(pos=1)	child/leaf
id= 6	(pid=4)	(pos=2)	child/leaf
id= 7	(pid=1)	(pos=2)	parent/child
id= 8	(pid=7)	(pos=1)	child/leaf
id= 9	(pid=7)	(pos=2)	parent/child
id=10	(pid=9)	(pos=1)	child/leaf



Execution Plan Reading

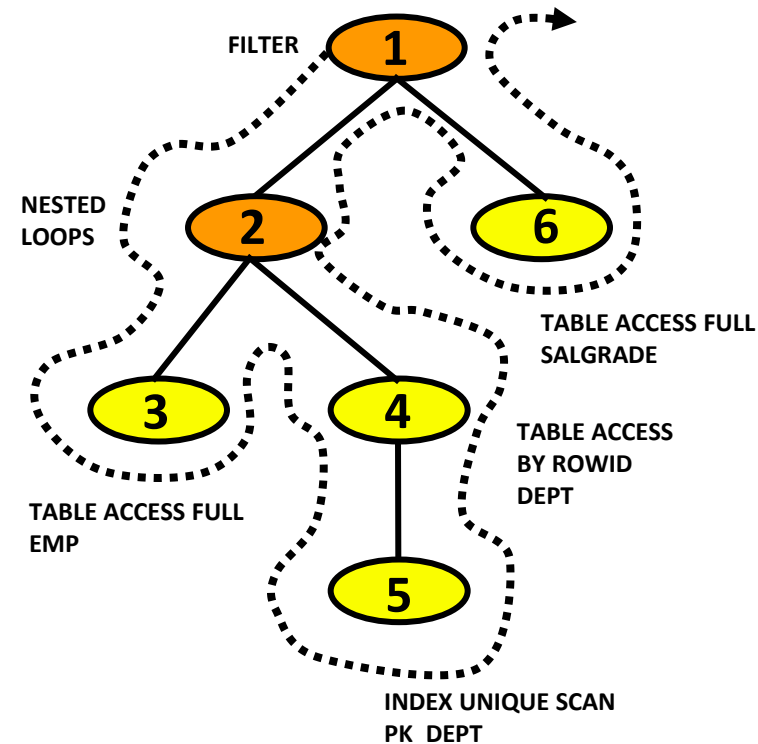
# Interpretation: Example 1

```
SELECT ename, job, sal, dname
FROM emp, dept
WHERE dept.deptno = emp.deptno AND NOT EXISTS (SELECT *
                                                FROM salgrade
                                                WHERE emp.sal BETWEEN losal AND hisal);
```

Id	Operation	Name
0	SELECT STATEMENT	
* 1	FILTER	
2	NESTED LOOPS	
3	TABLE ACCESS FULL	EMP
4	TABLE ACCESS BY INDEX ROWID	DEPT
* 5	INDEX UNIQUE SCAN	PK_DEPT
* 6	TABLE ACCESS FULL	SALGRADE

Predicate Information (identified by operation id):

- 1 - filter( NOT EXISTS  
(SELECT 0 FROM "SALGRADE" "SALGRADE" WHERE  
"HISAL">=:B1 AND "LOSAL"<=:B2))
- 5 - access("DEPT"."DEPTNO"="EMP"."DEPTNO")
- 6 - filter("HISAL">=:B1 AND "LOSAL"<=:B2)



# Interpretation: Example 1

```
SQL> ALTER SESSION SET statistics_level=ALL;
```

```
Session altered.
```

```
<execute query>
```

```
no rows selected
```

```
SQL> SELECT *
```

```
FROM TABLE(dbms_xplan.display_cursor(null,null,'TYPICAL IOSTATS LAST'));
```

```
SQL_ID 274019myw3vuf, child number 0
```

```
-----  
...  
Plan hash value: 1175760222
```

SQL Monitoring	
Executions	Actual Rows

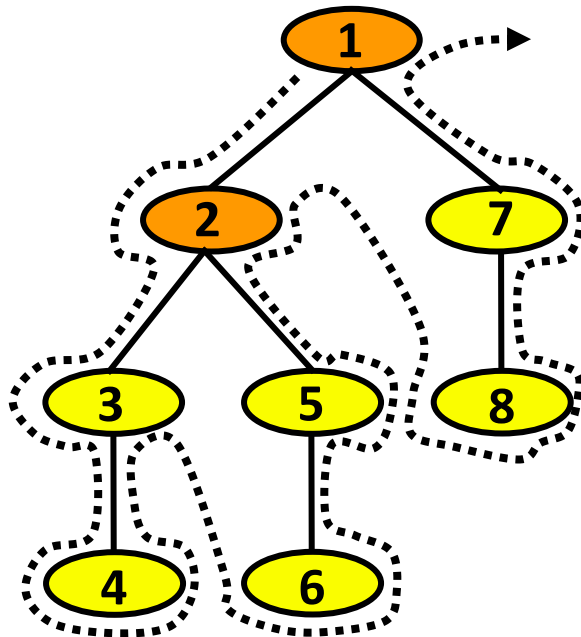
-----		-----		-----		-----		-----	
Id		Operation		Name		Starts		A-Rows	Buffers
-----									
*	1	FILTER				1		0	61
	2	NESTED LOOPS				1		14	25
	3	TABLE ACCESS FULL		EMP		1		14	7
	4	TABLE ACCESS BY INDEX ROWID		DEPT		14		14	18
*	5	INDEX UNIQUE SCAN		PK_DEPT		14		14	4
*	6	TABLE ACCESS FULL		SALGRADE		12		12	36
-----									

```
...
```



# Interpretation: Example 2

```
SQL> SELECT /*+ USE_NL(d) USE_NL(m) */  
        m.last_name AS dept_manager,  
2       d.department_name,  
3       l.street_address  
4 FROM hr.employees m  
5 JOIN hr.departments d ON (d.manager_id = m.employee_id)  
6 NATURAL JOIN hr.locations l  
7 WHERE l.city = 'Seattle';
```



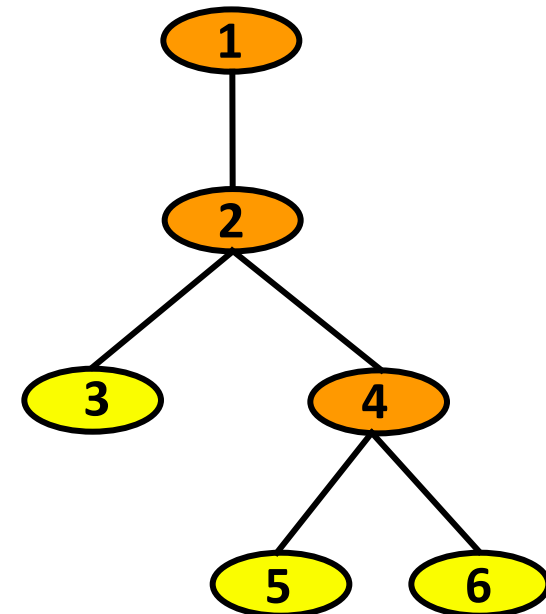
0	SELECT STATEMENT
1 0	NESTED LOOPS
2 1	NESTED LOOPS
3 2	TABLE ACCESS BY INDEX ROWID LOCATIONS
4 3	INDEX RANGE SCAN LOC_CITY_IX
5 2	TABLE ACCESS BY INDEX ROWID DEPARTMENTS
6 5	INDEX RANGE SCAN DEPT_LOCATION_IX
7 1	TABLE ACCESS BY INDEX ROWID EMPLOYEES
8 7	INDEX UNIQUE SCAN EMP_EMP_ID_PK

# Interpretation: Example 3

```
SELECT /*+ ORDERED USE_HASH(b) SWAP_JOIN_INPUTS(c) */  
      MAX(a.i)  
FROM t1 a, t2 b, t3 c  
WHERE a.i = b.i AND a.i = c.i;
```

```
0  SELECT STATEMENT  
1  SORT AGGREGATE  
2 1  HASH JOIN  
3 2  TABLE ACCESS FULL T3  
4 2  HASH JOIN  
5 4  TABLE ACCESS FULL T1  
6 4  TABLE ACCESS FULL T2
```

<a href="#">Expand All</a>	<a href="#">Collapse All</a>		
Operation	Object	Order	
▼ SELECT STATEMENT		7	
▼ SORT AGGREGATE		6	
▼ HASH JOIN		5	
TABLE ACCESS FULL	<a href="#">T3</a>	1	
▼ HASH JOIN		4	
TABLE ACCESS FULL	<a href="#">T1</a>	2	
TABLE ACCESS FULL	<a href="#">T2</a>	3	



Join order is: T1 - T2 - T3

# Q & A

## + Quiz

# Plan Reading Workshop



- Interpret plan for top dbtime SQLs
  - using EM SQL Monitoring
  - during worklod jobs executed
- Display top CPU SQL plan using dbms\_xplan
  - Minimize displayed information and later display all information
  - Analyze content of plan full report
- Use Autotrace to see statistics only (plan only later)
  - by execution labs/autotrace.sh script in virtual linux desktop

```
SELECT t.chanl_desc, p.prod_name, f.sold_amt
FROM sh_sales_fct f
     INNER JOIN sh_chanl_dim t ON t.chanl_id = f.chanl_id
     INNER JOIN sh_prod_dim p ON p.prod_id = f.prod_id
WHERE p.prod_id > 10;
```

- Change tables order in SQL text - is plan changed?
- Test following hints: ORDERED, FULL(p), LEADING(p)

# Q & A

## + Workshop Comp

# SQL Tuning Resources

- Oracle Database Documentation Library

[http://docs.oracle.com/cd/E11882\\_01/index.htm](http://docs.oracle.com/cd/E11882_01/index.htm)

- Database Performance Tuning Guide

[http://docs.oracle.com/cd/E11882\\_01/server.112/e41573/toc.htm](http://docs.oracle.com/cd/E11882_01/server.112/e41573/toc.htm)

- Oracle SQL Tuning

[http://docs.oracle.com/cd/E28271\\_01/server.1111/e16638/sql\\_overview.htm](http://docs.oracle.com/cd/E28271_01/server.1111/e16638/sql_overview.htm)