

Oracle SQL Tuning

3 parts (BASIC, INTERMEDIATE, ADVANCED)

www.lingaro.com



Training Agenda

- **BASIC**
 - SQL Tuning Introduction
 - Instrumentation & Rewrite SQL Text
 - SQL Plan Reading
- **INTERMEDIATE**
 - Optimizer & Statistics
 - SQL Plan Operations
 - SQL Plan Tuning
- **ADVANCED**
 - SQL Processing & Cursor Sharing
 - Advanced Plan Operations
 - SQL Plan Transformations
 - SQL Plan Management

Training Agenda

- For All Topics
 - Theory
 - Quiz Comp
 - Workshop
 - Workshop Comp
- Winner Awards

Training Agenda

- Tuning Techniques (not covered yet)
 - Using Automatic SQL Tuning
 - Using SQL Tuning & SQL Access Advisors (`dbms_sqltune` or OEM)
 - Using Oracle Enterprise Manager Console OEM
 - Using `DBA_HIST` and `V$` performance views (in BASIC and here)
 - Using Test Executions
 - Using SQL Decomposition
 - Test Execution on SQL parts

Oracle SQL Tuning

INTERMEDIATE

www.lingaro.com

Training Agenda

- Optimizer & Statistics
- SQL Plan Operations
- SQL Execution Plan Tuning

Topic Agenda

Oracle Optimizer

- **Optimizer Introduction**
- **Optimizer Components**
 - Calculator, Transformer, Estimator, Plan Generator
- **Optimizer Statistics**
 - Tables, Columns, Histograms and Indexes
- **Execution Cost Estimation**
 - Selectivity, Cardinality, Cost
 - Optimizer Estimation Examples
- **Dynamic Sampling**

Optimizer

Introduction

- SQL statements - description **what to do** with data
- Instruction **how to do** it - **not needed** inside SQL
- **Optimizer** - used to **chooses** SQL execution **plan** - from many possible
- Optimizer - components used **during hard parse** step

- Calculator - to simplify expressions

e.g. `SELECT sales_amt * 100 / (47 + 3)` -> `SELECT sales_amt * 2`

- Transformer - to rewrite statement on the fly

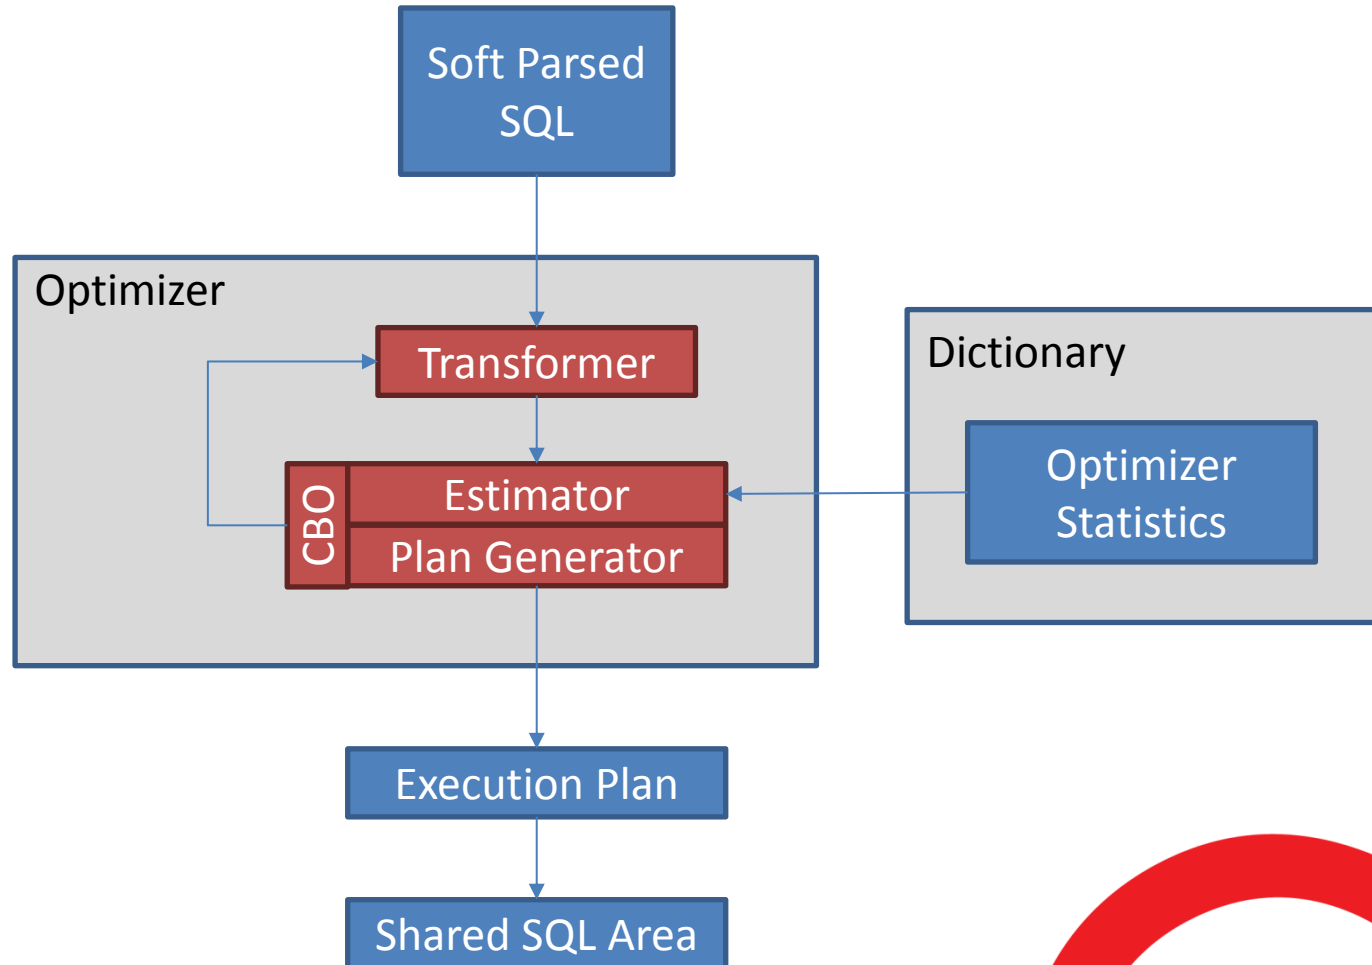
e.g. changes subquery into join

- **Estimator** - to estimate operations costs - visible in execution plan

e.g. number of rows, memory, time estimates for particular statement operations

- Plan Generator - to build SQL execution plan

Optimizer Components



Optimizer Statistics

Introduction

- Metadata - **describes data** and its storage (in tables, indexes)
- Used by Optimizer Estimator to **estimate** SQL Plan **Operations costs**.
- **Are static - not changed when data are changed - refresh needed.**
 - Gathered using **samples** - random part of data - to reduce refresh cost.
 - Database has (default enabled) **refresh job** - disabled in most DW.
 - If defined percent of data is modified then statistic is marked as **STALE**.
- ETL processes **should refresh** it just after data load or modification.
- **Not up to date statistics can lead to suboptimal SQL execution plans.**
 - If statistic is high **deviated** then **wrong plan** can make execution costs.
 - **Dynamic sampling** can be used **when** optimizer statistics are **not available** in data dictionary during SQL parse phase.
- Object, **partitions** and sub partitions **have individual** statistics.
- Optionally **column statistics and** column **histograms** can be added.

Optimizer Statistics

Statistics Maintenance

- Use **dbms_stats** package procedures

- Calculate/refresh

- `gather_table_stats` - for table & columns (optionally with indexes)
 - `gather_index_stats` - for index
 - `gather_schema_stats` - for all objects in schema

- Delete and lock

- `delete_table_stats`, `delete_column_stats`, `delete_index_stats`, `delete_schema_stats`
 - `lock_table_stats`, `lock_schema_stats`, `unlock_table_stats`, `unlock_schema_stats`

- Manual entering and reading

- `set_table_stats`, `set_col_usag`, `set_index_stats`
 - `get_table_stats`, `get_column_stats`, `get_index_stats`

- Restoring from history. - Get historical timestamps from `USER_TAB_STATS_HISTORY` view

- `restore_table_stats`, `restore_schema_stats`

- Export & import

- `create_stat_table`, `export_table_stats`, `import_table_stats`

Example:

```
dbms_stats.gather_table_stats(USER, 'SALES_FCT', estimate_percent => 5)
```

- Default values for few parameters are configured as preferences

Optimizer Statistics

Gather Table Statistics Procedure Parameters

- **partname** - Name of partition or subpartition
- **estimate_percent** - Percentage of rows to estimate. (100% if over 20%)
The default is DBMS_STATS.AUTO_SAMPLE_SIZE
- **block_sample**: TRUE - blocks sample, FALSE - rows sample
- **method_opt** - **columns statistics** and **histogram buckets** specification
 - 'FOR ALL [INDEXED | HIDDEN] COLUMNS [SIZE <buckets>] '
 - 'FOR COLUMNS [SIZE <buckets>] <column> [SIZE <buckets>] [, <column> [SIZE <buckets>] ...] '

<buckets>:

- **integer** - Number of histogram buckets. Must be in the range [1 - 254].
- **REPEAT** - Collects histograms only on the columns that already have histograms.
- **AUTO** - Oracle determines the columns to collect histograms based on data distribution and the workload of the columns.
- **SKEWONLY** - Oracle determines the columns to collect histograms based on the data distribution.

The default is 'FOR ALL COLUMNS SIZE AUTO'.

Optimizer Statistics

Gather Table Statistics Procedure Parameters (continued)

- **degree** - Degree of parallelism.
- **granularity**:
 - 'ALL' - Gathers all (sub partition, partition, and global) statistics
 - 'AUTO' - Determines the granularity based on the partitioning. This is the default.
 - 'GLOBAL' - Gathers global statistics
 - 'GLOBAL AND PARTITION' - Gathers the global and partition level statistics.
No sub partition level statistics are gathered.
 - 'PARTITION' - Gathers partition-level statistics
 - 'SUBPARTITION' - Gathers sub partition-level statistics.
- **cascade** - Gather statistics on the indexes for this table.
- **stattab** - User statistics table name
- **statid** - Identifier associate with these statistics within stattab
- **statown** - Schema containing stattab (if different than ownname)
- **no_invalidate** - TRUE - Does not invalidate the dependent cursors.
- **force** - Gather statistics of table even if it is locked

Optimizer Statistics

Statistics Maintenance Examples

- **Delete** table level or partition level or column level statistics

```
dbms_stats.delete_table_stats(USER, 'SALE_FCT', cascade_parts => FALSE, cascade_columns => TRUE)
dbms_stats.delete_table_stats(USER, 'SALE_FCT', partname => 'P201503', cascade_parts => FALSE)
dbms_stats.delete_column_stats(USER, 'SALE_FCT', 'PROD_ID', cascade_parts => FALSE,
                                col_stat_type => 'HISTOGRAM')
```

- **Compute** on sub partition **without columns** or **without histogram**

```
dbms_stats.gather_table_stats(USER, 'SALE_FCT', partname => 'P201503-WM',
                                granularity => 'SUBPARTITION', method_op => NULL)
dbms_stats.gather_table_stats(USER, 'SALE_FCT', method_op => method_op => 'FOR ALL COLUMNS SIZE 1')
```

- **Lock** statistic - `dbms_stats.lock_table_stats(USER, 'SALE_FCT')`

- **Restore** table statistics

```
SELECT table_name, partition_name, subpartition_name, stats_update_time FROM user_tab_stats_history;
dbms_stats.restore_table_stats(USER, 'SALE_FCT', as_of_timestamp => SYSDATE - 1)
```

- **Create stat table** - `dbms_stats.create_stat_table(USER, 'MY_STAT_TABLE')`

- **Export/Import** table statistics

```
dbms_stats.export_table_stats(USER, 'SALE_FCT', stattable => 'MY_STAT_TABLE')
```

- **Set/Get** table stats

```
dbms_stats.set_table_stats(USER, 'SALE_FCT', numrows => 1222000, numblks => 312000, avgrlen => 122300)
```

Optimizer Statistics

Statistics Preferences

- Used to balance accuracy and costs and define default gathering parameters
- Can be set on schema & table levels

```
dbms_stats.set_schema_prefs, dbms_stats.set_table_prefs
```

- Default values are set on database level by DBA

```
dbms_stats.get_database_prefs
```

- Preferences

- CASCADE - Collect index statistics with table statistics.
- DEGREE - Determines parallelism for gathering statistics.
- ESTIMATE_PERCENT - Percentage of rows to estimate.
- METHOD_OPT - Controls column statistics and histogram
- NO_INVALIDATE - If FALSE - invalidation of dependent cursors
- GRANULARITY - Gathering level for partitioned tables
- PUBLISH - If FALSE then statistic is pending
- INCREMENTAL - If TRUE Oracle update global statistics by scanning only changed partitions.
- STALE_PERCENT - Percentage of rows have to change before the statistics are deemed stale.
The default value is 10%. Blue - available only in preferences

- Example

```
dbms_stats.set_table_prefs(USER, 'SALES_FCT', 'STALE_PERCENT', '15')
```

Optimizer Statistics

Table Statistics

- Dictionary Views
 - USER_TABLES Only tables statistics
 - USER_TAB_STATISTICS Tables, partitions, and subpartitions statistics in separate rows
- Dictionary View Columns
 - NUM_ROWS Number of rows in the object
 - BLOCKS Number of used blocks in the object – **EXACT VALUE**
 - EMPTY_BLOCKS Number of empty blocks in the object – **EXACT VALUE**
 - AVG_SPACE Average available free space in the object
 - CHAIN_CNT Number of chained rows in the object
 - AVG_ROW_LEN Average row length, including row overhead
 - AVG_SPACE_FREELIST_BLOCKS Average freespace of all blocks on a freelist
 - NUM_FREELIST_BLOCKS Number of blocks on the freelist
 - AVG_CACHED_BLOCKS Average number of blocks in the buffer cache
 - AVG_CACHE_HIT_RATIO Average cache hit ratio for the object
 - SAMPLE_SIZE Sample size used in analyzing the table
 - LAST_ANALYZED Date of the most recent time the table was analyzed
 - GLOBAL_STATS Indicates calculated without merging underlying partitions YES/NO
 - USER_STATS Indicates whether statistics were entered directly by the user YES/NO
 - STATTYPE_LOCKED Type of lock: <null>, (DATA, CACHE) - internal lock, **ALL** - locked
 - STALE_STATS Indicates whether statistics for the object are stale YES/NO

Optimizer Statistics

Table Columns Statistics

- Dictionary Views
 - USER_TAB_COL_STATISTICS or USER_COLUMNS Column statistics for tables
 - USER_PART_COL_STATISTICS Column statistics for table partitions
 - USER_SUBPART_COL_STATISTICS Column statistics for table subpartitions
- Dictionary View Columns
 - NUM_DISTINCT **NDV** - Number of distinct values in the column
 - LOW_VALUE Low value in the column - **EXACT VALUE**
 - HIGH_VALUE High value in the column - **EXACT VALUE**
 - DENSITY Density of the column (similar to selectivity but not used from 10g)
 - NUM_NULLS Number of nulls in the column
 - NUM_BUCKETS Number of buckets in histogram for the column
 - LAST_ANALYZED Date on which this column was most recently analyzed
 - SAMPLE_SIZE Sample size used in analyzing this column
 - GLOBAL_STATS Indicates whether collected as a whole (YES) or partitions (NO)
 - USER_STATS Indicates whether entered directly by the user (YES) or not (NO)
 - AVG_COL_LEN Average length of the column (in bytes)
 - HISTOGRAM Indicates existence/type of histogram:
NONE, FREQUENCY, HEIGHT BALANCED

Optimizer Statistics

Histograms

- Are optional column statistics used to **describe skew distributed data**
- Dictionary Views
 - USER_TAB_HISTOGRAMS Column histograms for tables
 - USER_PART_HISTOGRAMS Column histograms for table partitions
 - USER_SUBPART_HISTOGRAMS Column histograms for table sub partitions
- Dictionary View Columns
 - ENDPOINT_NUMBER Histogram bucket number
 - ENDPOINT_VALUE Normalized endpoint value for this bucket
 - ENDPOINT_ACTUAL_VALUE Actual string value of the endpoint for this bucket





Example Table
Histogram.sql

```
dbms_stats.gather_table_stats(USER, '<table name>',  
    method_opt => 'for columns <column> size <buckets>');
```

- Types - Oracle decision
 - FREQUENCY
Typically used when NDV = buckets count (small NDV)
 - HEIGHT BALANCED
Typically used if NDV is higher then buckers count

Examples: 16 rows 8 values 1 - 8



Optimizer Statistics

Reading Histograms

- Checking histogram for columns

```
SELECT table_name, column_name, num_distinct, num_buckets, histogram
FROM user_tab_col_statistics
WHERE table_name = 'SEA_VALDN_FUNC_LKP' AND column_name = 'VALDN_ID' OR
      table_name = 'SEA_INTVN_FCT' AND column_name = 'INTVN_DESC' OR
      table_name = 'SEA_EXPSN_PI_AND_DV_DRIDN_IFCT' AND column_name = 'LAST_ONHND_QTY'
```

TABLE_NAME	COLUMN_NAME	NUM_DISTINCT	NUM_BUCKETS	HISTOGRAM
SEA_EXPSN_PI_AND_DV_DRIDN_IFCT	LAST_ONHND_QTY	214	71	FREQUENCY
SEA_INTVN_FCT	INTVN_DESC	13	147	HEIGHT BALANCED
SEA_VALDN_FUNC_LKP	VALDN_ID	14	14	FREQUENCY

- Display buckets values

FREQUENCY

```
SELECT endpoint_value          AS bckt_val,
       endpoint_number - lag(endpoint_number, 1, 0)
       OVER (ORDER BY endpoint_number) AS rows_cnt,
       endpoint_number          AS rows_cnt_cum
FROM user_histograms
WHERE table_name='<table name>'
AND column_name='<column name>'
ORDER BY bckt_val;
```

HEIGHT BALANCED

```
SELECT endpoint_number          AS bckt_num,
       <NDV>/<buckets> * endpoint_number AS bckt_last_row_num,
       endpoint_value            AS last_row_val
FROM user_histograms
WHERE table_name='<table name>'
AND column_name='<column name>'
ORDER BY bckt_num;
```

- Popular value - bucked ending value for more then one bucket

Optimizer Statistics

Index Statistics

- Dictionary Views
 - USER_INDEXES Only indexes statistics
 - USER_IND_STATISTICS Indexes, index partitions and index subpartitions statistics in separate rows
- Dictionary View Columns
 - BLEVEL B-Tree level – **EXACT VALUE**
 - LEAF_BLOCKS Number of leaf blocks in the index
 - DISTINCT_KEYS Number of distinct keys in the index
 - AVG_LEAF_BLOCKS_PER_KEY Average number of leaf blocks per key
 - AVG_DATA_BLOCKS_PER_KEY Average number of data blocks per key
 - CLUSTERING_FACTOR Indicates order of the rows in the table related to index order.
If the value is near the number of blocks, then the table is well ordered.
 - NUM_ROWS Number of rows in the index
 - AVG_CACHED_BLOCKS Average number of blocks in the buffer cache
 - AVG_CACHE_HIT_RATIO Average cache hit ratio for the object
 - SAMPLE_SIZE Sample size used in analyzing the index
 - LAST_ANALYZED Date of the most recent time the index was analyzed
 - GLOBAL_STATS Indicates calculated without merging underlying partitions YES/NO
 - USER_STATS Indicates whether statistics were entered directly by the user YES/NO
 - STATTYPE_LOCKED Type of statistics lock
 - STALE_STATS Whether statistics for the object are stale or not

Q & A

+ Quiz

Two large, thick white circles are positioned at the bottom right of the slide, partially overlapping each other.

www.lingaro.com

Optimizer Statistics Workshop



- Fact table statistics creation and reading
 - Use SQL_Tuning_INTERM_workshop.sql script in SQL-Developer
 - Connect to database on VBOX using train/oracle@<vbox IP>/orcl
 - Create table SALES_FCT from SH_SALES_FCT (only chanl_id=4)
and bitmap index on prod_id
and create product_dim with PK from sh_product_dim
 - Compute statistics on tables using default preferences
 - Display table and columns statistics, histograms and index statistics
 - Display histogram buckets for selected column

Q & A

+ Workshop Comp

www.lingaro.com

Optimizer Statistics

Refresh Strategy for DW

- Assumptions - typical DW workload
 - No data modifications, only load in ETL processes
 - Statistics calculated immediately after data load - in ETL process
- For **small** unpartitioned metadata, lookup and dimension **tables**
 - Recalculate all statistics on whole table if statistic not exists or STALE_STATS = 'YES'
 - Use columns statistics on all columns
- For **large partitioned** fact and dimension tables
 - Calculate only lowest level (e.g. sub partition) statistic after load and before exchange
 - Use one from **two available strategy** to have fresh statistics on higher levels
 - 1 - **Cheap** but problems with NDV quality - steps below for sub partitioned tables
 - Delete statistics on table and partition level
 - Keep all sub partitions statistics even are empty or not used
 - 2 - **Accurate NDV** but more expensive so use only for few tables with NDV problems
 - Use INCREMENTAL = TRUE in table statistics preferences
 - Calculate statistics using SKEWONLY in method_opt buckets count
- **BEWARE maximum, minimum values in column statistics & buckets values in histograms**
 - For dates and sequence numbers in columns stale statistics can make problems with estimates
 - Possible very low wrong cardinality estimation and very bad plans if statistics are not fresh
- Use SIZE = SKEWONLY (not AUTO) if no nominal workload was executed yet on tables

Optimizer Estimations

Terminology

- Estimated **Selectivity**

- Fraction of rows to be processed by 1 operation (from many statement operations) estimation
- Values between 0 and 1 where 1 means all rows
- For non-unique column equality filtering: $\text{Selectivity} = 1 / \text{NDV column stat}$
- For range filtering: $\text{Selectivity} = (\text{high filter} - \text{low filter}) / (\text{highest val stat} - \text{lowest val stat})$
- If histogram exists selectivity is calculated from histogram
- For many column filters (AND) selectivity = multiply of columns selectivity
 - It is not correct if columns are related like product category and sub category
 - if related use Extended (multicolumn-virtual column) Column Statistics

- Estimated **Cardinality**

- Number of rows processed by operation estimation
- $\text{Cardinality} = \text{Selectivity} * \text{Number of rows table stat}$
- $\text{Join Cardinality} = T1 \text{ Cardinality} * T2 \text{ Cardinality} * \text{Join Selectivity}$

- Estimated **Cost**

$$\frac{\text{single block reads cost} + \text{multi block reads cost} + \text{CPU cost}}{\text{1blk read cnt} * \text{1blk read time} + \text{mblk read cnt} * \text{mblk read time} + \text{CPU cycles} / \text{cycles per second}}$$

CPU speed
↓

- Is used in plan generation too choose plan with lowest cost

Optimizer Statistics

Cardinality Validation

- You can validate cardinality by **comparing estimated and actual rows count**

1)

- Query with count(*) instead of resulting column list and
- CARDINALITY column of plan operation visible after F10 pressed in SQL*Developer

2)

- E-Rows and A-Rows columns in plan from DBMS_XPLAN with 'ALLSTATS LAST' option
- Use STATISTIC_LEVEL = ALL on session or GATHER_PLAN_STATISTICS hint in SQL

3)

- Columns

Estimated Rows	Actual Rows
----------------	-------------

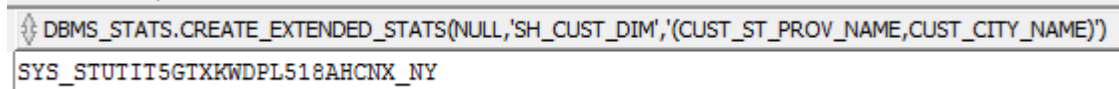
 from SQL Monitoring in:
- v\$sql_monitor and v\$sql_monitor_plan
- SQL Monitor Lingaro tool
- Enterprise Manager Console

Optimizer Statistics

Columns Extended Statistics

- If many **related columns** in filter
- Create implicit virtual column

```
SELECT dbms_stats.create_extended_stats(NULL,  
    '<table name>', '(<column1>, <column2>') FROM dual;
```



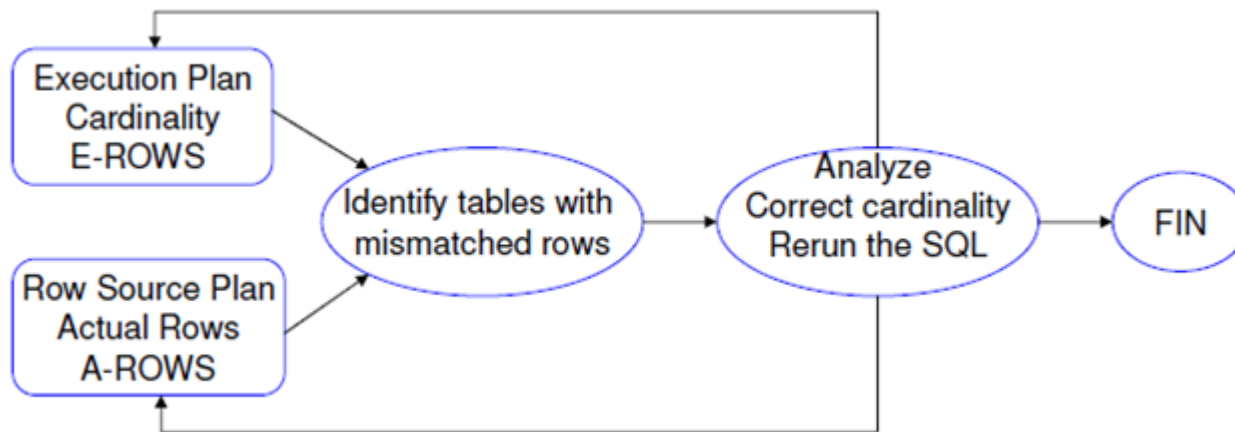
SQL*Plus output showing the execution of the DBMS_STATS.CREATE_EXTENDED_STATS procedure. The command is: DBMS_STATS.CREATE_EXTENDED_STATS(NULL,'SH_CUST_DIM','(CUST_ST_PROV_NAME,CUST_CITY_NAME)'). The output is: SYS_STUTIT5GTXXWDPL518AHCNX_NY.

- Compute all column statistics with SKEWONLY option
- Can create explicit virtual column with any expression
- Can calculate column **statistics on any virtual columns**
 - No expression based index needed

Optimizer Statistics

Cardinality Feedback

- Use **Statistic Level ALL** or GATHER_PLAN_STATISTICS hint
 - If you have problem with cardinality estimation
- **Cardinality** estimation will be **automatically corrected**
 - Correct value will take into account actual row number from first execution
- Plan will change on second execution



Dynamic Sampling

Overview

- When dictionary stored optimizer statistics not available (or insufficient)
- Calculate it during parse phase after recursive sample e.g. 10% table scans

```
SELECT /* OPT_DYN_SAMP */ ... FROM table SAMPLE(10) ...
```

- Criteria to use it and number of block sampled depends on sampling level
- Dynamic Sampling level can be modified:
 - On session - `OPTIMIZER_DYNAMIC_SAMPLING` parameter (default 2)
 - In statement - `DYNAMIC_SAMPLING[(level)]` query hint
- If used to build plan then plan is indicated by note

```
Note -----  
- dynamic sampling level 8 used for this statement
```

- Can produce better cardinality estimates
 - Dictionary statistics can be stale
 - But if number of sampled block is smaller then during dictionary statistic gathering then can be worse
- Can lead to performance problems due to
 - Costly sample scan for large number of block
 - CPU cost to process large sampled data
 - Bugs possible for PQ (if default level 2 used) - “library cache lock”

Dynamic Sampling

Levels

- Level 0: Do not use dynamic sampling.
- Level 1: Sample all tables not analyzed if the following criteria are met:
 - there is at least 1 unanalyzed table in the query;
 - this unanalyzed table is joined to another table or appears in a subquery or non-mergeable view;
 - this unanalyzed table has no indexes;
 - this unanalyzed table has more blocks than the number of blocks that would be used for dynamic sampling of this table. The number of blocks sampled is the **default number of dynamic sampling blocks = 32**.
- Level 2 (Default): Apply dynamic sampling to **all unanalyzed tables**.
 - The number of blocks sampled is two times the default number of dynamic sampling blocks.
- Level 3: Sampling to all tables that meet Level 2 criteria, **plus** all tables for which **standard selectivity estimation used a guess** for some predicate that is a potential dynamic sampling predicate.
 - The number of blocks sampled is the default number of dynamic sampling blocks. For unanalyzed tables, the number of blocks sampled is two times the default number of dynamic sampling blocks.
- Level 4: Sampling to all tables that meet Level 3 criteria, **plus** all tables that have single-table **predicates that reference 2 or more columns**.
 - The number of blocks sampled is the default number of dynamic sampling blocks. For unanalyzed tables, the number of blocks sampled is two times the default number of dynamic sampling blocks.
- Levels 5, 6, 7, 8, and 9: Previous level criteria
 - using 2, 4, 8, 32, or 128 times the default number of dynamic sampling blocks respectively.
- Level 10: Previous level criteria
 - using all blocks in the table.

Q & A

+ Quiz

www.lingaro.com

Cardinality Validation

Workshop



- Cardinality Validation

- Create table my_cust_dim as copy of sh_cust_dim
- Compare rows count and cardinality in plan (F10 in SQLdeveloper) for SQL

```
SELECT count(*) FROM my_cust_dim
WHERE cust_st_prov_name = 'CA'           -- filter f1
      AND cust_city_name LIKE 'Los Angeles' -- filter f2
```

- Enter cardinality results in following table

Cardinality	Count(*)	No stats	AUTO	SKEWONLY	Extended	Dynamic Sampling (2)
f1	3341	1213	3175	3455	3354	3064
f2	932	1213	874	874	1093	1108
f1 & f2	932	12	56	54	874	1108

Q & A

+ Workshop Comp

www.lingaro.com

Topic Agenda

SQL Plan Operations


- **Table Scans**
 - FULL, BY ROWID, SAMPLE
- **Index Scans**
 - UNIQUE, RANGE, FULL, FAST FULL, JOIN
- **Table Join Operations**
 - HASH, SORT MERGE, NESTED LOOPS

Table Operations

Full Table Scan

Hint: `FULL(<table alias>)`

- Is faster than index range scans **for large** amount of **data**
- Performs **multiblock** reads (here `DB_FILE_MULTIBLOCK_READ_COUNT = 4`)
- Reads all formatted blocks below the **high-water mark**
- May filter rows
- Can be **parallel**
- Uses private pool



The screenshot shows the SQL Developer interface. The query is `select * from emp where ename='King';`. The execution plan is displayed, showing the following operations:

OPERATION	OBJECT_NAME	OPTIONS
SELECT STATEMENT		
TABLE ACCESS	EMP	FULL
Filter Predicates		ENAME='King'

The 'FULL' option and the 'Filter Predicates' section are highlighted with red boxes.

Table Operations

ROWID Scan

```
SELECT * FROM emp WHERE rowid='AAQ+LAAEAAAAfAAJ';
```

Id	Operation	Name	Rows	Bytes	Cost	

0	SELECT STATEMENT		1	37	1	
1	TABLE ACCESS BY USER ROWID	EMP	1	37	1	

Hint: INDEX(<table alias>, [<index>])

```
SELECT * FROM emp WHERE ename = 'King';
```

Id	Operation	Name	Rows	Bytes	Cost	

0	SELECT STATEMENT		1	37	1	
1	TABLE ACCESS BY INDEX ROWID	EMP	1	37	1	
2	INDEX RANGE SCAN	EMPX	1	17	1	

Table Operations

Sample Table Scan

- Reads rows from **random** selected **blocks sample**
- Only **specified percent** of blocks are chosen
- Oracle uses it **to gather optimizer statistics and for dynamic sampling**

```
SELECT * FROM emp SAMPLE BLOCK (10) SEED (1);
```

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	

0	SELECT STATEMENT		4	99	2 (0)	
1	TABLE ACCESS SAMPLE	EMP	4	99	2 (0)	

Q & A

+ Table Scan Quiz

Index Types

- Storage **techniques**:
 - **B*-tree** indexes: The default and the most common
 - Normal (key on single or multiple columns)
 - Function based (key on expression)
 - Index-organized table (IOT)
 - **Bitmap** indexes
- Index **attributes**:
 - Key compression (**COMPRESS**)
 - Reverse key (**REVERSE**)
 - **Descending** sorted (DESC after key column name)
- **Domain** indexes: Specific to an application or cartridge

Index Types

Normal B-tree Index




Table data retrieved by using ROWID

Plan Operations

Index Operations

B-tree Index

- UNIQUE SCAN
- RANGE SCAN
- SKIP SCAN
- FULL SCAN
- FAST FULL SCAN
- JOIN SCAN
- NULL values in key
- Composite Indexes

Hint: INDEX(tab [indx])


Opp.: NO_INDEX(tab [indx])
FULL(tab)

B-tree Index Operations

UNIQUE SCAN

- Returns **at most single ROWID**
- Statement contains
UNIQUE
PRIMARY KEY
- All columns in unique (B*-tree) index
- Equality condition in filter

index UNIQUE Scan PK_EMP



```
CREATE UNIQUE INDEX pk_emp ON emp(empno)
```

```
SELECT * FROM emp WHERE empno = 7835;
```

Explain Plan x				
0.024 seconds				
OPERATION	OBJECT_NAME	OPTIONS	COST	CARDINALITY
SELECT STATEMENT			1	1
TABLE ACCESS	EMP	BY INDEX ROWID	1	1
INDEX	PK_EMP	UNIQUE SCAN	0	1
Access Predicat				
EMPNO=7835				

Index Operations

RANGE SCAN

Hint: INDEX_ASC

For filters:


col1 = :b1 (only **none unique** index)

col1 < :b1 col1 > :b2

col1 **LIKE** 'ABC%'.

- **Can avoid sorting** operation
ORDER BY/GROUP BY
Only **when** key **columns** are **NOT NULL**

Index Range SCAN I_DEPTNO



```
CREATE INDEX i_deptno ON emp(deptno);  
  
SELECT /*+ INDEX(emp i_deptno) */ * FROM emp  
WHERE deptno = 10 AND sal > 1000;
```

Script Output x Explain Plan x

0.012 seconds

OPERATION	OBJECT_NAME	OPTIONS	COST	CARDINALITY
SELECT STATEMENT			2	4
TABLE ACCESS	EMP	BY INDEX ROWID	2	4
Filter Predicates				
SAL>1000				
INDEX	I_DEPTNO	RANGE SCAN	1	5
Access Predicat				
DEPTNO=10				

Plan Operations

www.lingaro.com

Index Operations

RANGE SCAN DESCENDING

Hint: INDEX_DESC


When scan order is different then index order

```
SELECT * FROM emp WHERE deptno > 20  
ORDER BY deptno DESC;
```

Script Output x Explain Plan x

0.426 seconds

OPERATION	OBJECT_NAME	OPTIONS	COST	CARDINALITY
SELECT STATEMENT			2	7
TABLE ACCESS	EMP	BY INDEX ROWID	2	7
INDEX	I_DEPTNO	RANGE SCAN DESCENDING	1	7
Access Predicate		DEPTNO>20		




```
DROP INDEX i_deptno; CREATE INDEX ix_d ON emp(deptno DESC);  
SELECT * FROM emp WHERE deptno < 30;
```

Script Output x Explain Plan x

0.011 seconds

OPERATION	OBJECT_NAME	OPTIONS	COST
SELECT STATEMENT			2
TABLE ACCESS	EMP	BY INDEX ROWID	2
INDEX	IX_D	RANGE SCAN	1
Access Predicates		SYS_OP_DESCEND(DEPTNO)>HEXTORAW('3EE0FF')	
Filter Predicates		SYS_OP_UNDESCEND(SYS_OP_DESCEND(DEPTNO))<30	




Plan Operations

B-tree Index Operations

FULL SCAN

All leafs read in index order (**eliminate sort**) when:

- predicate references column from index.
- no predicate if both the conditions are met
all referenced columns included in the index
at least one of the index columns is not null



```
CREATE INDEX i_deptno ON emp(deptno);  
  
SELECT * FROM emp  
  WHERE sal > 1000 AND deptno IS NOT NULL  
  ORDER BY deptno;
```

Script Output x Explain Plan x				
0.012 seconds				
OPERATION	OBJECT_NAME	OPTIONS	COST	CARDINALITY
SELECT STATEMENT			2	13
TABLE ACCESS	EMP	BY INDEX RO...	2	13
Filter Predicates SAL>1000				
INDEX	I_DEPTNO	FULL SCAN	1	14
Filter Predicates DEPTNO IS NOT NULL				

Index Operations

FAST FULL SCAN


Hint: INDEX_FFS
Opp.: NO_INDEX_FFS

- All leafs read in storage order.
- Not eliminate sort operation.
- Alternative to full table scans – multi-block and parallel
- When index contains all the columns needed.
 - at least one column in key has a NOT NULL.
- Without accessing the table (index only scan).

db_file_multiblock_read_count = 4

LEGEND:

SH=segment header
R=root block
B=branch block
L=leaf block



```
SELECT /*+ INDEX_FFS(EMP I_DEPTNO) */ deptno
FROM emp WHERE deptno IS NOT NULL;
```

Script Output x Explain Plan x				
0.01 seconds				
OPERATION	OBJECT_NAME	OPTIONS	COST	CARDINALITY
SELECT STATEMENT			2	14
INDEX	I_DEPTNO	FAST FULL SCAN	2	14
Filter Predicates				
DEPTNO IS NOT NULL				


Plan Operations

B-tree Index Operations

SKIP SCAN

Hints: INDEX_SS
INDEX_SS_ASC
INDEX_SS_DESC
Opp.: NO_INDEX_SS

- If **predicate** exists on **not first column** in index **key**.
and **first column** has **small NDV** (number of distinct values)
- Not interesting **leafs** are **skipped**
after checked in **brunch** blocks



```
CREATE INDEX ix_ss ON emp(deptno, sal);
```

```
SELECT /*+ INDEX_SS(emp ix_ss) */ * FROM emp WHERE sal < 1500;
```

Script Output x Explain Plan x

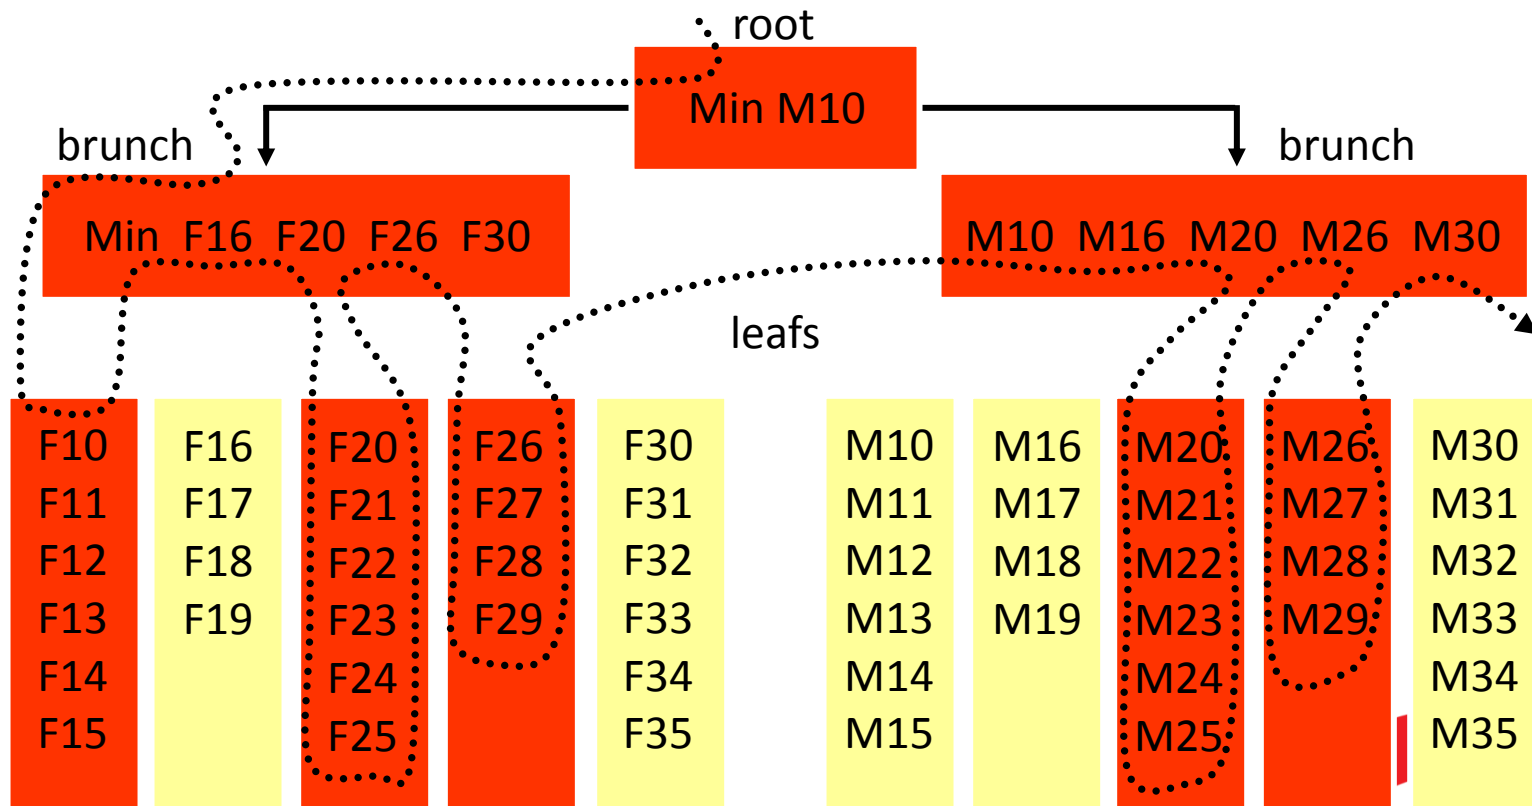
0.01 seconds

OPERATION	OBJECT_NAME	OPTIONS	COST	CARDINALITY
SELECT STATEMENT			4	2
TABLE ACCESS	EMP	BY INDEX ROWID	4	2
INDEX	IX_SS	SKIP SCAN	3	2
Access Predicates				
SAL<1500				
Filter Predicates				
SAL<1500				

B-tree Index Operations

SKIP SCAN - Example

```
SELECT * FROM employees WHERE age BETWEEN 20 AND 29
```



Index Operations

JOIN SCAN

Hint: INDEX_JOIN

- Hash join of several indexes to gather all the table columns referenced
- No table access needed
- Cannot be used to eliminate a sort operation.

```
ALTER TABLE emp MODIFY (sal NOT NULL, ename NOT NULL);  
CREATE INDEX i_ename ON emp(ename);  
CREATE INDEX i_sal ON emp(sal);
```

select /*+ INDEX_JOIN(e) */ ename, sal from emp e;

Statement Output x Autotrace x

1.563 seconds

OPERATION	OBJECT_NAME	COST
SELECT STATEMENT		3
VIEW	index\$_join\$_001	3
type="db_version"		
11.2.0.1		
HASH JOIN		
Access Predicates		
ROWID=ROWID		
INDEX FAST FULL SCAN	I_ENAME	1
INDEX FAST FULL SCAN	I_SAL	1

Plan Operations

B-tree Index Operations

B-tree Do Not Store NULL Values

Null elimination needed
to use index in plan

```
CREATE TABLE nulltest (col1 NUMBER NULL, col2 NUMBER NOT NULL);  
CREATE INDEX nullind1 ON nulltest (col1);  
CREATE INDEX notnullind2 ON nulltest (col2);
```

COL1 is defined NULL
index cannot be used

select /*+ index(t nullind1) */ col1 from nulltest t;

Script Output x Explain Plan x

0.864 seconds

OPERATION	OBJECT_NAME	OPTIONS	COST	CARDINALITY
SELECT STATEMENT			2	1
TABLE ACCESS	NULLTEST	FULL	2	1

Predicate against COL1
this eliminates NULLs
index is used.

select col1 from nulltest t where col1=10;

Script Output x Explain Plan x

1.105 seconds

OPERATION	OBJECT_NAME	OPTIONS	COST	CARDINALITY
SELECT STATEMENT			1	1
INDEX	NULLIND1	RANGE SCAN	1	1

Access Predicates
COL1=10

COL2 is NOT NULL
index is used

select /*+ index(t notnullind2) */ col2 from nulltest t;

Script Output x Explain Plan x


0.034 seconds

OPERATION	OBJECT_NAME	OPTIONS	COST	CARDINALITY
SELECT STATEMENT			1	1
INDEX	NOTNULLIND2	FULL SCAN	1	1

Plan Operations

Index Operations

Composite Indexes



```
CREATE INDEX cars_make_model_idx ON cars(make, model);
```

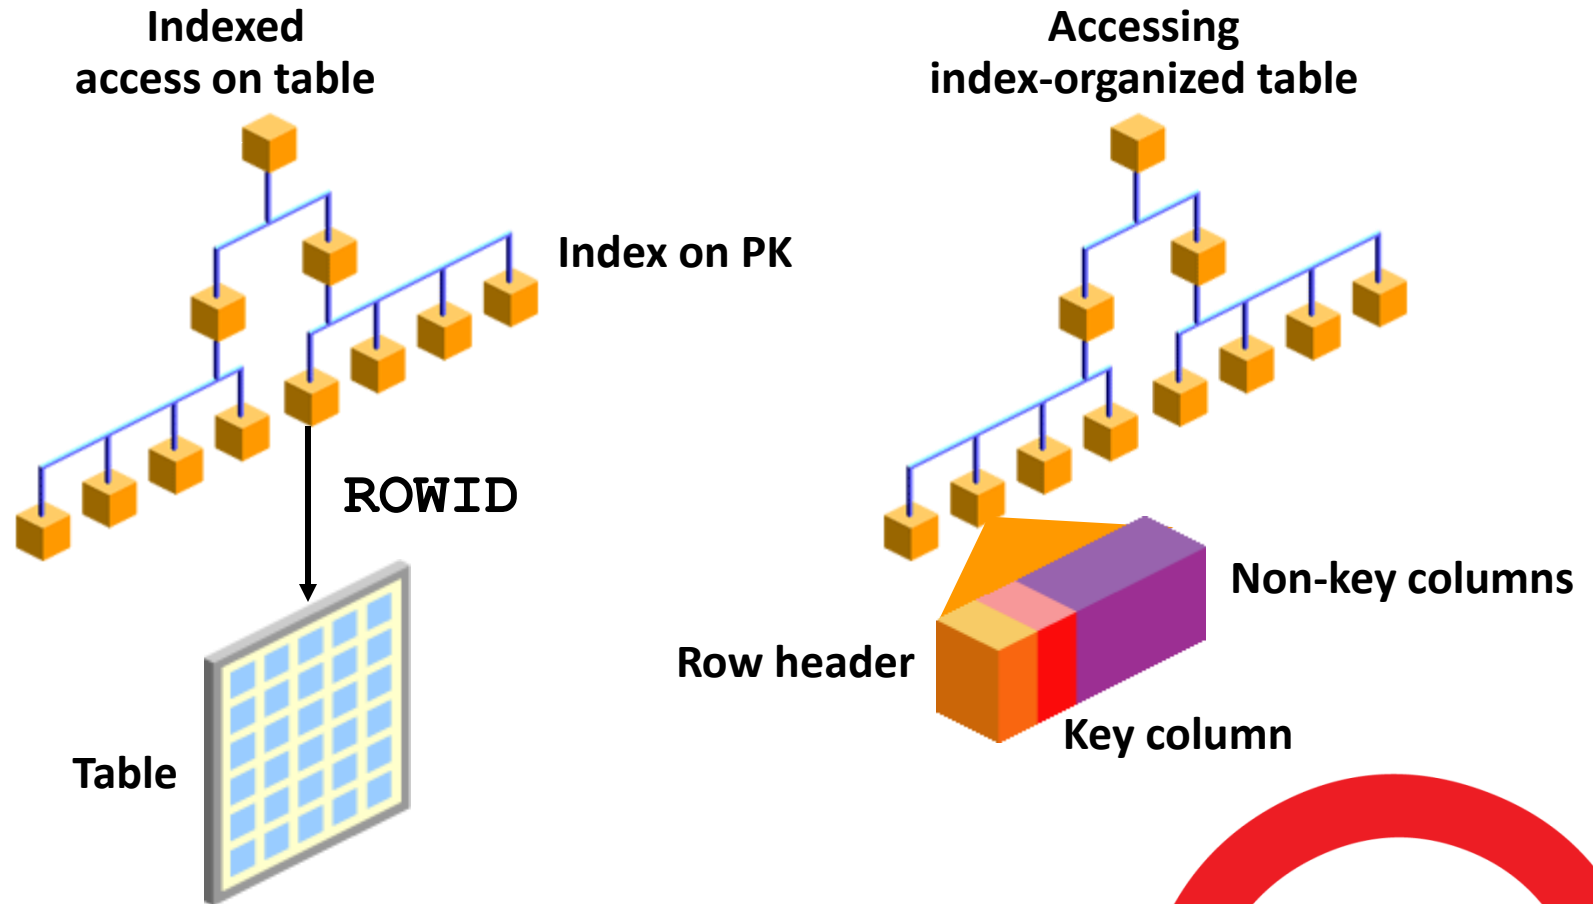
```
SELECT *  
FROM cars  
WHERE make = 'CITROEN' AND model = '2CV';
```

Id	Operation	Name	

0	SELECT STATEMENT		
1	TABLE ACCESS BY INDEX ROWID	CUSTOMERS	
*	INDEX RANGE SCAN	CARS_MAKE_MODEL_IDX	

Index Organized Table

IOT Overview



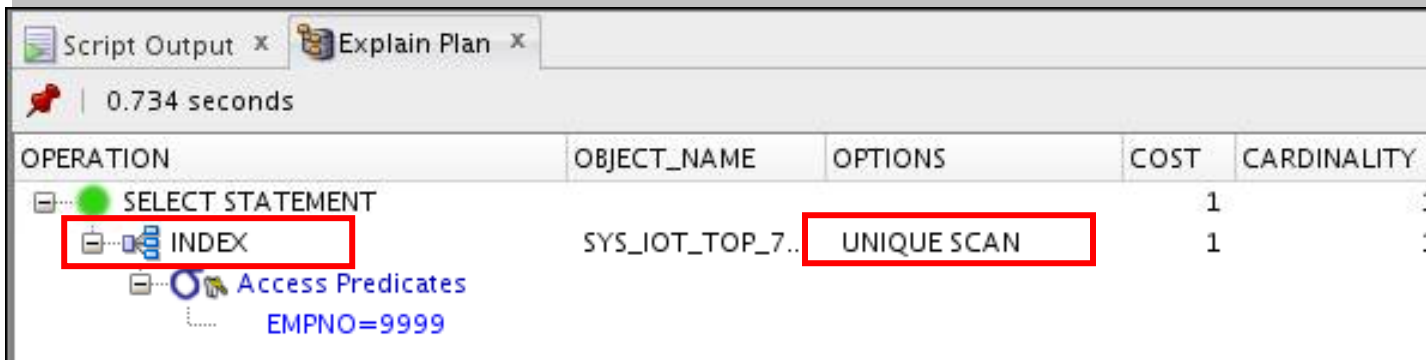
Plan Operations

Index Organized Table

IOT Operations

```
CREATE TABLE iotemp  
( empno NUMBER(4) PRIMARY KEY, ...)  
ORGANIZATION INDEX;
```

```
select * from iotemp where empno = 9999;
```



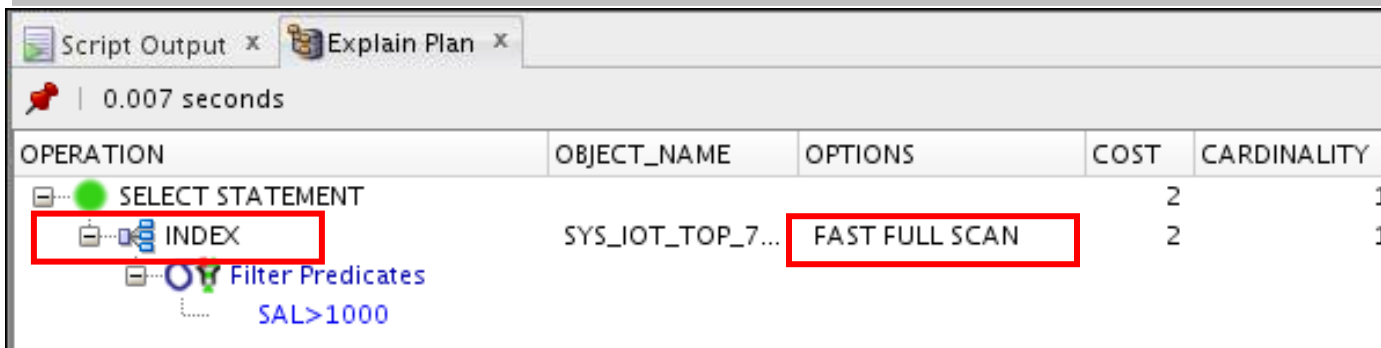
Script Output x Explain Plan x

0.734 seconds

OPERATION	OBJECT_NAME	OPTIONS	COST	CARDINALITY
SELECT STATEMENT			1	1
INDEX	SYS_IOT_TOP_7...	UNIQUE SCAN	1	1

Access Predicates
EMPNO=9999

```
select * from iotemp where sal > 1000;
```



Script Output x Explain Plan x

0.007 seconds



OPERATION	OBJECT_NAME	OPTIONS	COST	CARDINALITY
SELECT STATEMENT			2	1
INDEX	SYS_IOT_TOP_7...	FAST FULL SCAN	2	1

Filter Predicates
SAL>1000

Plan Operations

Bitmap Index

Overview



Key	Start ROWID	End ROWID	Bitmap		
<Blue,	10.0.3,	12.8.3,	100010000	010000000	010010100>
<Green,	10.0.3,	12.8.3,	000101000	000000000	100100000>
<Red,	10.0.3,	12.8.3,	010000000	001100000	000001001>
<Yellow,	10.0.3,	12.8.3,	001000000	100000000	001000010>

Plan Operations

Bitmap Index

Operations

- BITMAP CONVERSION:
 - TO ROWIDS
 - FROM ROWIDS
 - COUNT
- BITMAP INDEX:
 - SINGLE VALUE (similar to a B-tree unique scan)
 - RANGE SCAN (similar to a B-tree range scan)
 - FULL SCAN (all used columns are in index, NOT NULL not needed)
- Bitmap Combining
 - BITMAP AND / OR / MINUS
- Star Transformation (in next topic)
 - BITMAP MERGE
 - BITMAP KEY ITERATION

Bitmap Index

Counting Rows

```
SQL> CREATE TABLE t1 AS select * from all_objects;  
SQL> CREATE BITMAP INDEX t1_b1 on t1(owner);
```

```
SELECT COUNT(*) FROM t1
```

Id	Operation	Name	Starts	E-Rows	A-Rows
0	SELECT STATEMENT		1		1
1	SORT AGGREGATE		1	1	1
2	BITMAP CONVERSION COUNT		1	84499	31
3	BITMAP INDEX FAST FULL SCAN	T1_B1	1		31

Bitmap Index

Access

```
SELECT * FROM cust_dim WHERE cntry_code = 'FR';
```

Id	Operation	Name	Rows	Bytes
0	SELECT STATEMENT		1	45
1	TABLE ACCESS BY INDEX ROWID	CUST_DIM	1	45
2	BITMAP CONVERSION TO ROWIDS			
3	BITMAP INDEX SINGLE VALUE	IX_B2		

```
Predicate: 3 - access("COUNTRY"='FR')
```

```
SELECT * FROM cust_dim WHERE cntry_code > 'FR';
```

Id	Operation	Name	Rows	Bytes
0	SELECT STATEMENT		1	45
1	TABLE ACCESS BY INDEX ROWID	CUST_DIM	1	45
2	BITMAP CONVERSION TO ROWIDS			
3	BITMAP INDEX RANGE SCAN	IX_B2		

```
Predicate: 3 - access("CNTRY_CODE">'FR') filter("CNTRY_CODE">'FR')
```

Plan Operations

Bitmap Index

Combining - Overview

```
SELECT * FROM cust_dim WHERE cntry_code in ('FR','DE');
```

FR	0	0	1	1	1	1	0	0	0	0	0	0
----	---	---	---	---	---	---	---	---	---	---	---	---

OR

0	1	1	1	1	1	0	0	0	0	0	0
---	---	---	---	---	---	---	---	---	---	---	---

DE	0	1	0	0	0	0	0	0	0	0	0	0
----	---	---	---	---	---	---	---	---	---	---	---	---

```
SELECT * FROM cust_dim WHERE cntry_code = 'FR' and gndr_code = 'M';
```

FR	0	0	1	1	1	1	0	0	0	0	0	0
----	---	---	---	---	---	---	---	---	---	---	---	---

AND

0	0	1	0	1	1	0	0	0	0	0	0
---	---	---	---	---	---	---	---	---	---	---	---

M	1	1	1	0	1	1	0	1	0	1	1	1
---	---	---	---	---	---	---	---	---	---	---	---	---

Bitmap Index

Combining - Example

Hint: INDEX_COMBINE

```
SELECT * FROM cust_dim WHERE cntry_code IN ('FR','DE');
```

Id	Operation	Name	Rows	Bytes
0	SELECT STATEMENT		1	45
1	INLIST ITERATOR			
2	TABLE ACCESS BY INDEX ROWID	CUST_DIM	1	45
3	BITMAP CONVERSION TO ROWIDS			
4	BITMAP INDEX SINGLE VALUE	IX_B2		

Predicate: 4 - access("CNTRY_CODE"='DE' OR "CNTRY_CODE"='FR')

```
SELECT * FROM cust_dim WHERE cntry_code = 'FR' AND gndr_code = 'M';
```

Id	Operation	Name	Rows	Bytes
0	SELECT STATEMENT		1	45
1	TABLE ACCESS BY INDEX ROWID	CUST_DIM	1	45
2	BITMAP CONVERSION TO ROWIDS			
3	BITMAP AND			
4	BITMAP INDEX SINGLE VALUE	IX_B1		
5	BITMAP INDEX SINGLE VALUE	IX_B2		

Predicate: 4 - access("GNDR_CODE"='M') 5 - access("CNTRY_CODE"='FR')

Invisible Indexes

- Index is altered as not visible to the optimizer:

```
ALTER INDEX ind1 INVISIBLE;
```

- Optimizer does not consider this index:

```
SELECT /*+ index(TAB1 IND1) */ COL1 FROM TAB1 WHERE ...;
```

- Optimizer considers this index:

```
ALTER INDEX ind1 VISIBLE;
```

- Create an index as invisible initially:

```
CREATE INDEX IND1 ON TAB1(COL1) INVISIBLE;
```

Indexes Recommendations

- Index should be used on large tables filter columns.
- Are not recommended when processing over 10% rows.
- Create indexes after inserting table data.
- Index the correct tables and columns.
- Order index columns for performance.
- Limit the number of indexes for each table (DMLs).
- Drop indexes that are no longer required.
- Consider parallelizing index creation.
- Consider creating indexes with `NOLOGGING`.
- Consider costs and benefits of coalesce or rebuild.
- Consider cost before off or drop PK, UK constraints.

Index Investigation

Why Is Not Used ?

- There are functions being applied to the predicate.
- There is a data type mismatch.
- Statistics are old.
- The column can contain NULLs.
- Using index would actually be slower than not using it.

Why Is Used When Not Needed?

- Wrong low number of rows estimation computed
 - Index statistics are wrong

Q & A

+ Index Scan Quiz

Table Join Operations


HASH JOIN

Hint: `USE_HASH(tab1 tab2 ...)`

- Good in DW star schema – fact join with dimensions
- The small row source (dim) is used to build a hash table.
- The second source (fct) is hash checked against the hash table.

```
SELECT /*+ USE_HASH(f p t) */
  p.prod_name, t.fisc_yr_num, f.sold_amt
FROM sh_sales_fct f
  INNER JOIN sh_time_dim t
    ON t.time_id = f.time_id
  INNER JOIN sh_prod_dim p
    ON p.prod_id = f.prod_id
```


OPERATION	OBJECT_NAME	CARDINALITY	COST
SELECT STATEMENT		918843	1414
HASH JOIN		918843	1414
Access Predicates			
T.TIME_ID=F.TIME_ID			
PART JOIN FILTER (CREATE)	:BF0000	1826	17
TABLE ACCESS (FULL)	SH_TIME_DIM	1826	17
HASH JOIN		918843	1392
Access Predicates			
P.PROD_ID=F.PROD_ID			
TABLE ACCESS (FULL)	SH_PROD_DIM	72	3
PARTITION RANGE (JOIN-FILTER)		918843	1385:
TABLE ACCESS (FULL)	SH_SALES_FCT	918843	1385:



NESTED LOOPS

Hint: `USE NL (tab1 tab2)`

- Good when driving row source is very small.
- All driving row source is scanned.
- Each row returned drives a lookup in inner row source using needed index on join key column
- If driving source is large - many index scans in loop !



```
SELECT /*+ USE_NL(e d) */ e.employee_id, d.department_name
FROM employees e, departments d
WHERE e.department_id = d.department_id AND first_name LIKE 'A%'
```

	Id	Operation	Name	Starts	A-Rows
	0	SELECT STATEMENT		1	10
	1	NESTED LOOPS		1	10
	2	NESTED LOOPS		1	10
*	3	TABLE ACCESS FULL	EMPLOYEES	1	10
*	4	INDEX UNIQUE SCAN	DEPT_ID_PK	10	10
	5	TABLE ACCESS BY INDEX ROWID	DEPARTMENTS	10	10

```
3 - filter("FIRST_NAME" LIKE 'A%')
```

```
4 - access("E"."DEPARTMENT ID"="D"."DEPARTMENT ID")
```


Table Join Operations

SORT MERGE JOIN

Hint: `USE_MERGE (tab1 tab2)`

- Good when join condition is not equality or results are ordered by join key.
- Can generate large and expensive sort.
- First and second row sources are sorted by same join key.
- Sorted rows from both tables are merged.

```
SELECT /*+ USE_MERGE(d e) NO_INDEX(d) */  
       ename, e.deptno, d.deptno, dname  
FROM emp e, dept d  
WHERE e.deptno = d.deptno AND ename > 'A'
```



OPERATION	OBJECT_NAME	OPTIONS	COST	CARDINALI...
SELECT STATEMENT			8	14
MERGE JOIN			8	14
SORT		JOIN	4	4
TABLE ACCESS	DEPT	FULL	3	4
SORT		JOIN	4	14
Access Predicates				
AND				
E.DEPTNO=D.DEPTNO				
SYS_OP_DESCEND(E.DEPT				
Filter Predicates				
E.DEPTNO=D.DEPTNO				
TABLE ACCESS	EMP	FULL	3	14
Filter Predicates				
ENAME>'A'				

Q & A

+ Table Join Quiz

Topic Agenda

Execution Plan Tuning

- Plan Generation
- Plan Evaluation
- Influence on Plan

Execution Plan

Optimizer Decisions

- Selectivity and cardinality derived from statistics or sampling
- Cost calculated for possible plan variants and **smallest cost** variant **used**
- **Choosing objects** to scan:
 - Tables - for large results volume
 - Indexes - for small results volume from big tables and to **avoid sorts or table scans**
- Choosing objects **scan methods**
 - Full scan for large volume reads
- Partition **Pruning** needed ? - static or dynamic if filtering on partition key exists ?
- Choosing **Access Predicate** - first filter - most restrictive
- Choosing Filter **Predicates Order** - most restrictive first
- **PQ plan** needed? Degree of parallelism choice - on execution time
 - Depends on time estimation and number of rows processed
- Row **Distribution Methods** (PQ plan only) - broadcast for small row sources
- **Sorting** and Grouping **Methods** - hash for large volume
- **Join Methods** - hash join for large volume but nested loop for very small
- **Join Order** - small row sources first

Evaluation

Where is the problem?

- Check if plan is changed to worse
- Make SQL Operations ranking (by CPU, time)
- Look for the following:
 - Driving table has the best filter
 - Fewest number of rows are returned to the next step
 - The join method is correct for number of rows
 - Parallel row distribution is good for number of rows
 - Views are correctly used
 - Unintentional Cartesian products
 - Tables/indexes accessed efficiently (also pruning)

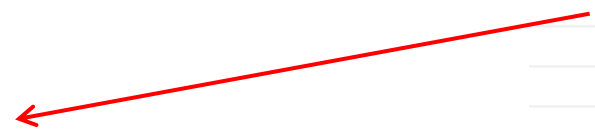
Evaluation

Top SQL Plan Operations

```
SELECT r.samples, u.username, r.sql_id, r.event, r.in_parse,
       r.sql_exec_id, r.sql_exec_start, r.sql_plan_hash_value,
       r.sql_plan_operation, r.sql_plan_options,
       p.object_owner, p.object_name, p.cardinality
FROM ( SELECT * FROM ( SELECT
      count(*) samples, user_id, sql_id, event, in_parse,
      row_number() OVER (PARTITION BY sql_id ORDER BY count(*) DESC) AS rn,
      sql_exec_id, sql_exec_start, sql_plan_hash_value,
      sql_plan_line_id, sql_plan_operation, sql_plan_options
    FROM dba_hist_active_sess_history
    WHERE sample_time > SYSDATE - 8
      AND sql_id = '6mm9w58yznv1k'
      AND sql_plan_hash_value = 1652160185
      AND session_state = 'ON CPU'
      --AND session_state = 'WAITING' and event <> 'resmgr:cpu quantum'
    GROUP BY sql_id, session_state, event, in_parse,
             sql_plan_operation, sql_plan_options, sql_plan_line_id,
             sql_exec_id, sql_exec_start, sql_plan_hash_value, user_id
    ORDER BY samples DESC
  ) WHERE ROWNUM <= 10 ) r
LEFT JOIN dba_hist_sql_plan p ON p.ID = r.sql_plan_line_id
  AND p.sql_id = '6mm9w58yznv1k'
  AND p.plan_hash_value = 1652160185
JOIN all_users u ON r.user_id = u.user_id
ORDER BY samples DESC
```

Plan Change History

```
SELECT COUNT(*) samples,
       sql_exec_start,
       sql_plan_hash_value
FROM dba_hist_active_sess_history
WHERE sample_time > SYSDATE - 8
  AND sql_id = '6mm9w58yznv1k'
  AND session_state = 'ON CPU'
GROUP BY sql_exec_start,
         sql_plan_hash_value
ORDER BY sql_exec_start
```



SMPS	SQL_EXEC_S	HASH_VALUE
3 29	13:13:56	1652160185
6 29	13:24:07	1652160185
4 29	14:05:38	1652160185
5 29	14:13:44	1652160185
4 29	17:12:49	1652160185
3 29	17:19:02	1652160185
4 29	17:25:12	1652160185
4 30	12:40:08	1481012807
3 30	15:32:57	3896590179

Execution Plan Tuning

Plan Modification

- What has influence on Execution Plan ?
 - Optimizer Statistics (described in previous topic)
 - Optimizer Sensitive Parameters (change on session)
 - Stored Outline (depreciated)
 - SQL Profile (main method)
 - Plan Baseline (new method)
 - Hints (as last resort)
- } SQL Plan Management topic

Optimizer Sensitive Parameters


- **CURSOR_SHARING:** SIMILAR, EXACT, FORCE
- **DB_FILE_MULTIBLOCK_READ_COUNT**
- **PGA_AGGREGATE_TARGET** (for DBA but *_AREA_SIZE can change on session)
- **STAR_TRANSFORMATION_ENABLED**
- **RESULT_CACHE_MODE:** MANUAL, FORCE
- **RESULT_CACHE_MAX_SIZE** (for DBA)
- **RESULT_CACHE_MAX_RESULT** (for DBA)
- **RESULT_CACHE_REMOTE_EXPIRATION** (for DBA)
- **OPTIMIZER_INDEX_CACHING**
- **OPTIMIZER_INDEX_COST_ADJ**
- **OPTIMIZER_FEATURES_ENABLED**
- **OPTIMIZER_MODE:** ALL_ROWS, FIRST_ROWS, FIRST_ROWS_n
- **OPTIMIZER_CAPTURE_SQL_PLAN_BASELINES**
- **OPTIMIZER_USE_SQL_PLAN_BASELINES**
- **OPTIMIZER_DYNAMIC_SAMPLING**
- **OPTIMIZER_USE_INVISIBLE_INDEXES**
- **OPTIMIZER_USE_PENDING_STATISTICS**

Q & A

www.lingaro.com


Plan Tuning Workshop

- Run Workload
- Check operation ranking



Plan Tuning

Workshop Comp Round 1-3



- Run statement from script
- Answer quiz questions

SQL Tuning Resources

- Oracle Database Documentation Library

http://docs.oracle.com/cd/E11882_01/index.htm

- Database Performance Tuning Guide

http://docs.oracle.com/cd/E11882_01/server.112/e41573/toc.htm

- Oracle SQL Tuning

http://docs.oracle.com/cd/E28271_01/server.1111/e16638/sql_overview.htm

- Optimizer Statistics

<http://www.oracle.com/technetwork/database/bi-datawarehousing/twp-optimizer-stats-concepts-110711-1354477.pdf>