

Oracle SQL Tuning

3 parts (BASIC, INTERMEDIATE, ADVANCED)

www.lingaro.com



Training Agenda

- **BASIC**
 - SQL Tuning Introduction
 - Instrumentation & Rewrite SQL Text
 - SQL Plan Reading
- **INTERMEDIATE**
 - Optimizer & Statistics
 - SQL Plan Tuning
 - SQL Plan Operations
- **ADVANCED**
 - SQL Processing & Cursor Sharing
 - Advanced Plan Operations
 - SQL Plan Transformations
 - SQL Plan Management

Training Agenda

- Theory
- Quiz Comp
- Workshop
- Workshop Comp
- Winner Awards On WINS

Topic Agenda

Tuning Techniques (not prepared yet)

- Using Automatic SQL Tuning
- Using SQL Tuning & SQL Access Advisors (**dbms_sqltune** or OEM)
- Using Oracle Enterprise Manager Console OEM
- **Using DBA_HIST and V\$ performance views (described here)**
- Using Test Executions
- Using SQL Decomposition
 - Test Execution on SQL parts

Oracle SQL Tuning

ADVANCED

www.lingaro.com



Topic Agenda

SQL Processing

- Introduction
- Processing Steps
- Workload Characteristics
- SQL Memory Structures
- Cursor Sharing

Introduction

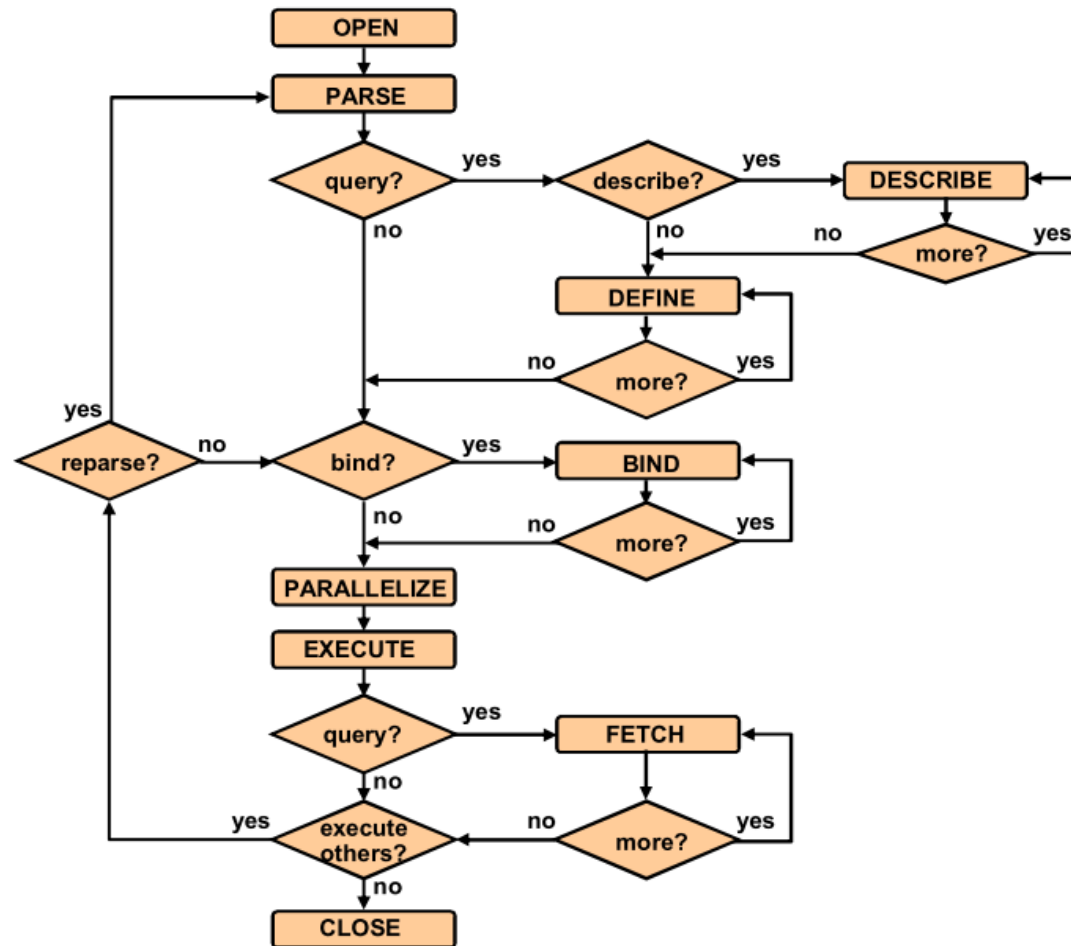
SQL Processing

- Is done inside Oracle Database Server SQL engine
- Oracle client send SQL text & receive results from server
 - by network - connected via database service
 - by local connection - connected to local instance using ORACLE_SID
- **Server steps:**

	Query	DML	DDL
– Parse	v	v	v
– Bind	v	v	
– Execute	v	v	
– Fetch	v		
- **SQL optimization** is done by Oracle Optimizer **during parse**
 - Has significant **influence on execution and fetch phases costs**
 - **Parse** phase cost is **more important in OLTP** workload

Introduction

SQL Processing Steps



Workload characteristic

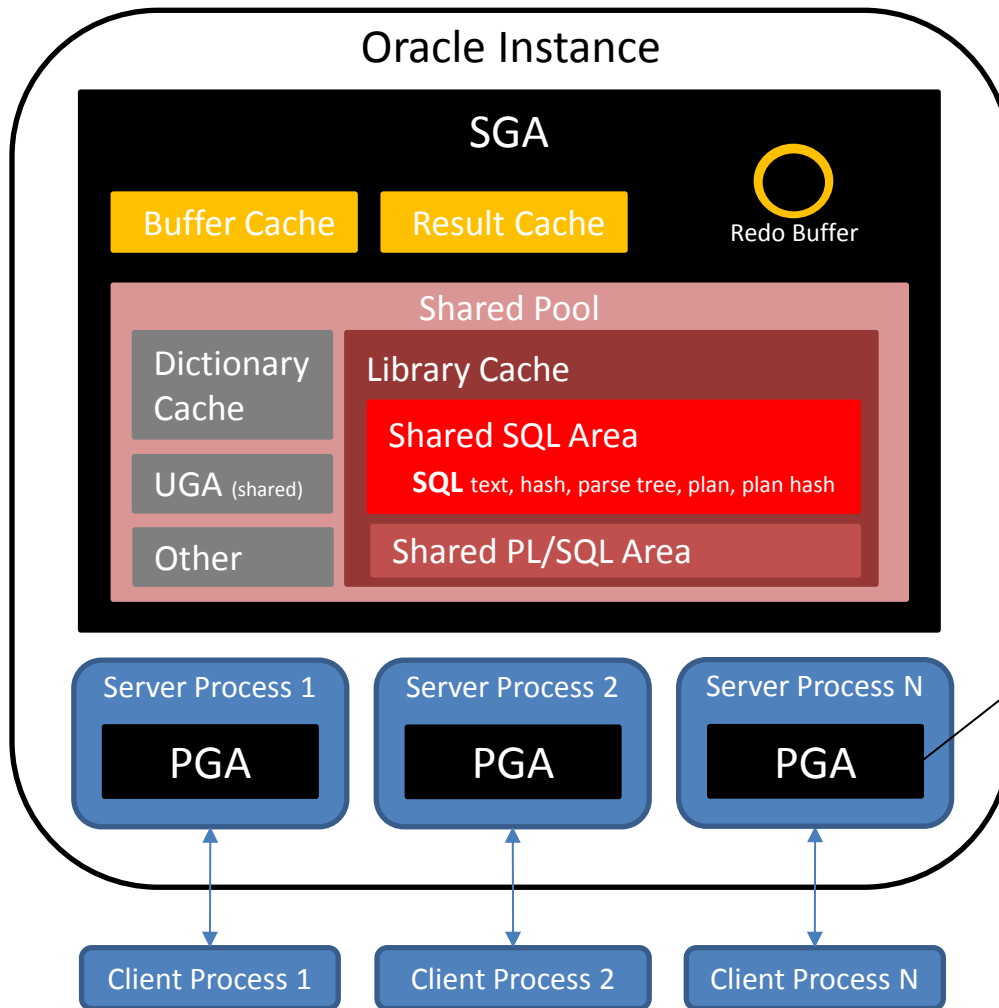
- OLTP workload

- Large number of short time **single/few rows query and DML**
- All SQLs summary **parse** cost is big and **more important**
- SQL **optimization is less important**
- Main tuning **goal - response time**
- **Critical** resource - **CPU, memory**
- Top blocking **event - contention**

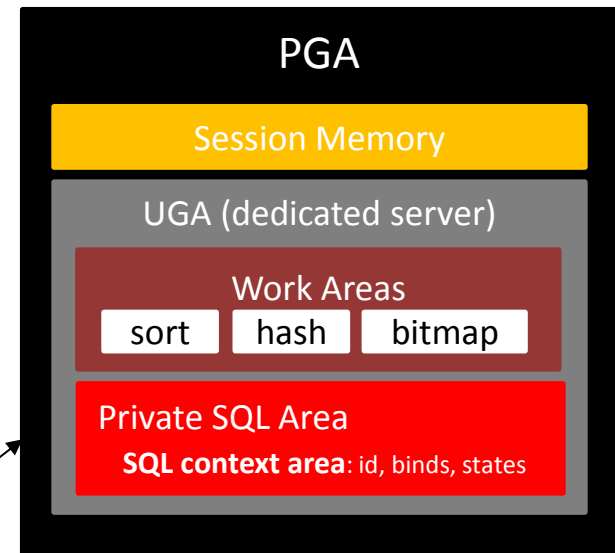
- DSS/Warehouse workload

- average number of **long running** large data volume **reports and ETL**
- **execution** and fetch cost is bigger and **more important**
- SQL **optimization is very important**
- Main tuning **goal - throughput**
- **Critical** resource - **CPU, storage IO**
- Top blocking **event - ON CPU - processing**

SQL Memory Structures




Cursor is
name, pointer or handle
to SQL context area
(linked to child of shared SQL area)



SQL Processing

Cursor Sharing

One SQL
Example



- Multiple executions of the same SQL uses the same shared SQL area
- Reduces hard parses - plan generation memory & CPU costs
- Can be done by
 - Reusing the same SQL statement many times using code standardization and libraries
 - Changing many similar SQLs to one SQL using Bind Variables instead of literal values in SQL text
 - by developer - PL/SQL source code change - use of PL/SQL variables
 - e.g. `chanl_id = 102, chanl_id = 105, ... -> chanl_id = l_chanl_id`
 - or SQL script session variables - host variables
 - e.g. `chanl_id = 102, chanl_id = 105, ... -> chanl_id = :g_chanl_id`
 - by DBA or developer - set **CURSOR_SHARING** parameter - bind variables automatically generated
 - EXACT - default value - SQL text need to be identical to share cursor
 - if text is different separate SQL area is allocated in library cache
 - SIMILAR - bind variables used if plan not need to be changed when bind value is changed
 - FORCE - deterministic bind variables usage
 - e.g. `ALTER SESSION SET cursor_sharing='SIMILAR';`

Cursor Sharing

Reusing the same SQL statement many times

- SQL source code standardization
 - Like: keywords uppercase, identifiers lowercase, single space separator, two space indentation
 - Standards make SQLs to be more likely reused - Example:

```
SELECT  /*+PARALLEL(2)*/
        TO_CHAR(f.sales_date,'YY/MM/DD') AS wk_end_date,
        f.gln                               AS store_id,
        SUM(pos_amt)                        AS sales_amt
FROM    gpos_cp_fads PARTITION("3663_CP") f,
        prod_extrn_dads p
WHERE   f.srce_sys_id = 3663
        AND f.prod_extrn_id = p.prod_extrn_id
        AND f.srce_sys_id = p.srce_sys_id
        AND ((f.pos_amt != 0 AND f.pos_amt IS NOT NULL)
              OR (f.pos_qty != 0 AND f.pos_qty IS NOT NULL))
GROUP BY
        f.sales_date,
        f.gln;
```

- Using shared library modules
 - PLSQL procedures containing static SQL statements
 - Avoid using dynamic SQL
 - If dynamic SQL is necessary then use bind variables inside dynamic SQL code

Bind Variables




Not used

~~Cursor sharing~~

```
SELECT * FROM prod_dim WHERE min_price_amt > 11000;
```

```
SELECT * FROM prod_dim WHERE min_price_amt > 9000;
```

```
SELECT * FROM prod_dim WHERE min_price_amt > 6500;
```



Library cache

- **Pros**

- Each filter constant value have its own execution plan

- **Cons**

- Each constant value produce **separate SQL statement with different SQL_ID**
- **Many plan tuning techniques needs constant SQL_ID for tuned statement**
- Large number of sql statements
- Large memory usage and statements compilation CPU costs

Bind Variables

Used

Cursor sharing

11000


```
SELECT * FROM prod_dim WHERE min_price_amt > :min_price;
```

9000

```
SELECT * FROM prod_dim WHERE min_price_amt > :min_price;
```

6500

```
SELECT * FROM prod_dim WHERE min_price_amt > :min_price;
```



Library cache

- **Pros**

- One statement executed many times
- One plan for all executions - lower CPU cost for compilation
- **All plan tuning techniques possible**

- **Cons**

- One SQL plan can be not optimal for all bind variable values
but Adaptive Cursor Sharing makes many plans for 1 SQL possible

Cursor Sharing

Matching Signatures

```
select 0 from dual where dummy = 'A';  
select 0 from dual where dummy = 'B';  
select 0 from dual where dummy = 'C';  
select 0 from dual where dummy = 'A';
```

```
SELECT  force_matching_signature, exact_matching_signature, sql_text  
FROM    v$sqlarea  
WHERE   sql_text like 'select 0%';
```

⚡ FORCE_MATCHING_SIGNATURE	⚡ EXACT_MATCHING_SIGNATURE	⚡ SQL_TEXT
13154199455204052618	8255937935626560799	select 0 from dual where dummy = 'A'
13154199455204052618	8255937935626560799	select 0 from dual where dummy = 'A'
13154199455204052618	4228015713888331644	select 0 from dual where dummy = 'C'
13154199455204052618	16821029007583452703	select 0 from dual where dummy = 'B'

Cursor Sharing

Finding SQL candidates to use

– source code standardization

```
SELECT
  '||exact_matching_signature||',
FROM v$sqlarea
WHERE exact_matching_signature <> 0
GROUP BY exact_matching_signature
HAVING count(*) > 2;

SELECT exact_matching_signature, sql_id,
       count(*) OVER (PARTITION BY
                       exact_matching_signature) cnt
FROM v$sqlarea
WHERE exact_matching_signature IN (
  '16523011601676190874',
  '2945800283839423321',
  '2672114946588399948',
  '3854345704683617056',
  '14461745472358877479'
) ORDER BY cnt DESC,1;
```

– bind variables implementation

```
SELECT
  '||force_matching_signature||',
FROM v$sqlarea
WHERE force_matching_signature <> 0
GROUP BY force_matching_signature
HAVING count(*) > 100;

SELECT force_matching_signature, sql_id,
       count(*) OVER (PARTITION BY
                       force_matching_signature) cnt
FROM v$sqlarea
WHERE force_matching_signature IN (
  '7881417843405219978',
  '16841543523900648019',
  '1106612903482690771',
  '12296313255013650141',
  '6556405521552985841',
  '8931422929029301538'
) ORDER BY cnt DESC,1;
```


Cursor Sharing

SQL Information from V\$SQLSTAT

```
SELECT sql_id, version_count, executions, invalidations, parse_calls,  
       rows_processed, cpu_time, elapsed_time, avg_hard_parse_time  
FROM v$sqlstats WHERE sql_text LIKE 'update ENGINE_INSTANCES%'
```

SQL_ID	VERSION_COUNT	EXECUTIONS	INVALIDATIONS	PARSE_CALLS	ROWS_PROCESSED	CPU_TIME	ELAPSED_TIME	AVG_HARD_PARSE_TIME
akbss8gya5a28	5	182558	0	182558	182558	167678549	392783430	19182

- executions = parse calls (soft parse)
 - client doing single SQL execution in one database call
- executions > parse calls
 - single anonymous PL/SQL call run doing many executions of the same SQL
 - only first run of PL/SQL procedure increase PARSE_CALLS of included SQL
- executions < parse calls
 - statement parsed but not executed
 - statement additionally parsed during execution (purged or invalidated)

Cursor Sharing

SQL Cursor Versions (Childs) Information from V\$SQL

```
SELECT sql_id, hash_value, address, plan_hash_value,  
       child_number, executions, parsing_schema_name  
FROM v$sql WHERE sql_id = 'akbss8gya5a28' ORDER BY child_number;
```

SQL_ID	HASH_VALUE	ADDRESS	PLAN_HASH_VALUE	CHILD_NUMBER	EXECUTIONS	PARSING_SCHEMA_NAME
akbss8gya5a28	4238518344	00000018C23FEBB8	3804546214	0	33843	ADWGP_REA
akbss8gya5a28	4238518344	00000018C23FEBB8	3804546214	1	94318	ADWGP_REA
akbss8gya5a28	4238518344	00000018C23FEBB8	2055260210	2	13945	ADWGP_REA_LOAD_BAL
akbss8gya5a28	4238518344	00000018C23FEBB8	2055260210	3	34772	ADWGP_REA_LOAD_BAL
akbss8gya5a28	4238518344	00000018C23FEBB8	4278141064	4	6037	ADWG_IBA_CRS_ETL

```
SELECT child_number, auth_check_mismatch, bind_mismatch  
FROM v$sql_shared_cursor WHERE sql_id='akbss8gya5a28';
```

CHILD_NUMBER	AUTH_CHECK_MISMATCH	BIND_MISMATCH
0	N	N
1	N	Y
2	Y	N
3	Y	Y
4	Y	N

Cursor Sharing

Adaptive Cursor Sharing: Overview

- Allows for **intelligent cursor sharing** for statements that use **bind variables**
- Works for bind variables produced **in source code** by developers **or** produced automatically using **CURSOR_SHARING parameter**.
- Is used **to trade off** between cursor sharing and plan optimization
- Has the following benefits:
 - Automatically **detects** when different executions would **benefit from different execution plans** - statements marked as “**bind aware**”
 - **Limits** the number of generated **child cursors to a minimum**
 - Provides an automated mechanism that **cannot be turned off**

Cursor Sharing

Adaptive Cursor Sharing: Architecture

Bind-sensitive
cursor







System
observes
statement



```
SELECT * FROM emp WHERE sal = :1 and dept = :2
```

Bind-aware
cursor



Cursor Sharing

Adaptive Cursor Sharing: Views

V\$SQL

Two columns show whether a cursor is **bind sensitive** or **bind aware**.

V\$SQL_CS_HISTOGRAM

Shows the distribution of the **execution count** across the execution **history histogram**

V\$SQL_CS_SELECTIVITY

Shows the **selectivity cubes** stored for **every** predicate containing a **bind variable** and whose selectivity is used in the cursor sharing checks

V\$SQL_CS_STATISTICS

Shows **execution statistics** of a cursor using different bind sets

SQL Processing Steps

Parse Step

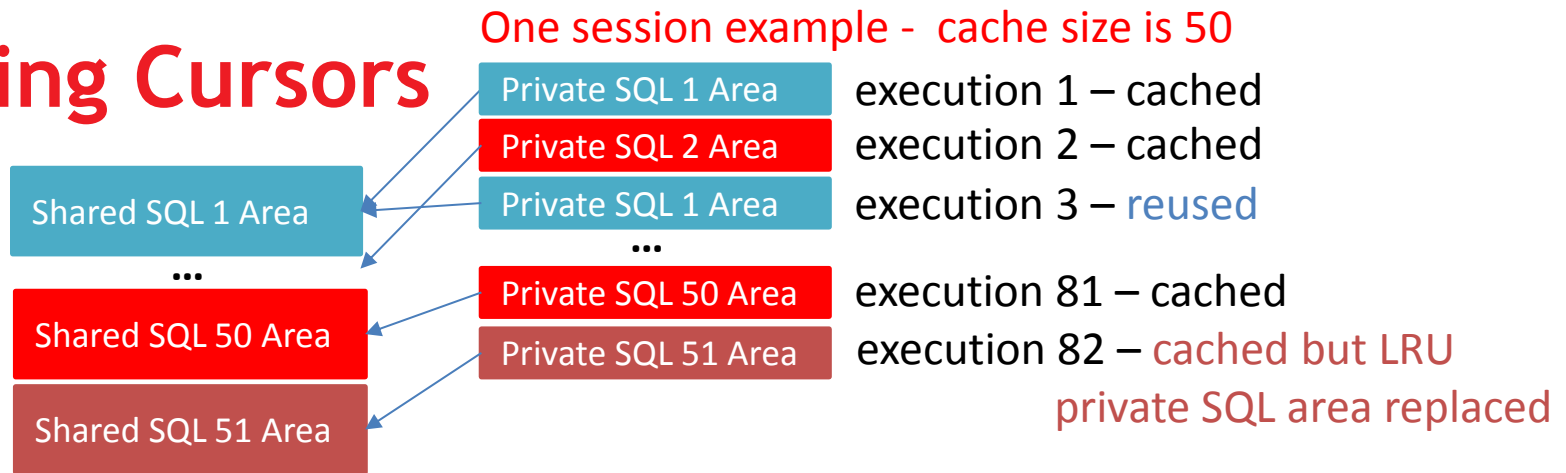
Soft Parse

- Determining statement **type** (query, DML, DDL)
- **Syntax** Checks: keywords, grammar, rules
- **Semantic** Analysis Checks:
 - referenced objects exists
 - have access to referenced objects
 - ambiguities
- Finding if same statement **was executed** (Shared Pool Check)
 - **Using hash** from SQL text to find same statement was executed
 - * on same session (exists in private SQL area if cursor is cached)
 - * on different sessions (exists in shared SQL areas)

Hard Parse

- Optimization & Estimation
 - Execution Plan Generation
- done by Oracle Optimizer
(next topic)

Caching Cursors



- Multiple executions of the same SQL in the same session
 - Can use the same session cached cursor
- Persistent part of private area is retained even cursor is closed
 - (after 3-th same SQL execution on the session)
- Can reduce private area generation costs - soft parse will be cheaper
- Enable to find shared SQL area handle in private SQL area
 - decrease library cache latch wait events and parsing serialization
 - increase scalability
- Size can be modified by

```
ALTER SESSION SET session_cached_cursors = 100 -- default 50 cursors cached
```

- Diagnostic

```
SELECT name, value FROM v$statname NATURAL JOIN v$mystat  
WHERE LOWER(name) LIKE '%cursor ca%';
```

NAME	VALUE
session cursor cache hits	49
session cursor cache count	37

SQL Processing Steps

Open Step (environment and language dependent)

- Generating cursor identifier before parse phase
- PL/SQL example

```
c_cur_num INTEGER;  
BEGIN  
    c_cur_num := dbms_sql.open_cursor;
```

- PL/SQL OPEN statement

```
c_cur CURSOR IS SELECT ...;  
BEGIN  
    OPEN c_cur;
```

- Is doing almost all processing steps
open, parse, bind, execution
- Is not doing fetch step

SQL Processing Steps

Other (less important) Steps

- Describe results (query only)
 - Datatypes, lengths, and names of columns for query result determination
 - Only necessary if the characteristics of a query's result are not known
- Define output (query only)
 - Location, size, and datatype of variables used to receive fetched value.
 - Oracle performs datatype conversion here if necessary.
- Bind
 - Assign value to bind variable
- Close cursor
 - Release none persistent part of private SQL area


Q & A

+ Quiz

Q & A

+ Workshop Comp

Cursor Sharing Workshop



- Implement bind variables

- Use bind variables in `wksh_ts_cur_sh_pro` procedure
- See statistics for test execution after reconnect before and after change

```
SELECT name, value FROM v$mystat NATURAL JOIN v$statname
WHERE name IN ('parse time cpu', 'parse count (hard)');
```

- before

parse time cpu	33
parse count (hard)	504

- after

parse time cpu	0
parse count (hard)	2

Topic Agenda

Advanced Plan Operations

- Sort & Grouping Operations
- Partition Pruning
- Parallel Plan
 - Introduction
 - Interpretation
 - Rows Distribution
 - Tuning

Sort Operations

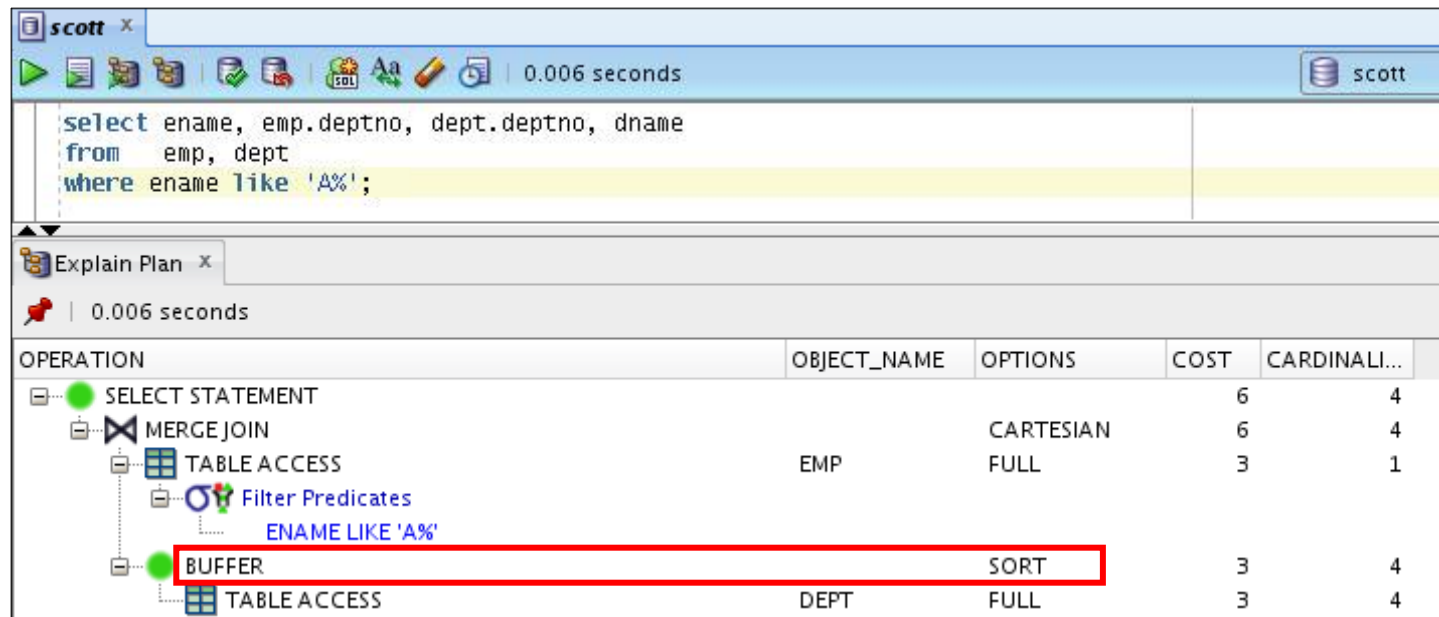
List

- **SORT** operators
 - **AGGREGATE** Single row from group function
 - **UNIQUE** To eliminate duplicates (DISTINCT, UNION, MINUS, INTERSECT)
 - **JOIN** Precedes a merge join
 - **GROUP BY**
 - **ORDER BY**
- **HASH** operators
 - **UNIQUE** Equivalent to **SORT UNIQUE**
 - **GROUP BY** Equivalent to **SORT GROUP BY** but not sorted results
 - If you want ordered results, always use **ORDER BY**
- **BUFFER SORT**

Sort Operations

BUFFER SORT

- Used **when** operation **needs all** the **input** data **before** it can **start**.
- Used to **avoid multiple table scans** against real data blocks.
- **Uses a temporary table or a UGA sort area** to store intermediate data.
- Swaps sort area to temporary tablespace when UGA sort area is full.
- **Data is not necessarily sorted - it is not sort itself.**



scott x | 0.006 seconds

```
select ename, emp.deptno, dept.deptno, dname
from   emp, dept
where  ename like 'A%';
```

Explain Plan x | 0.006 seconds

OPERATION	OBJECT_NAME	OPTIONS	COST	CARDINALI...
SELECT STATEMENT			6	4
MERGE JOIN		CARTESIAN	6	4
TABLE ACCESS	EMP	FULL	3	1
Filter Predicates ENAME LIKE 'A%'				
BUFFER		SORT	3	4
TABLE ACCESS	DEPT	FULL	3	4

Other Operations

Min/Max FIRST ROW

- Retrieves **only the first row** selected by a query.
- **Stops** accessing the data **after** the **first value** is returned.
- Works with the **index range** scan and the index **full scan**.
- On slide - there is an index on the quantity_on_hand column.

```
SELECT MIN(quantity_on_hand)
FROM INVENTORIES
WHERE quantity_on_hand < 500;
```

OPERATION	OBJECT_NAME	OPTIONS	COST	CARDINALI...
SELECT STATEMENT			2	1
SORT		AGGREGATE		1
FIRST ROW			2	1
INDEX	INV_QTY_INDEX	RANGE SCAN ...	2	1
Access Predicates				
QUANTITY_ON_HAND<500				

Other Operations

Hint: **RESULT_CACHE**

RESULT CACHE

- Use for SQL executed frequently on large volume but returning small results
- SQL results are stored in server memory and used on next execution
- Cached results are invalidated after source table modification
- Monitoring
 - V\$RESULT_CACHE_STATISTICS - result cache settings and memory usage statistics
 - V\$RESULT_CACHE_MEMORY - memory blocks in result cache
 - V\$RESULT_CACHE_OBJECTS - objects whose results are in
 - V\$RESULT_CACHE_DEPENDENCY - what results are dependent on
 - USER_TABLES - RESULT_CACHE column that shows the result cache mode: DEFAULT/FORCE/MANUAL

```
SELECT /*+ RESULT_CACHE */ deptno, AVG(sal)
FROM emp GROUP BY deptno;
```

OPERATION	OBJECT_NAME	OPTIONS	COST	CARDINALI...
SELECT STATEMENT			4	3
RESULT CACHE	54fsvk2n0a...			
HASH		GROUP BY	4	3
TABLE ACCESS	EMP	FULL	3	14

Partition Pruning

Examples

- **Static**

```
SELECT * FROM sales_fct WHERE time_id = TO_DATE('2001.01.01');
```

In Pstart, Pstop:

partition numbers

Operation for range filter:

PARTITION RANGE ITERATOR

All partitions chosen:

PARTITION RANGE ALL

Id	Operation	Name	Pstart	Pstop	

0	SELECT STATEMENT				
1	PARTITION RANGE SINGLE		17	17	
* 2	TABLE ACCESS FULL	SALES	17	17	

- **Dynamic**

```
SELECT SUM(sold_amt) FROM sales_fct WHERE time_id IN
(SELECT time_id FROM time_dim WHERE fisc_yr = 2000);
```

In Pstart, Pstop:

KEY – data convert

KEY(SQ) – subquery

KEY(I) – bind variable

Id	Operation	Name	Pstart	Pstop	

0	SELECT STATEMENT				
1	SORT AGGREGATE				
* 2	HASH JOIN				
* 3	TABLE ACCESS FULL	TIMES			
4	PARTITION RANGE SUBQUERY		KEY (SQ)	KEY (SQ)	
5	TABLE ACCESS FULL	SALES	KEY (SQ)	KEY (SQ)	

Partition Pruning Bloom Filter

Hint: PX_JOIN_FILTER
Opp.: NO_PX_JOIN_FILTER

```

-----
| Id | Operation                                | Name      | Pstart | Pstop |
-----
|  0 | SELECT STATEMENT                        |           |        |        |
|  1 | PX COORDINATOR                          |           |        |        |
|  2 | PX SEND QC (RANDOM)                     | :TQ10002  |        |        |
|*  3 | FILTER                                  |           |        |        |
|  4 | HASH GROUP BY                           |           |        |        |
|  5 | PX RECEIVE                              |           |        |        |
|  6 | PX SEND HASH                            | :TQ10001  |        |        |
|  7 | HASH GROUP BY                           |           |        |        |
|*  8 | HASH JOIN                               |           |        |        |
|  9 | PART JOIN FILTER CREATE                 | :BF0000   |        |        |
| 10 | PX RECEIVE                              |           |        |        |
| 11 | PX SEND PARTITION (KEY)                 | :TQ10000  |        |        |
| 12 | PX BLOCK ITERATOR                       |           |        |        |
| 13 | TABLE ACCESS FULL                     | CUSTOMERS |        |        |
| 14 | PX PARTITION HASH JOIN-FILTER           |           | :BF0000 | :BF0000 |
|* 15 | TABLE ACCESS FULL                     | SALES     | :BF0000 | :BF0000 |
-----

```

Predicate Information (identified by operation id):

```

-----
3 - filter(COUNT(SYS_OP_CSR(SYS_OP_MSR(COUNT(*)),0))>100)
8 - access("S"."CUST_ID"="C"."CUST_ID")
15 - filter("S"."TIME_ID"=TO_DATE(' 1999-07-01 00:00:00',
                                     'syyy-mm-dd hh24:mi:ss'))

```

Bloom Filters

- Data structure (array)
 - Tell you rapidly whether an element is present in a set
 - Memory efficient
- Contains Bit Vector and uses Hash Functions
 - Modified during each set element addition
 - Used to check - element is in set ?
 - False positive possible (can return more rows then needed but do not lose any)
 - Good for large number of rows pre elimination
- Oracle can use it to:
 - reduce data communication between PX slaves (good for RAC)
 - join-filter in partition pruning
 - support result caches
 - filter data from storage server cells in Exadata

Parallel Plan

```
SELECT /*+ PARALLEL(4) */ t.fisc_yr_num, SUM(f.sold_amt) AS sold_amt
FROM sh_sales_fct f INNER JOIN sh_time_dim t ON t.time_id = f.time_id
GROUP BY t.fisc_yr_num
```

Id	Operation	Name	TQ	IN-OUT	PQ Distrib
0	SELECT STATEMENT				
1	<u>PX COORDINATOR</u>				
2	PX SEND QC (RANDOM)	:TQ10004	Q1,04	P->S	QC (RAND)
3	HASH GROUP BY		Q1,04	PCWP	
4	<u>PX RECEIVE</u>		Q1,04	PCWP	
5	PX SEND HASH	:TQ10003	Q1,03	P->P	HASH
6	HASH GROUP BY		Q1,03	PCWP	
7	HASH JOIN		Q1,03	PCWP	
8	<u>PX RECEIVE</u>		Q1,03	PCWP	
9	PX SEND HASH	:TQ10001	Q1,01	P->P	HASH
10	VIEW	VW_GBC_5	Q1,01	PCWP	
11	HASH GROUP BY		Q1,01	PCWP	
12	<u>PX RECEIVE</u>		Q1,01	PCWP	
13	PX SEND HASH	:TQ10000	Q1,00	P->P	HASH
14	HASH GROUP BY		Q1,00	PCWP	
15	PX BLOCK ITERATOR		Q1,00	PCWC	
16	TABLE ACCESS FULL	SH_SALES_FCT	Q1,00	PCWP	
17	<u>PX RECEIVE</u>		Q1,03	PCWP	
18	PX SEND HASH	:TQ10002	Q1,02	P->P	HASH
19	PX BLOCK ITERATOR		Q1,02	PCWC	
20	TABLE ACCESS FULL	SH_TIME_DIM	Q1,02	PCWP	

Parallel Plan

Row Distribution

- HASH
 - Hash function applied to value of the join column
 - Distribute to the workload on the corresponding hash value
- Round Robin
 - Randomly but evenly distributes the data among the consumers
- **Broadcast**
 - **The size of one of the result sets is small**
 - Sends a copy of the data to all consumers
- Range
 - Typically used for parallel sort operations
 - Individual parallel servers work on data ranges
 - QC doesn't sort just present the parallel server results in the correct order
- Partitioning Key Distribution - PART (KEY)
 - Assumes that the target table is partitioned
 - Partitions of the target tables are mapped to the parallel servers
 - Producers map each row to consumer based on partitioning column

Parallel Plan & PX Tuning

Hint: `STATEMENT_QUEUEING`
Opp.: `NO_STATEMENT_QUEUEING`

- Validate if correct distribution method is used
 - E.g. BROADCAST only for small number of rows
- Check if parallel reduction or serialization is not used
 - for intermediate operations on large volume (e.g. serial PL/SQL fn)
- Find SQL candidates for tuning by comparing
 - SQL elapsed time is very far from DEGREE times smaller from
 - SQL database time
- Check if one SQL slaves are not locking each other
- Set DOP on large tables instead of using hints if possible
- Set `PARALLEL_DEGREE_POLICY=AUTO` on session
 - To enable DEGREE auto-tuning, [queuing](#) and in-memory PX

Parallel Plan & PX Tuning

Hint: `STATEMENT_QUEUEING`
Opp.: `NO_STATEMENT_QUEUEING`

- Validate if correct distribution method is used
 - E.g. BROADCAST only for small number of rows
- Check if parallel reduction or serialization is not used
 - for intermediate operations on large volume (e.g. serial PL/SQL fn)
- Find SQL candidates for tuning by comparing
 - SQL elapsed time is very far from DEGREE times smaller from
 - SQL database time
- Check if one SQL slaves are not locking each other
- Set DOP on large tables instead of using hints if possible
- Set `PARALLEL_DEGREE_POLICY=AUTO` on session
 - To enable DEGREE auto-tuning, [queuing](#) and in-memory PX

Q & A

+ Quiz

Advanced Plan Operations Workshop



- Check if partition pruning is working for SQL

```
SELECT /*+GATHER_PLAN_STATISTICS*/  
      t.fisc_mth_num, SUM(f.sold_amt) AS sold_amt  
FROM sh_sales_fct f INNER JOIN sh_time_dim t  
      ON t.time_id = f.time_id  
WHERE t.fisc_yr_num = 2000 GROUP BY t.fisc_mth_num
```

- Try to use parallel(4) hint - variant 1
and parallel with pq_distribute hints - variant 2
 - Compare SQL plan and execution time

Topic Agenda

Plan Transformations

- Introduction
- Star Transformation
- OR expansion
- Simple / Complex(11g) View Merging
- Filter Push Down (Non-mergable View)
- Subquery Unnesting
- Join Predicate Push Down - 11g
- Cost Based Transformation - 11g
 - Join Factorization & Elimination - 11g
 - COUNT(col) to COUN(*) - 11g
 - Transitive Closure
 - Table Elimination (PK-FK, Outer Join)
 - Predicate Move Around
 - Select, Project, Join
 - Order By Elimination
 - Set Join Conversion

Plan Transformation

Introduction


- **Oracle** can automatically **rewrite SQL** statements
- e.g. subquery can be modified to join
- New statement performance should be better
- Usage **decision** is based **on estimated cost**
- Results are the same
- Very hard to correlate plan operations
with original SQL text parts

Star Join

Hint: NO_STAR_TRANSFORMATION

No transformation

```
SELECT ch.chanl_class_name, c.cust_city_name, t.cal_qrt_code,
       SUM(f.sold_amt) sales_amt
FROM sales_fct f, time_dim t, cust_dim c, chanl_dim ch
WHERE f.time_id = t.time_id
      AND f.cust_id = c.cust_id
      AND f.chanl_id = ch.chan_id
      AND c.cust_st_pro_code = 'CA'
      AND ch.chanl_name IN ('Internet','Catalog')
      AND t.cal_qrt_code IN ('1999-Q1','1999-Q2')
GROUP BY ch.chanl_class_name, c.cust_city_name, t.cal_qrt_code;
```



Plan Transformations



Id	Operation	Name	

0	SELECT STATEMENT		
1	HASH GROUP BY		
* 2	HASH JOIN		
* 3	TABLE ACCESS FULL	CHANL_DIM	
* 4	HASH JOIN		
* 5	TABLE ACCESS FULL	TIME_DIM	
* 6	HASH JOIN		
* 7	TABLE ACCESS FULL	CUST_DIM	
8	TABLE ACCESS FULL	SALES_FCT	

Star Transformation

Overview

Hint: `STAR_TRANSFORMATION`
`FACT(tab) NO_FACT(tab)`

- Carried out in two phases:
 - First, **identify** interesting **fact rows** using bitmap indexes based on dimensional filters (phase 1). 
 - **Join them** to the dimension tables (phase 2). 
- Requirements
 - Create **bitmap indexes** on **fact** tables **foreign keys**.
 - Set `STAR_TRANSFORMATION_ENABLED` to `TRUE`
 - **At least two dimensions** and one fact table
 - Gather **statistics** on all corresponding objects.

Star Transformation

Restrictions

Hint: `STAR_TRANSFORMATION`

- Queries are **not transformed** if containing:
 - bind variables
 - references to remote fact tables
 - antijointed tables
 - references to unmerged views
 - `NO_STAR_TRANSFORMATION` hint used

Star Transformation

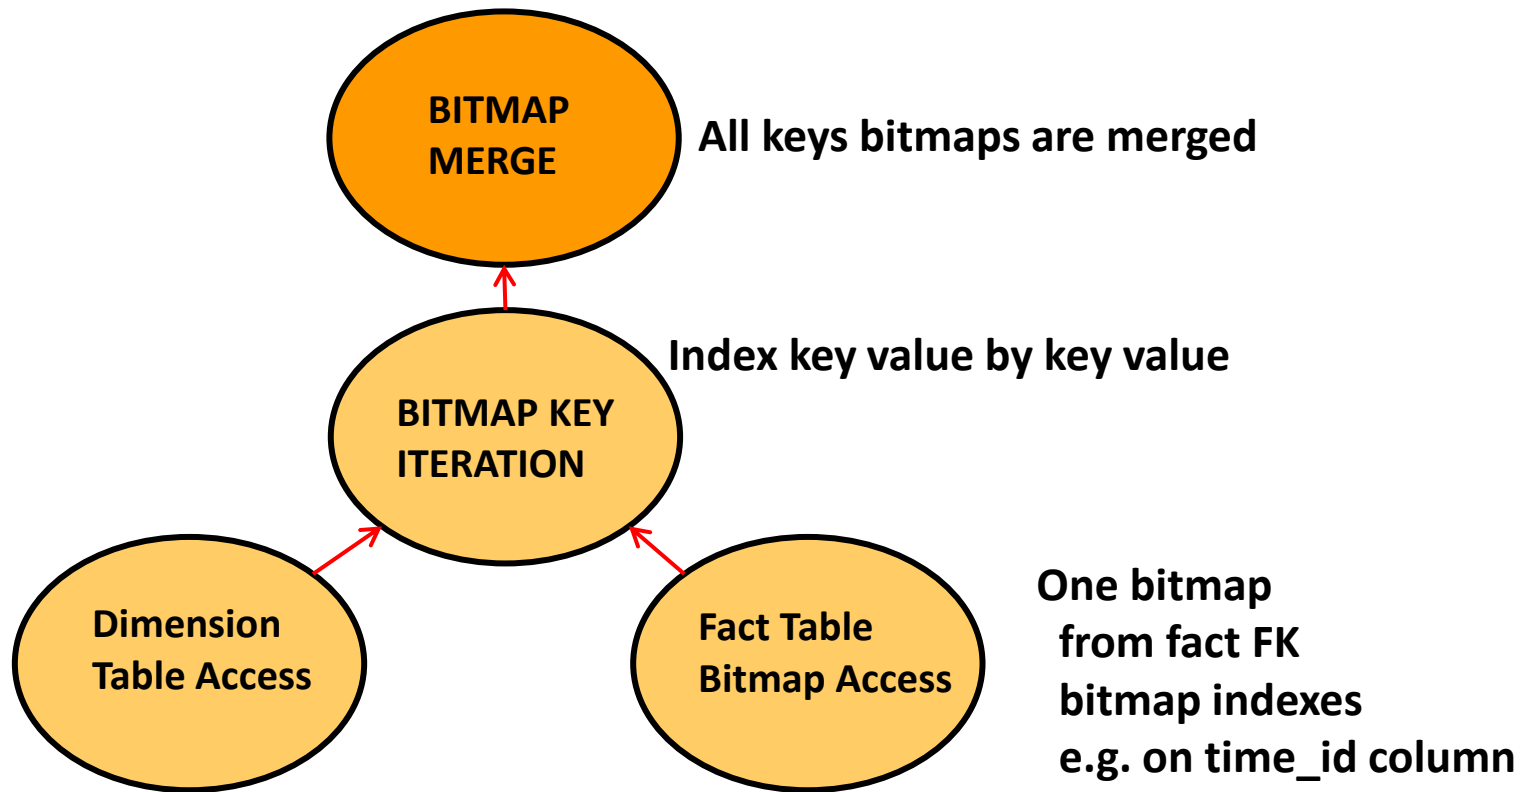
Phase 1

```
SELECT f.sold_amt, f.time_id, f.cust_id, f.chanl_id
FROM sales_fct f
WHERE time_id IN (SELECT time_id
                   FROM time_dim
                   WHERE cal_qrt_code IN
                      ('1999-Q1', '1999-Q2'))
AND cust_id IN (SELECT cust_id
                FROM cust_dim
                WHERE cust_st_prov = 'CA')
AND chanl_id IN (SELECT chanl_id
                 FROM chanl_dim
                 WHERE chanl_name IN
                    ('Internet', 'Catalog'));
```


Star Transformation

Phase 1

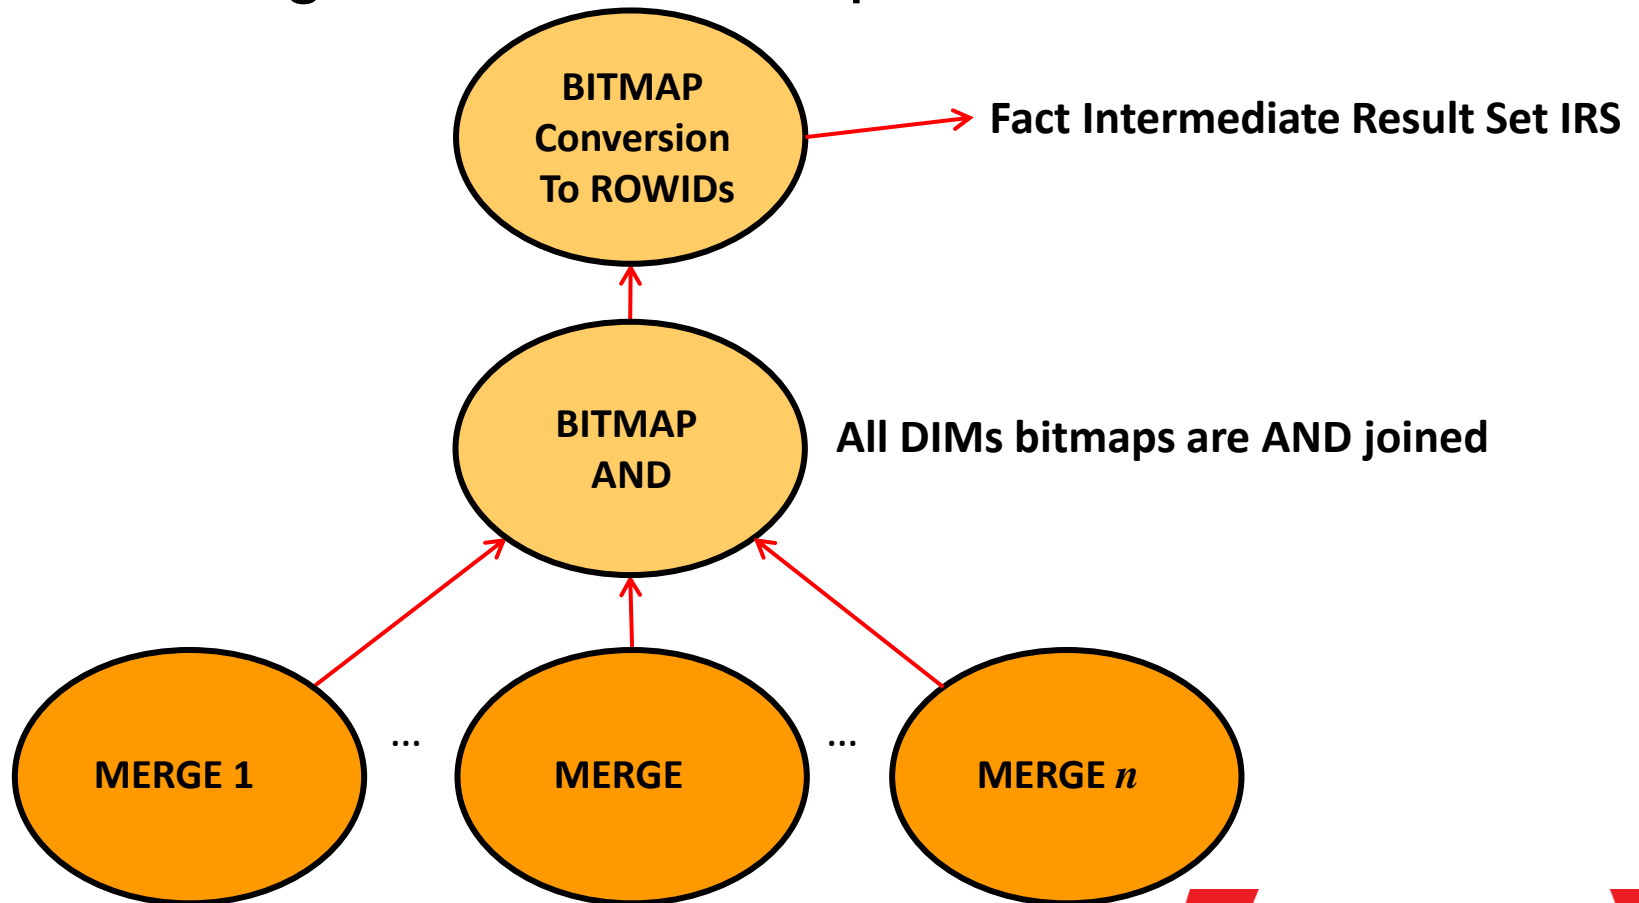
- Retrieving fact rows bitmap - one dimension



Star Transformation

Phase 1


- Retrieving fact rows bitmap - all dimensions



Star Transformation

Phase 2

- Join selected fact rows with dimension data



Star Transformation

Plan

```
SORT GROUP BY
  HASH JOIN
    HASH JOIN
      TABLE ACCESS BY INDEX ROWID SALES_FCT
        BITMAP CONVERSION TO ROWIDS
          BITMAP AND
            BITMAP MERGE
              BITMAP KEY ITERATION
                BUFFER SORT
                  TABLE ACCESS FULL CHANL_DIM
                    BITMAP INDEX RANGE SCAN SALES_CHANL_BX
                      BITMAP MERGE
                        BITMAP KEY ITERATION
                          BUFFER SORT
                            TABLE ACCESS FULL TIME_DIM
                              BITMAP INDEX RANGE SCAN SALES_TIME_BX
                                ...
                                  TABLE ACCESS FULL CHANL_DIM
                                    TABLE ACCESS FULL TIME_DIM
```

Star Transformation

Large Dimension Table

- Dimension tables are accessed twice
 - once for each phase.
 - can performance issue for big dimension tables (low selectivity)
- Oracle might decide to use temporary table
 - instead of accessing the same dimension table twice

```
LOAD AS SELECT      SYS_TEMP_0FD9D6720_BEBDC
TABLE ACCESS FULL   CUSTOMERS
...
filter("C"."CUST_STATE_PROVINCE"='CA')
```

- Can be disabled on session

```
ALTER SESSION SET star_transformation_enabled = temp_disable;
```

OR Expansion CONCATENATION

Hint: `USE_CONCAT`
Opp.: `NO_EXPAND`

- Concatenates the rows returned by two or more row sets.
- Works like **UNION ALL** but removes duplicate rows.
 - So appends a negation of the previous components using LNNVL function
 - filter (LNNVL(SAL=2)) returns all rows for which SAL != 2 or SAL is NULL.

```
SELECT /*+ USE_CONCAT */ * FROM emp WHERE deptno = 1 OR sal = 2;
```


	0		SELECT STATEMENT				8		696	
	1		CONCATENATION							
	2		TABLE ACCESS BY INDEX ROWID		EMP		4		348	
	3		INDEX RANGE SCAN		I_SAL		2			
	4		TABLE ACCESS BY INDEX ROWID		EMP		4		348	
	5		INDEX RANGE SCAN		I_DEPTNO		2			

3 - access("SAL"=2)
4 - filter(**LNNVL**("SAL"=2))
5 - access("DEPTNO"=1)

Preventing OR Expansion

INLIST ITERATOR

- When many values are used inside IN clause or between OR operators
- Works as a FOR LOOP statement in PL/SQL and similar to NESTED LOOPS
- Alternative for UNION ALL of each value or a FILTER against all the rows
- Optimizer use it when selective index exists for that column rather than CONCATENATION or UNION ALL because it is more efficient.



```
SELECT * FROM emp WHERE empno IN (7876, 7900, 7902);
```

OPERATION	OBJECT_NAME	OPTIONS	COST	CARDINALI...
SELECT STATEMENT			2	3
INLIST ITERATOR				
TABLE ACCESS	EMP	BY INDEX RO...	2	3
INDEX	PK_EMP	UNIQUE SCAN	1	3
Access Predicates				
OR				
EMPNO=7876				
EMPNO=7900				
EMPNO=7902				

View Merge

Simple

```
Hint: MERGE [ (vw) ]  
Opp.: NO_MERGE [ (vw) ]
```

- View merging
 - View SQL text and query using view is transformed to one query (view disappear)
 - New query have one more effective plan then two separate plans
- Simple view merging is possible if view not contain
GROUP BY, Agg. Functions, DISTINCT, Outer join, MODEL, CONNECT BY, Set operations
NO_MERGE hint
- Simple merging is still possible even view contain
 - view appears on the right side of a semijoin or antijoin.
 - view contains subqueries in the SELECT list.
 - outer query block contains PL/SQL functions.
 - view participates in an outer join,

View Merge

Simple

OPERATION	OBJECT_NAME	CARD...
SELECT STATEMENT		3
TABLE ACCESS (BY INDEX ROWID)	EMP	3
INDEX (RANGE SCAN)	I_DEPTNO	3
Access Predicates		
DEPTNO=10		

- Created view can be merged into query.

```
CREATE VIEW vw AS SELECT /*+ MERGE */ deptno, sal FROM emp;
SELECT * FROM vw WHERE deptno = 10;
```

- Inline view also can be merged

```
SELECT /*+ NO_MERGE(i_vw) or MERGE(...) */ e.first_name, locs_v.street_address FROM employees e,
(SELECT d.department_id, d.department_name, l.street_address
FROM departments d, locations l WHERE d.location_id = l.location_id) i_vw
WHERE dept_locs_v.department_id = e.department_id
AND e.last_name = 'Smith';
```

- Not merged – VIEW Operation

Id	Operation	Name	Cost
0	SELECT STATEMENT		7
* 1	HASH JOIN		7
2	TABLE ACCESS BY INDEX ROWID	EMPLOYEES	2
* 3	INDEX RANGE SCAN	EMP_NAME_IX	1
4	VIEW		5
* 5	HASH JOIN		5
6	TABLE ACCESS FULL	LOCATIONS	2
7	TABLE ACCESS FULL	DEPARTMENTS	2

- Merged view – 3 tables join

Id	Operation	Name	Cost (%CPU)
0	SELECT STATEMENT		4 (0)
1	NESTED LOOPS		
2	NESTED LOOPS		4 (0)
3	NESTED LOOPS		3 (0)
4	TABLE ACCESS BY INDEX ROWID	EMPLOYEES	2 (0)
* 5	INDEX RANGE SCAN	EMP_NAME_IX	1 (0)
6	TABLE ACCESS BY INDEX ROWID	DEPARTMENTS	1 (0)
* 7	INDEX UNIQUE SCAN	DEPT_ID_PK	0 (0)
* 8	INDEX UNIQUE SCAN	LOC_ID_PK	0 (0)
9	TABLE ACCESS BY INDEX ROWID	LOCATIONS	1 (0)

Plan Transformations

View Merge

Complex

- Optimizer **merges** complex views
 - containing **GROUP BY** and **DISTINCT**
- **Unable** to perform when:
 - **outer** query tables do **not have** a **ROWID** or **unique column**
 - view appears in a **CONNECT BY** query block
 - view contains **GROUPING SETS**, **ROLLUP**, or **PIVOT** clauses
 - view or outer query block contains the **MODEL** clause

```
SELECT c.cust_id, c.cust_first_name, c.cust_email
FROM cust_dim c, prod_dim p,
     (SELECT DISTINCT s.cust_id, s.prod_id
      FROM sales_fct f) cust_prod_ivw
WHERE c.cntry_id = 52790
      AND c.cust_id = cust_prod_v.cust_id
      AND cust_prod_ivw.prod_id = p.prod_id
      AND p.prod_name = 'T3 Faux Fur-Trimmed
Sweater';
```

Id	Operation	Name
0	SELECT STATEMENT	
1	VIEW	VM_NWVW_1
2	HASH UNIQUE	
* 3	HASH JOIN	
* 4	HASH JOIN	
* 5	TABLE ACCESS FULL	PROD_DIM
6	TABLE ACCESS FULL	SALES_FCT
* 7	TABLE ACCESS FULL	CUST_DIM


Not Mergable View

Filter Push Down

```
CREATE VIEW v2 AS SELECT /*+ NO_MERGE */ deptno, sal FROM emp;
```

```
SELECT * FROM v2 WHERE deptno = 10;
```

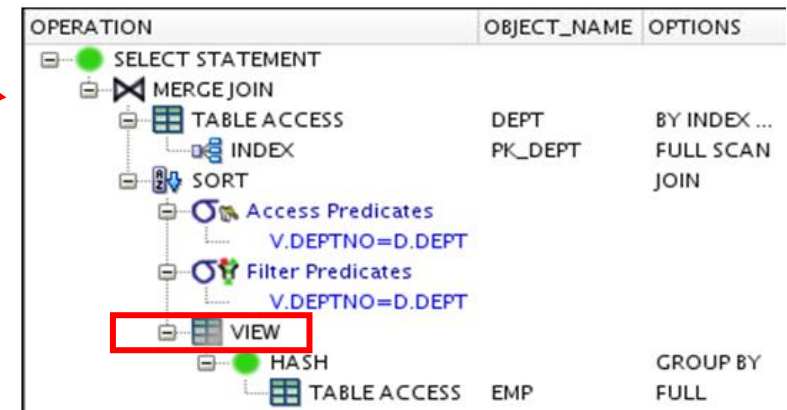
- If not merged appear as VIEW operator
 - executed separately
 - all rows from the view are returned
 - but predicates can be **pushed** or pulled



Nonmergeable inline view

```
SELECT v.*, d.dname  
FROM (SELECT deptno, SUM(sal) sum_sal  
      FROM emp GROUP BY deptno) v,  
      dept d  
WHERE v.deptno = d.deptno;
```

Why this inline view is not mergable?



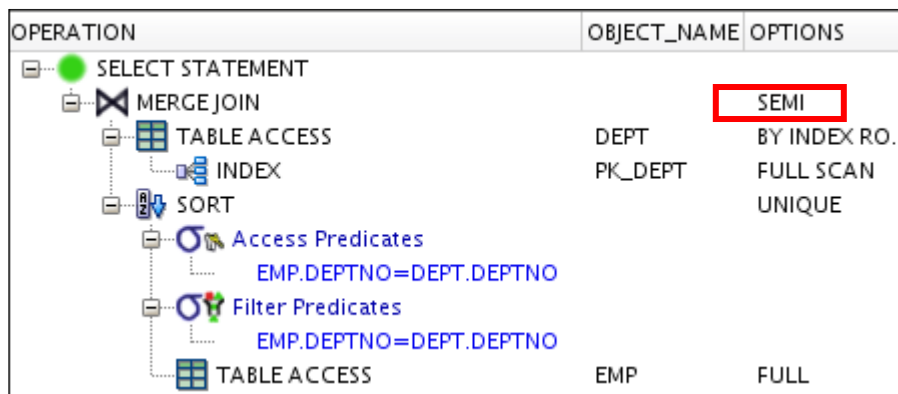
Subquery Unnesting

Semi-Join

Hints: UNNEST MERGE_SJ
HASH_SJ NL_SJ
Opp.: NO_UNNEST

- Transforming an **EXISTS** or **IN subquery** into a join
- Look only for the first match - return rows without duplicating rows
 - even subquery or normal join returns duplicates
- For each DEPT record, only the first matching EMP record is returned
- Merge, nested loops or hash semi joins can be used.

```
SELECT deptno, dname
FROM dept
WHERE EXISTS (SELECT 1 FROM emp WHERE emp.deptno = dept.deptno);
```




Subquery Unnesting

Anti-Join

Hints: UNNEST MERGE_AJ
HASH_AJ NL_AJ
Opp.: NO_UNNEST

- Transforming an **NOT EXISTS** or **NOT IN** subquery into a join
- Similar to Semi-Join but
 - Reverse of what would have been returned by a join

```
SELECT deptno, dname FROM dept
WHERE deptno IS NOT NULL
AND deptno NOT IN
(SELECT /*+ HASH_AJ */ deptno
FROM emp WHERE deptno IS NOT NULL);
```



OPERATION	OBJECT_NAME	OPTIONS
SELECT STATEMENT		
HASH JOIN		ANTI
Access Predicates		DEPTNO=DEPTNO
TABLE ACCESS	DEPT	FULL
TABLE ACCESS	EMP	FULL
Filter Predicates		DEPTNO IS NOT NULL

Cost Based Transform

Group by Placement

Hints: `PLACE_GROUP_BY(@qb)`
 Opp.: `NO_PLACE_GROUP_BY(@qb)`

- Aggregate table before joining if volume from this table is large

```
SELECT /*+ qb_name(main)
         no_place_group_by(@main)
       */
       t.fisc_yr_num, SUM(f.sold_amt)
FROM   sh_sales_fct f
       INNER JOIN sh_time_dim t
         ON t.time_id = f.time_id
       GROUP BY t.fisc_yr_num
```

Id	Operation	Name
0	SELECT STATEMENT	
1	HASH GROUP BY	
2	HASH JOIN	
3	PART JOIN FILTER CREATE	:BF0000
4	TABLE ACCESS FULL	SH_TIME_DIM
5	PARTITION RANGE JOIN-FILTER	
6	TABLE ACCESS FULL	SH_SALES_FCT

```
/*+ qb_name(main)
     place_group_by(@main)
  */
```

Id	Operation	Name
0	SELECT STATEMENT	
1	HASH GROUP BY	
2	HASH JOIN	
3	VIEW	VW_GBC_5
4	HASH GROUP BY	
5	PARTITION RANGE ALL	
6	TABLE ACCESS FULL	SH_SALES_FCT
7	TABLE ACCESS FULL	SH_TIME_DIM

Query Block Name / Object Alias

```
1 - SEL$B359EB6E
3 - SEL$8FF313FB / VW_GBC_5@SEL$CBB12FA2
4 - SEL$8FF313FB
6 - SEL$8FF313FB / F@SEL$1
7 - SEL$B359EB6E / T@SEL$1
```

QB_NAME hint

```
SELECT /*+ QB_NAME(QB_MAIN) */
      prod_name, SUM(sold_amt) AS chanl2_amt, chanl3_amt
FROM   sh_sales_fct f2
JOIN   sh_prod_dim p ON f2.prod_id = p.prod_id
JOIN   (SELECT /*+ QB_NAME(QB_SUB) */
        prod_id, SUM(sold_amt) AS chanl3_amt
        FROM   sh_sales_fct WHERE chanl_id = 3
        GROUP BY prod_id) f3
      ON f2.prod_id = f3.prod_id WHERE chanl_id = 2
GROUP BY p.prod_name, chanl3_amt
```

```
dbms_xplan.display_cursor ('976cmcmdpj9pd', format => 'basic +alias +outline');
```

Id	Operation	Name
0	SELECT STATEMENT	
1	HASH GROUP BY	
2	HASH JOIN	
3	VIEW	VW_GBC_9
4	HASH GROUP BY	
5	HASH JOIN	
6	VIEW	
7	HASH GROUP BY	
8	PARTITION RANGE ALL	
9	TABLE ACCESS FULL	SH_SALES_FCT
10	PARTITION RANGE ALL	
11	TABLE ACCESS FULL	SH_SALES_FCT
12	TABLE ACCESS FULL	SH_PROD_DIM

Query Block Name / Object Alias

```
-----
1 - SEL$2F9F79D7
3 - SEL$C6F03CC3 / VW_GBC_9@SEL$CA133882
4 - SEL$C6F03CC3
6 - QB_SUB      / F3@SEL$2
7 - QB_SUB
9 - QB_SUB      / SH_SALES_FCT@QB_SUB
11 - SEL$C6F03CC3 / F2@SEL$1
12 - SEL$2F9F79D7 / P@SEL$1
```

Outline Data

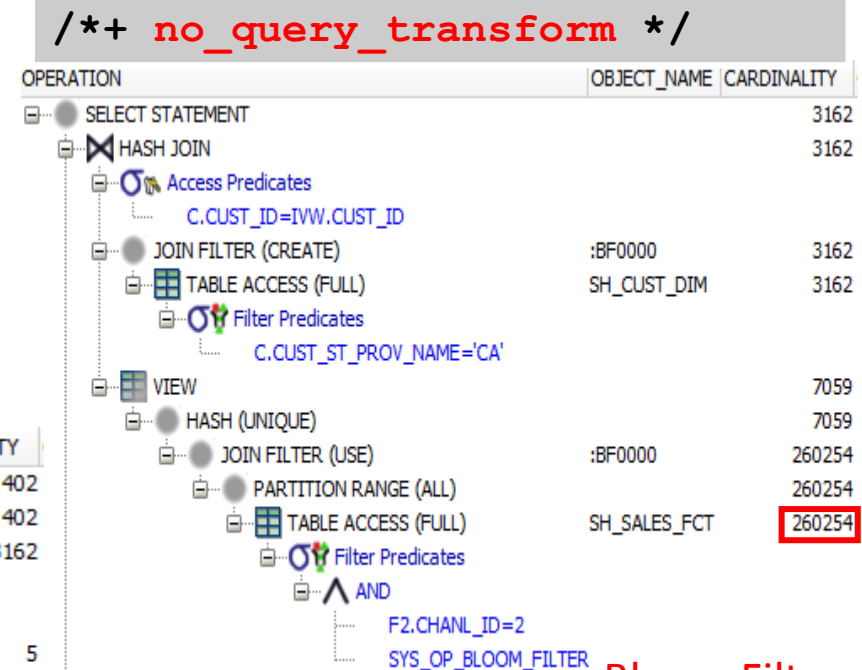
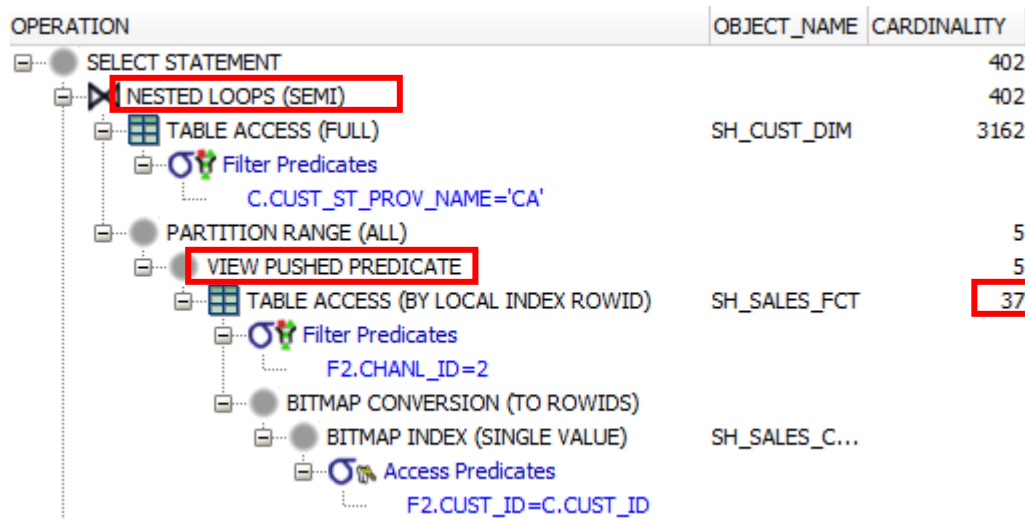
```
-----
/*+
  BEGIN_OUTLINE_DATA
  ...
  MERGE(@"SEL$58A6D7F6")
  OUTLINE(@"QB_MAIN")
  OUTLINE(@"SEL$58A6D7F6")
  MERGE(@"SEL$1")
  OUTLINE(@"SEL$2")
  OUTLINE(@"SEL$1")
  FULL(@"SEL$2F9F79D7" "P"@"SEL$1")
  LEADING(@"SEL$2F9F79D7"
  ...
  END_OUTLINE_DATA
*/
```

Join Predicate Pushdown

Hints: `PUSH_PRED (vw)`
Opp.: `NO_PUSH_PRED (vw)`

- View to be joined with nested-loop and pushed filter from outside

```
SELECT /*+ push_pred(ivw) */
       c.cust_last_name, c.cust_city_name
FROM   sh_cust_dim c,
       (SELECT DISTINCT f2.cust_id
        FROM sh_sales_fct f2
        WHERE f2.chanl_id = 2) ivw
WHERE  c.cust_st_prov_name = 'CA'
       AND c.cust_id = ivw.cust_id
```




Bloom Filter

Q & A

+ Quiz

Plan Transformations Workshop



- Check execution time and plan for star query case
 - Compare second variant by uncomment second -SELECT
- Compare SQL plans same way for next few cases

Topic Agenda

SQL Plan Management

- Overview
- Using SQL Profiles
- SQL Plan Baselines
- Documented 11.2 Hints

Plan Management

Overview

- Use **hints** embedded in SQL source code only **if**
 - Each statement execution uses **different sql_id**
 - Examples: dynamic SQL used, reports with substitution variables
 - To avoid this you can use
 - * bind variables in source code
 - * CURSOR_SHARING parameter
- Use SQL **Profile** where you need only SQL **plan correction**
- Use SQL Plan **Baseline** if
 - Need execution **plan stability**
 - Many good validated plans exist for SQL
 - To **avoid surprise** with new **wrong plan**
- SQL Profile and SQL Plan Baseline
 - Requires access from DBA (to select_catalog_role, ADMINISTER SQL MANAGEMENT OBJECT, dbms_sqltune. ADVISOR)
 - Need same sql_id for all executions of one SQL
 - **Can apply hints without source code modification**

SQL Profile - Accepted - SQL Tuning Advisor - Example

```
DECLARE
l_task_name VARCHAR2(30);
l_sql_ftext  VARCHAR2(32000) :=
'SELECT /*+ ORDERED USE_NL(c) FULL(c) FULL(s) */ COUNT(*)
  FROM sh_sales_fct s, sh_cust_dim c
 WHERE c.cust_id = s.cust_id
 AND c.cust_first_name = :cname
 AND time_id > :tid
 ORDER BY time_id';
l_binds sql_binds := sql_binds(
  anydata.ConvertVARCHAR2('Dina'),
  anydata.ConvertDATE(SYSDATE-100000)
);
BEGIN
  l_task_name := dbms_sqltune.create_tuning_task(
    sql_text      => l_sql_ftext,
    bind_list     => l_binds,
    user_name     => 'SH',
    scope         => 'COMPREHENSIVE',
    time_limit    => 60,
    task_name     => 'SH SQL-tune task 1');
END;
/
```

```
SELECT task_id, task_name
FROM dba_advisor_log
WHERE owner = 'SH';
```

TASK_ID	TASK_NAME
5744	SH SQL-tune task 1

FINDINGS SECTION (2 findings)

1- SQL Profile Finding

Recommendation (estimated benefit: 99.87%)

	Original Plan	With SQL Profile	% Improved
Completion Status:	PARTIAL	COMPLETE	
Elapsed Time(us):	14370351	1938775	86.5 %
CPU Time(us):	14343750	171875	98.8 %
User I/O Time(us):	0	175859	
Buffer Gets:	4778546	6159	99.87 %

2- Index Finding

Recommendation (estimated benefit: 78.42%)

create index ... on SH_SALES_FCT("CUST_ID");

```
dbms_sqltune.set_tuning_task_parameter(task_name=> 'SH SQL-tune task 1', parameter=> 'TIME_LIMIT', value=> 300)
dbms_sqltune.execute_tuning_task(task_name=> 'SH SQL-tune task 1')
SELECT dbms_sqltune.report_tuning_task('SH SQL-tune task 1') FROM dual;
dbms_sqltune.accept_sql_profile(task_name=> 'SH SQL-tune task 1', task_owner=> 'SH', replace=> TRUE)
dbms_sqltune.drop_tuning_task(task_name=> 'SH SQL-tune task 1')
SELECT name, sql_text, category, status FROM dba_sql_profiles;
```

NAME	SQL_TEXT	CATEGORY	STATUS
SYS_SQLPROF_015213a88b700000	SELECT /*+ ORDERED USE_NL(c)...	DEFAULT	ENABLED

Plan Management

SQL Profile - Create Manually - Example

```
set serveroutput on
exec profile_test('SUM(sal)', 10);
```

```
SELECT sql_id, plan_hash_value, sql_text
FROM v$sqlarea WHERE upper(sql_text) LIKE '%DEPTNO = :%'
AND sql_text NOT LIKE '%v$%';
```

SQL_ID	PLAN_HASH_VALUE	CHILD_NUMBER	SQL_TEXT
b12dj8f8ftmdk	1364575507	0	SELECT SUM(sal) FROM emp where deptno = :deptno

```
SELECT * FROM TABLE(dbms_xplan.display_cursor('b12dj8f8ftmdk', 0, format => 'basic +alias +outline')) ;
```

Id	Operation	Name
0	SELECT STATEMENT	
1	SORT AGGREGATE	
2	TABLE ACCESS BY INDEX ROWID	EMP
3	INDEX RANGE SCAN	I_DEPTNO

Id	Operation	Name
0	SELECT STATEMENT	
1	SORT AGGREGATE	
2	TABLE ACCESS FULL	EMP

SQL_ID	PLAN_HASH_VALUE	CHILD_NUMBER
b12dj8f8ftmdk	2083865914	0

```
DECLARE
  in_hints VARCHAR2(100)
    := 'FULL(@"SEL$1" "EMP"@"SEL$1")';
  in_sql_id VARCHAR2(100) := 'b12dj8f8ftmdk';
  l_sql_ftext CLOB;
BEGIN
  SELECT SQL_FULLTEXT INTO l_sql_ftext
  FROM V$SQLAREA
  WHERE SQL_ID = in_sql_id AND ROWNUM <= 1;
  DBMS_SQLTUNE.IMPORT_SQL_PROFILE(
    SQL_TEXT => l_sql_ftext,
    PROFILE => SQLPROF_ATTR(in_hints),
    NAME => 'PROFILE_' || in_sql_id,
    REPLACE => TRUE,
    FORCE_MATCH => TRUE
  );
END;
/
```

Plan Management

SQL Profile Management

- Viewing

```
SELECT name, sql_text, category, status FROM dba_sql_profiles;
```

NAME	SQL_TEXT	CATEGORY	STATUS
PROFILE_b12dj8f8ftmdk	SELECT SUM(sal) FROM emp where deptno = :deptno	DEFAULT	ENABLED

- Dropping

```
exec dbms_sqltune.drop_sql_profile('PROFILE_b12dj8f8ftmdk');
```

- Edit in stage table

```
exec dbms_sqltune.create_stgtab_sqlprof('my_sql_profiles_stage');
```

```
exec dbms_sqltune.pack_stgtab_sqlprof
```

```
('PROFILE_b12dj8f8ftmdk', staging_table_name => 'my_sql_profiles_stage')
```

```
select obj_name, comp_data from my_sql_profiles_stage;
```

OBJ_NAME	COMP_DATA
PROFILE_b12dj8f8ftmdk	<outline_data><hint><![CDATA[FULL(@"SEL\$1" "EMP"")]></hint></outline_data>

Now xml column comp_data can be updated.

```
exec dbms_sqltune.unpack_stgtab_sqlprof
```

```
(replace => TRUE, staging_table_name => 'my_sql_profiles_stage');
```

SQL Plan Baseline

Overview

- This is optional set of one SQL plans
 - Can be create, captured, exported, imported
- Optimizer can choose only from baseline plans in if
 - parameter `OPTIMIZER_USE_SQL_PLAN_BASELINES` must be TRUE (on session)
 - baseline exists with at least one enabled plan
- Used to guarantee known optimizer behavior
- Examples of usage:
 - changing database version,
 - code migration DEV -> QA -> UAT -> PROD
 - problems with deep varying statistics and cardinality
- If SQL profile is accepted then is added to baseline
- (Not like SQL profile) SQL Plan Baseline consists of outlines
 - Outline is set off all hints needed to reproduce particular plan

SQL Plan Baseline

- Create

```
DROP INDEX i_deptno;  
exec profile_test('SUM(sal)', 10);
```

```
SELECT sql_handle, origin, enabled,  
       accepted, fixed, autopurge  
FROM dba_sql_plan_baselines;
```

```
SELECT sql_id, plan_hash_value, operation, options sql_text  
FROM v$sql s JOIN v$sql_plan p using(sql_id,plan_hash_value)  
WHERE upper(sql_text) LIKE '%DEPTNO = :%'  
AND sql_text NOT LIKE '%v$%' AND object_name = 'EMP';
```

SQL_ID	PLAN_HASH_VALUE	OPERATION	OPTIONS	SQL_TEXT
b12dj8f8ftmdk	2083865914	TABLE ACCESS	FULL	SELECT SUM(sal) FROM emp where deptno = :deptno

```
BEGIN  
  dbms_output.put_line('Loaded: ' || dbms_spm.load_plans_from_cursor_cache(  
    sql_id => 'b12dj8f8ftmdk', plan_hash_value => '2083865914'));  
END;  
/
```

SQL_HANDLE	ORIGIN	ENABLED	ACCEPTED	FIXED	AUTOPURGE
SYS_SQL_f437939c5404f8eb	MANUAL-LOAD	YES	YES	NO	YES

- Use

```
CREATE INDEX i_deptno ON emp(deptno);  
ALTER SESSION SET optimizer_use_sql_plan_baselines = TRUE;
```

SQL_ID	PLAN_HASH_VALUE	OPERATION	OPTIONS	SQL_TEXT
b12dj8f8ftmdk	2083865914	TABLE ACCESS	FULL	SELECT SUM(sal) FROM emp where deptno = :deptno

```
ALTER SESSION SET optimizer_use_sql_plan_baselines = FALSE;
```

SQL_ID	PLAN_HASH_VALUE	OPERATION	OPTIONS	SQL_TEXT
b12dj8f8ftmdk	2083865914	TABLE ACCESS	FULL	SELECT SUM(sal) FROM emp where deptno = :deptno
b12dj8f8ftmdk	1364575507	TABLE ACCESS	BY INDEX ROWID	SELECT SUM(sal) FROM emp where deptno = :deptno

Documented 11.2 Hints

Access Paths

FULL
CLUSTER
HASH
INDEX and NO_INDEX
INDEX_ASC and INDEX_DESC
INDEX_COMBINE and INDEX_JOIN
INDEX_JOIN
INDEX_FFS and NO_INDEX_FFS
INDEX_SS and NO_INDEX_SS
INDEX_SS_ASC and INDEX_SS_DESC

Join Orders

LEADING
ORDERED

Join Operations

USE_NL and NO_USE_NL
USE_NL_WITH_INDEX
USE_MERGE and NO_USE_MERGE
USE_HASH and NO_USE_HASH
NO_USE_HASH

Application Upgrade

CHANGE_DUPKEY_ERROR_INDEX
IGNORE_ROW_ON_DUPKEY_INDEX
RETRY_ON_ROW_CHANGE

Parallel Execution

PARALLEL and **NO_PARALLEL**
PARALLEL_INDEX and **NO_PARALLEL_INDEX**
PQ_DISTRIBUTE

Query Transformations

NO_QUERY_TRANSFORMATION
USE_CONCAT
NO_EXPAND
REWRITE and **NO_REWRITE**
MERGE and **NO_MERGE**
STAR_TRANSFORMATION and **NO_STAR_TRANSFORMATION**
FACT and **NO_FACT**
UNNEST and **NO_UNNEST**

Additional Hints

APPEND, **APPEND_VALUES**, and **NOAPPEND**
CACHE and **NOCACHE**
PUSH_PRED and **NO_PUSH_PRED**
PUSH_SUBQ and **NO_PUSH_SUBQ**
QB_NAME
CURSOR_SHARING_EXACT
DRIVING_SITE
DYNAMIC_SAMPLING
MODEL_MIN_ANALYSIS

Undocumented 11.2 Hints

CARDINALITY

OPT_ESTIMATE

SWAP_JOIN_INPUTS

NO_SWAP_JOIN_INPUTS

CONNECT_BY_COST_BASED

NO_CONNECT_BY_COST_BASED

NATIVE_FULL_OUTER_JOIN

NO_NATIVE_FULL_OUTER_JOIN

USE_HASH_AGGREGATION

NO_USE_HASH_AGGREGATION

...

Check `v$sql_hint`

Nested Views Hints


Example

```
SELECT /*+ FULL(vw1.vw2.f1) USE_HASH(vw1.vw2, vw1.vw3) */
      prod_name, SUM(sold_amt), chanl3_amt
FROM sh_sales_fct f2
JOIN sh_prod_dim p ON f2.prod_id = p.prod_id
JOIN (SELECT prod_id, SUM(sold_amt) AS chanl3_amt
      FROM (SELECT prod_id, SUM(sold_amt) AS chanl3_amt
            FROM sh_sales_fct f1 WHERE chanl_id = 3
            WHERE cal_yr_num > 2000
            GROUP BY prod_id, chanl_id) vw2
      JOIN (SELECT prod_id, SUM(sold_amt) AS chanl3_amt
            FROM sh_sales_fct WHERE chanl_id = 3
            WHERE cal_yr_num <= 2000
            GROUP BY prod_id, chanl_id) vw3
      GROUP BY prod_id) vw1,
ON f2.prod_id = f3.prod_id WHERE chanl_id = 2
GROUP BY p.prod_name, chanl3_amt
```

Q & A

+ Quiz

Plan Management Workshop



- Create SQL profile using manual and advisor method.
- Test Profile.
- Create SQL plan baseline and test it.
- All statements available in workshop script.

Topic Agenda

Other Performance Tips (not prepared yet)

- Exadata Tips
 - Smart Scan
 - Bloom Filtering
 - Storage Index
 - Flash Cache